

A Machine Learning Based Parallelization Assistant

Aleksandr Maramzin, Christos Vasiladiotis, Roberto Castañeda Lozano,
Björn Franke, Murray Cole

The University of Edinburgh
United Kingdom



AI-SEPS 2019 workshop

A Machine Learning Based Parallelization Assistant

Aleksandr Maramzin, Christos Vasiladiotis, Roberto Castañeda Lozano,
Björn Franke, Murray Cole

“It Looks Like You’re Writing a Parallel Loop”



AI-SEPS 2019 workshop



SPLASH
ATHENS 2019


Presentation Structure

- Problem statement
- How to use the assistant tool
- Assistant internals
- Parallelization of SNU NPB benchmarks

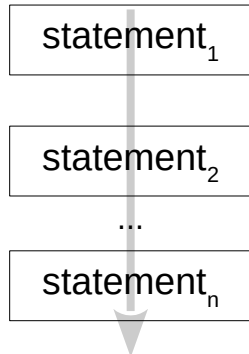


Problem Statement

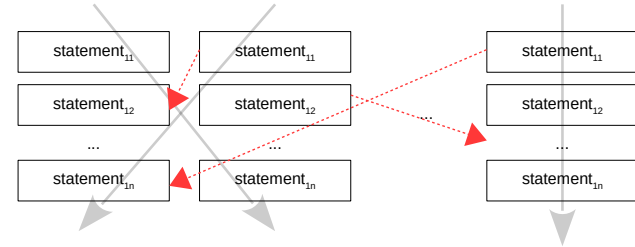
**Parallel Hardware
is Ubiquitous**



**Software
is Sequential**



**Automatic Parallelization
is Limited**



**Manual Parallelization
is Hard**



Problem Statement

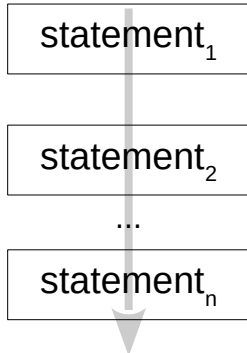
Parallel Hardware
is Ubiquitous

Automatic Parallelization
is Limited


The **assistant solution** we propose alleviates
the process of manual parallelization



Software
is Sequential



Manual Parallelization
is Hard



Problem Statement


NAS Parallel
Benchmarks
(NPB)



**Parallelization Speedup
(Gmean)**

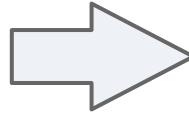

Manual: **1.73x**

Automatic: **0.79x - 1.10x**


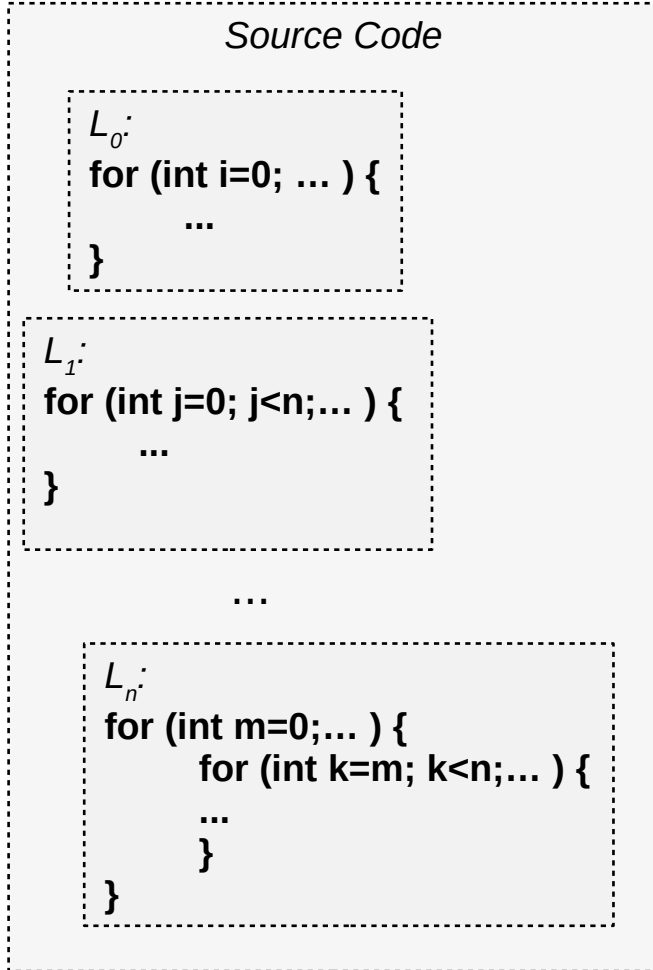


**Programmer still plays the major role
in the hard task of software
parallelization**

Manual Parallelization: Where Should I Start?

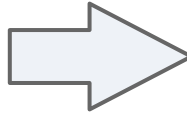
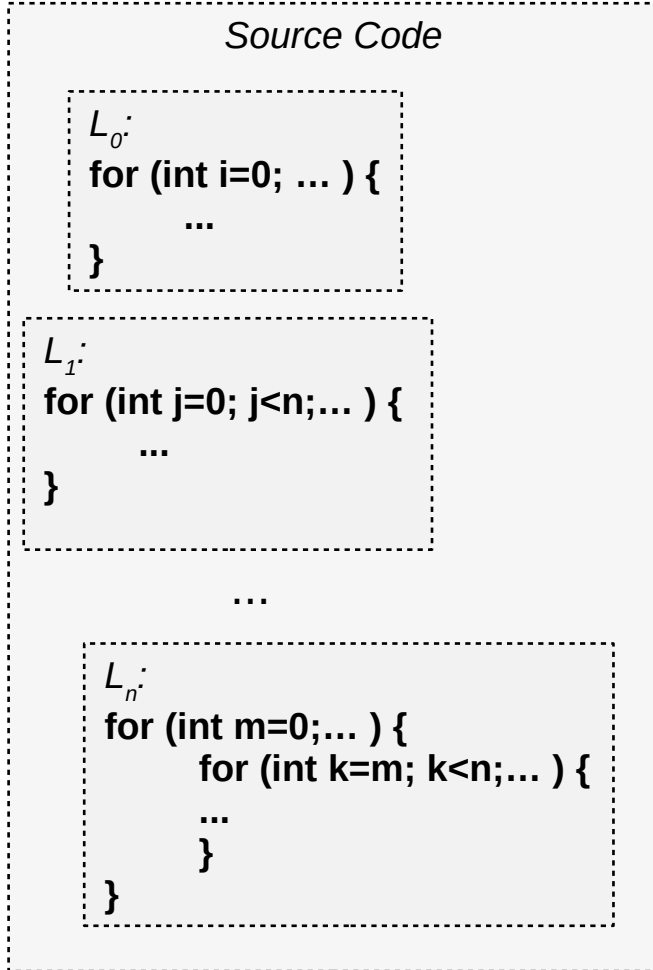


Manual Parallelization: Where Should I Start?



Where Should I Start?

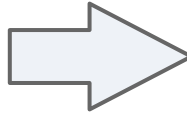
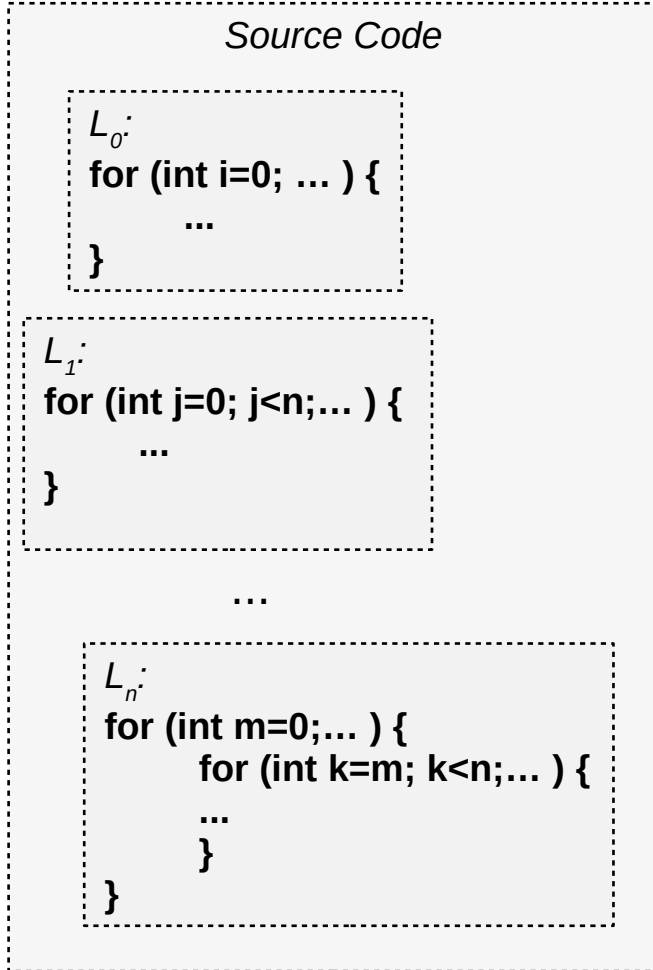
Manual Parallelization: Where Should I Start?



Program Loop Rankings

Source Order
 L_0, L_1, \dots, L_n

Manual Parallelization: Where Should I Start?

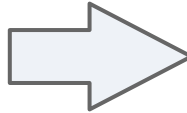
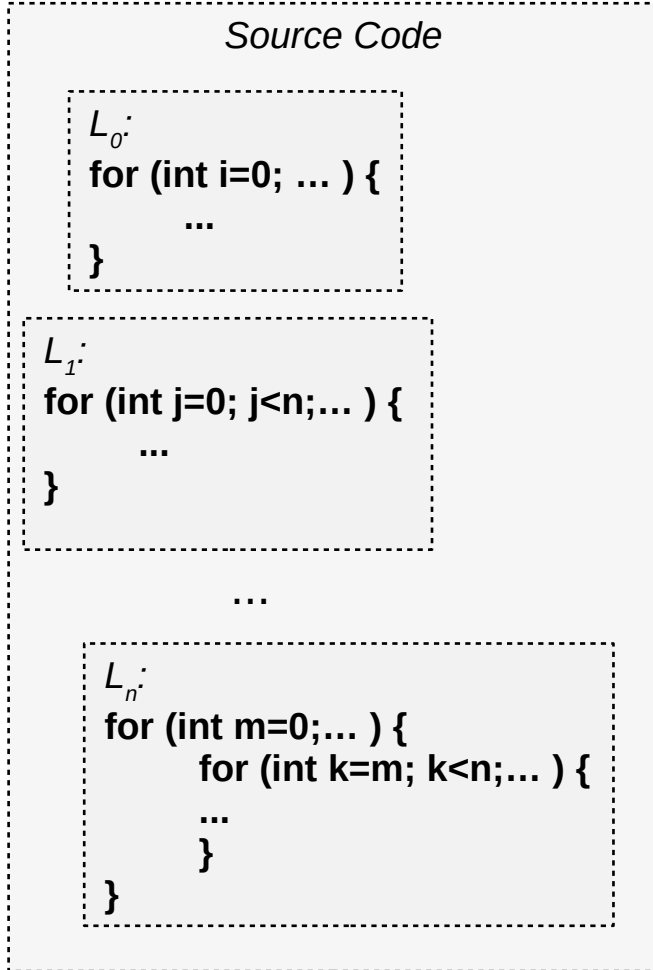


Program Loop Rankings



Source Order
 L_0, L_1, \dots, L_n

Manual Parallelization: Where Should I Start?




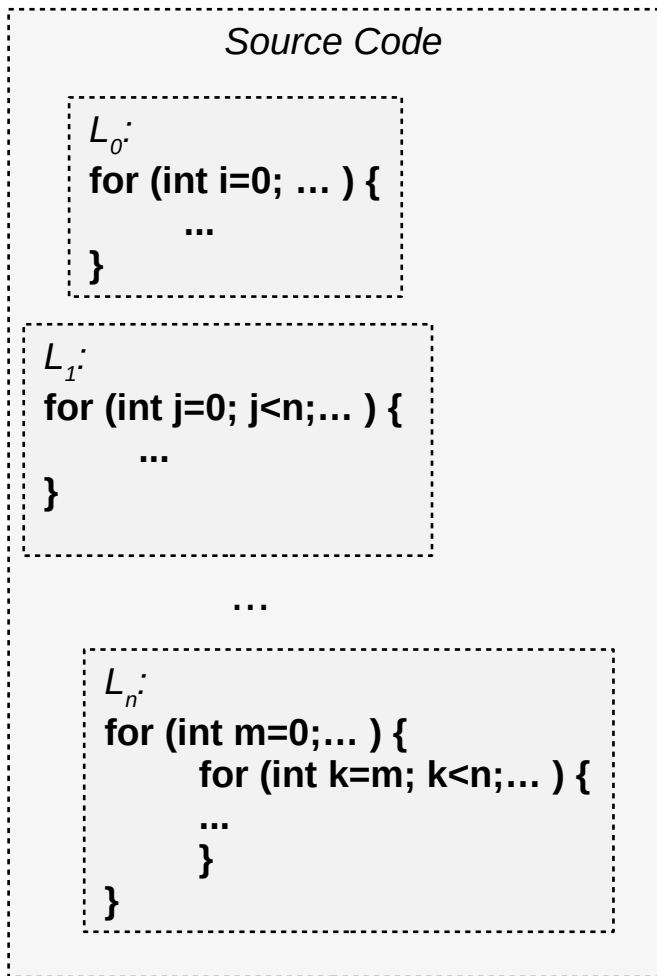
Program Loop Rankings



Source Order
 L_0, L_1, \dots, L_n



Manual Parallelization: Where Should I Start?



Traditional Empirical and
Experimental Method

Program Loop Rankings

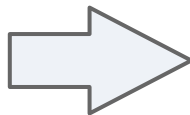
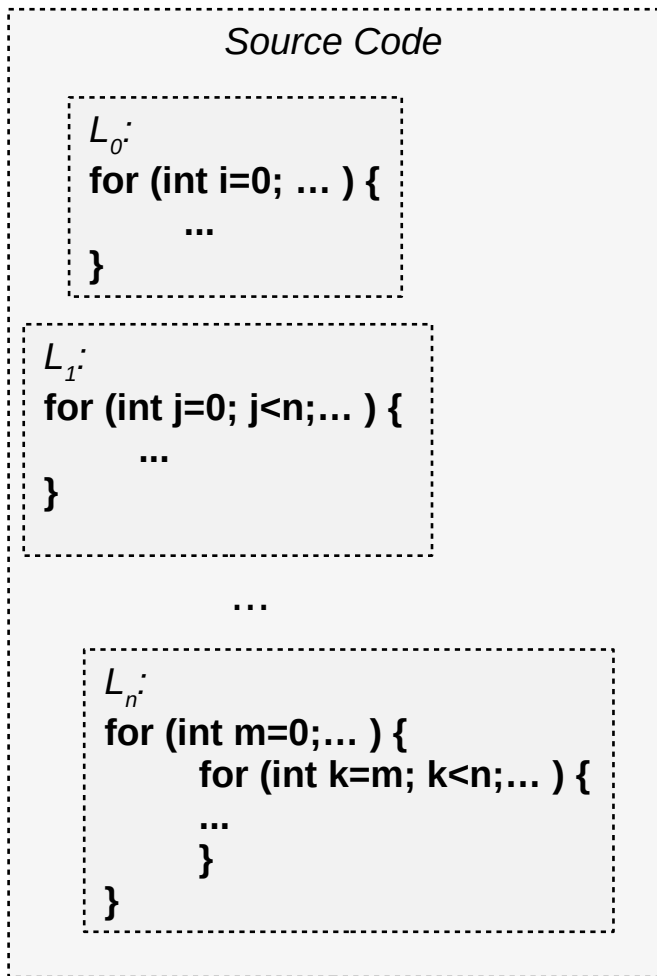


Source Order
 L_0, L_1, \dots, L_n



Profile Order
 $L_{17}, L_{12}, \dots, L_1$

Manual Parallelization: Where Should I Start?



Traditional Empirical and
Experimental Method

Program Loop Rankings



Source Order

L_0, L_1, \dots, L_n





Profile Order

$L_{17}, L_{12}, \dots, L_1$



Manual Parallelization: Where Should I Start?



Traditional Empirical and
Experimental Method

Program Loop Rankings



Source Order

L_0, L_1, \dots, L_n



Profile Order

$L_{17}, L_{12}, \dots, L_1$

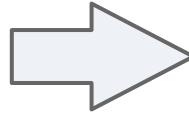



Assistant Order

L_{12}, L_7, \dots, L_3

AI-inspired Method

Manual Parallelization: Where Should I Start?



Traditional Empirical and
Experimental Method

Program Loop Rankings



Source Order
 L_0, L_1, \dots, L_n



Profile Order
 $L_{17}, L_{12}, \dots, L_1$




Assistant Order
 L_{12}, L_7, \dots, L_3




AI-inspired Method

Manual Parallelization: Where Should I Start?



**Final Programmer
Decision**



**Traditional Empirical and
Experimental Method**

AI-inspired Method

Program Loop Rankings



Source Order
 L_0, L_1, \dots, L_n



Profile Order
 $L_{17}, L_{12}, \dots, L_1$



Assistant Order
 L_{12}, L_7, \dots, L_3




Proposed Solution

(ML based software parallelization assistant)




Assistant Training

Assistant Training



1415 loops

Assistant Training



Training Data Set

Parallelizable loops

```
for (int i=...) { ... } [ f1, f2, ..., fn ]  
...  
for (int j=...) { ... } [ f1, f2, ..., fn ]
```

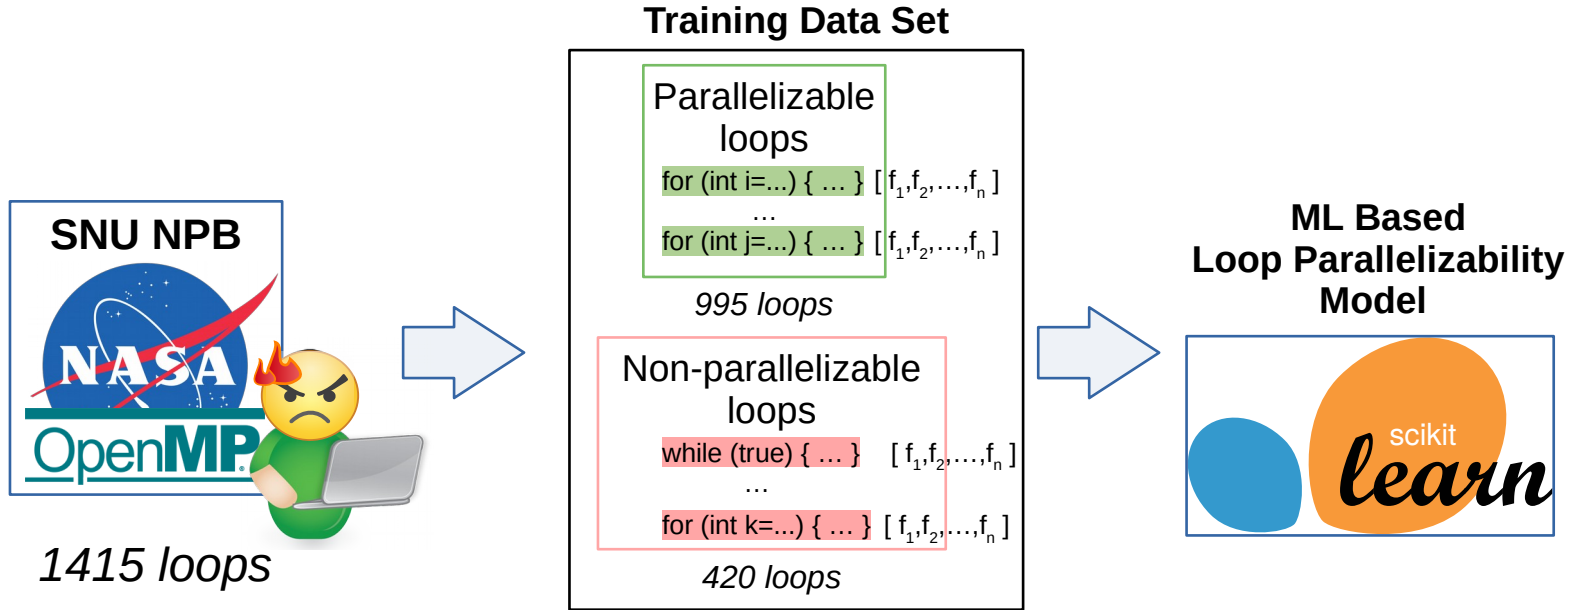
995 loops

Non-parallelizable loops

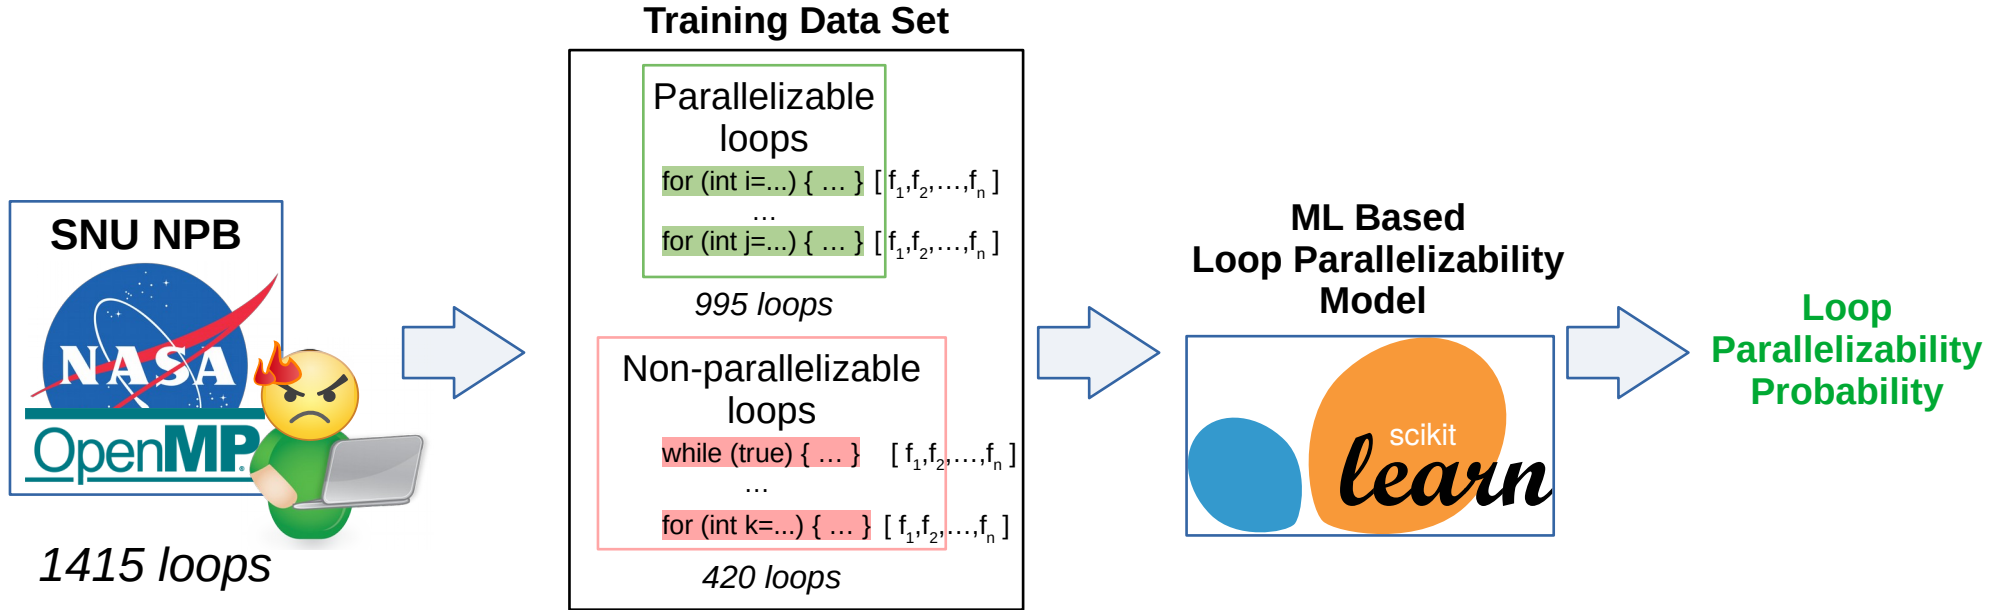
```
while (true) { ... } [ f1, f2, ..., fn ]  
...  
for (int k=...) { ... } [ f1, f2, ..., fn ]
```

420 loops

Assistant Training

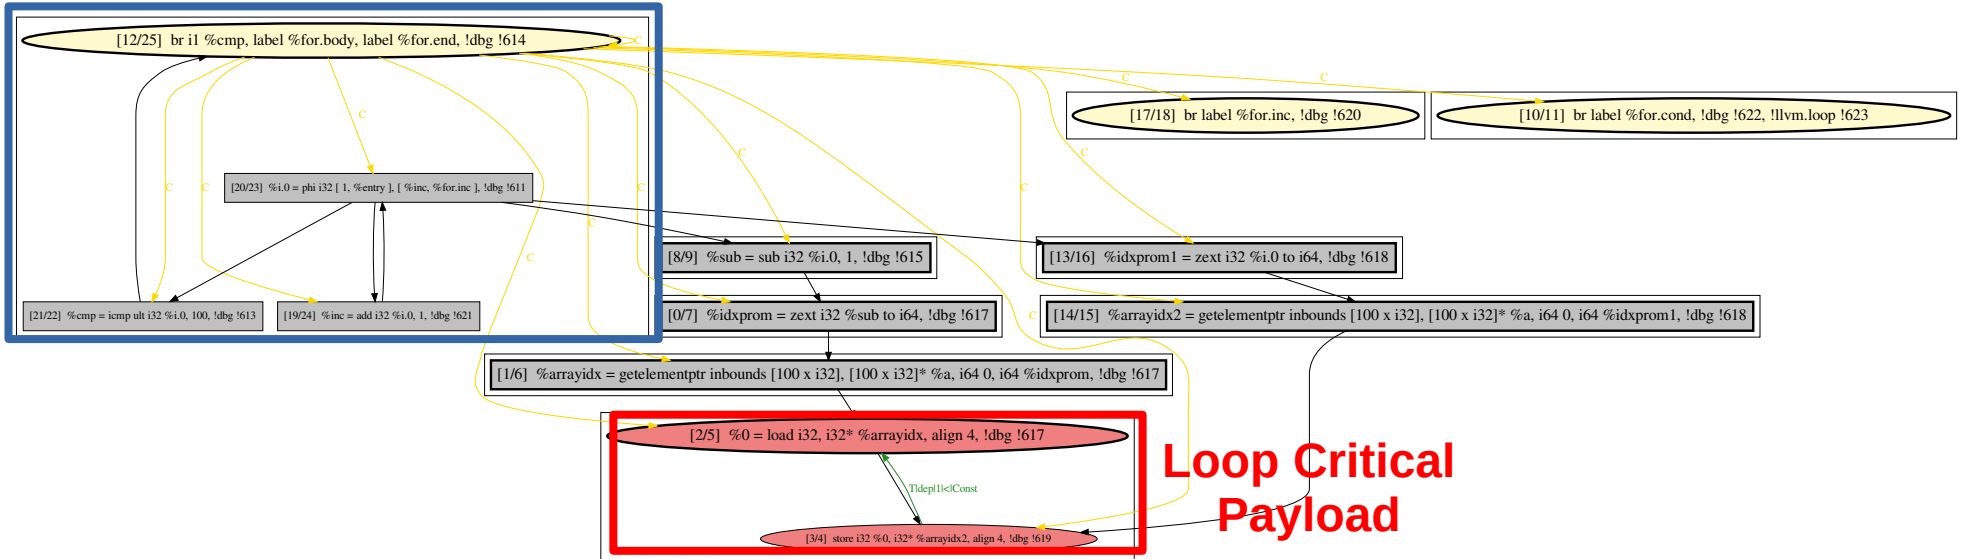


Assistant Training



Loop Features [74]

Loop Iterator



Loop features are based on static structural properties of loop program dependence graphs (PDGs):

- Absolute size
- Loop Iterator/Payload cohesion
- Number of dependence edges
- Instruction types (calls, loads/stores, etc.)
- etc.

Loop Ranking Score

ML Based
Loop Parallelizability
Model



Loop Ranking Score

ML Based
Loop Parallelizability
Model



Not Aware of Loop
Running Time

Loop Ranking Score

Profiler



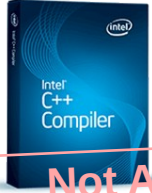
**ML Based
Loop Parallelizability
Model**



**Not Aware of Loop
Running Time**

Loop Ranking Score

Profiler



Not Aware of Loop
Parallelizability

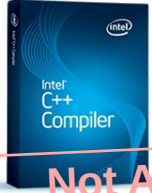
ML Based Loop Parallelizability Model



Not Aware of Loop
Running Time

Loop Ranking Score

Profiler



Not Aware of Loop
Parallelizability

**ML Based
Loop Parallelizability
Model**



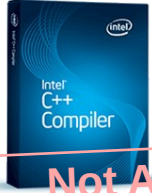
Programmer Effort



Not Aware of Loop
Running Time

Loop Ranking Score

Profiler



Not Aware of Loop
Parallelizability

Performance Gain



ML Based
Loop Parallelizability
Model



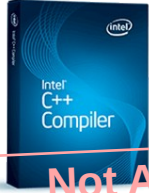
Not Aware of Loop
Running Time

Programmer Effort



Loop Ranking Score

Profiler



Not Aware of Loop Parallelizability

Performance Gain



ML Based
Loop Parallelizability
Model




Not Aware of Loop Running Time

Programmer Effort



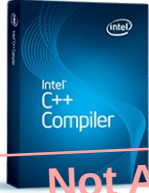
Aware of Loop Parallelizability



Aware of Loop Running Time

Loop Ranking Score

Profiler



Not Aware of Loop Parallelizability

Performance Gain



ML Based
Loop Parallelizability
Model




Not Aware of Loop Running Time

Programmer Effort




Aware of Loop Parallelizability



Aware of Loop Running Time

Improved Loop Ranking



Results



Deployment of our assistant on
SNU NPB benchmarks



Predictive performance of our ML
based loop parallelizability model




AI-SEPS 2019 workshop



Loop Rankings: Profiler vs. Assistant




FT benchmark



Loop Rankings: Profiler vs. Assistant

FT benchmark


Profiler Loop
Ranking




■ parallelizable □ non-parallelizable

Loop Rankings: Profiler vs. Assistant


FT benchmark



Profiler Loop
Ranking




Transformation




■ parallelizable ■ non-parallelizable

Loop Rankings: Profiler vs. Assistant


FT benchmark



Profiler Loop
Ranking



Transformation

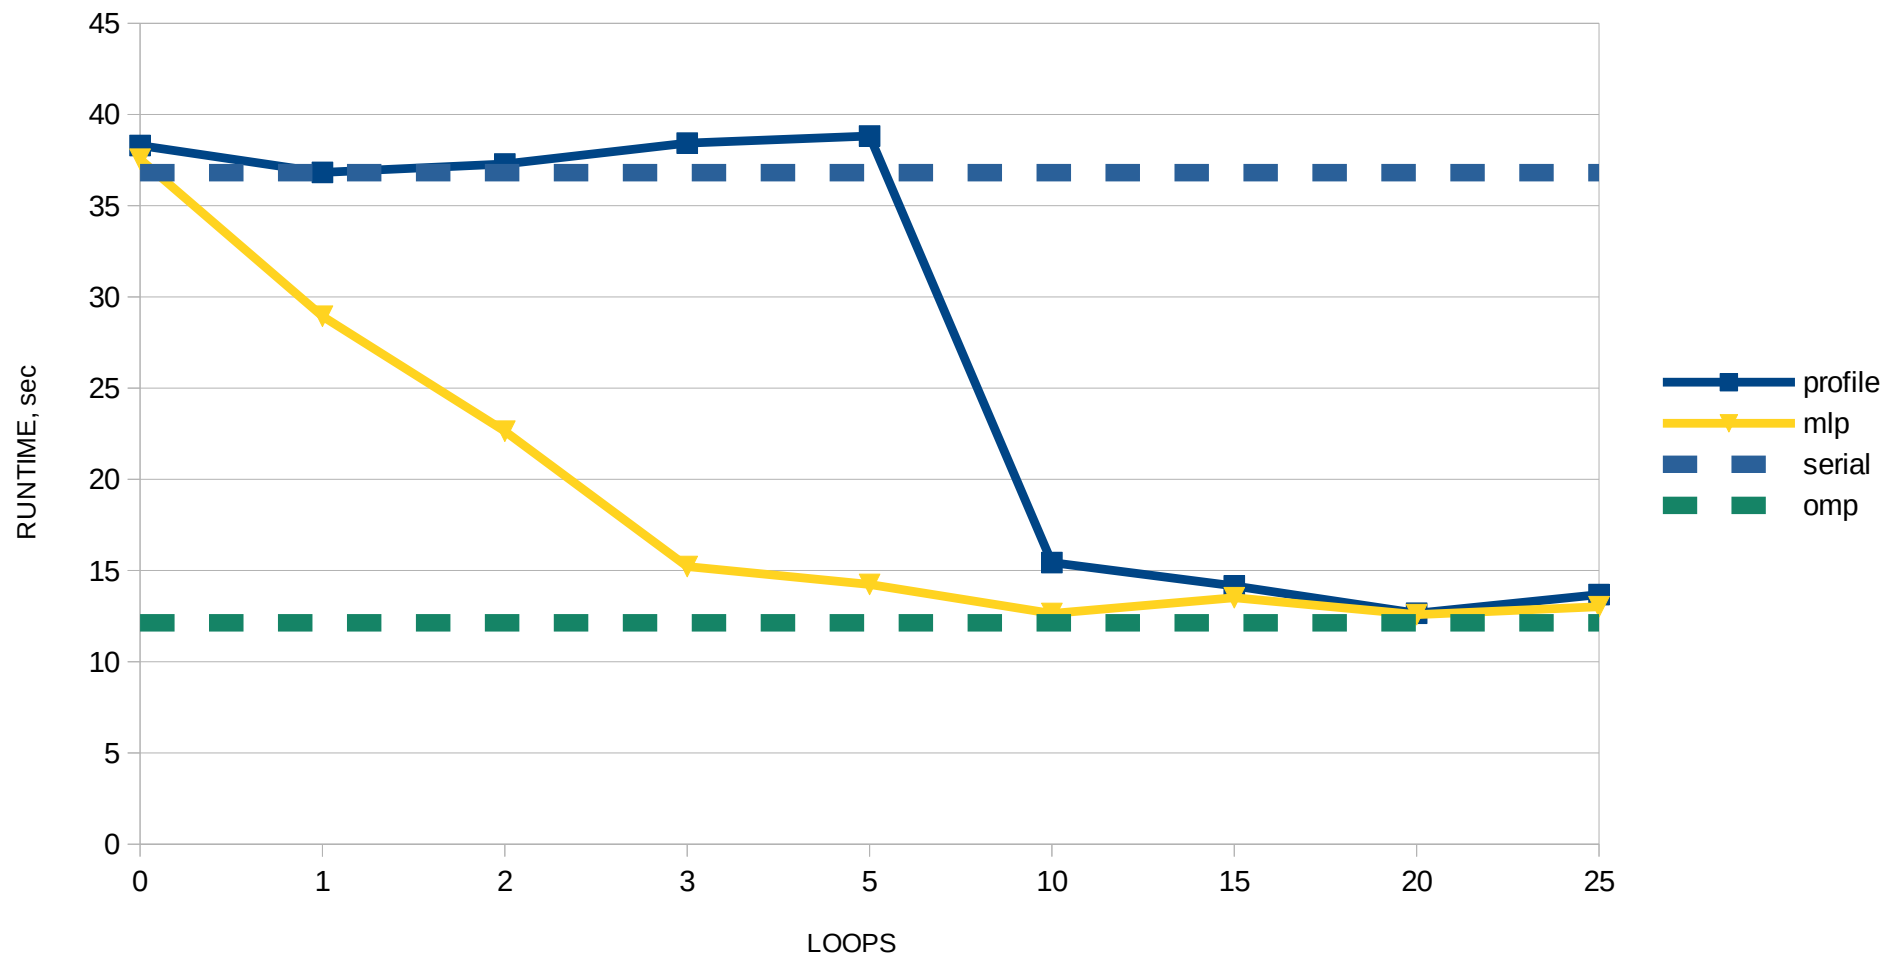


Improved Loop
Ranking

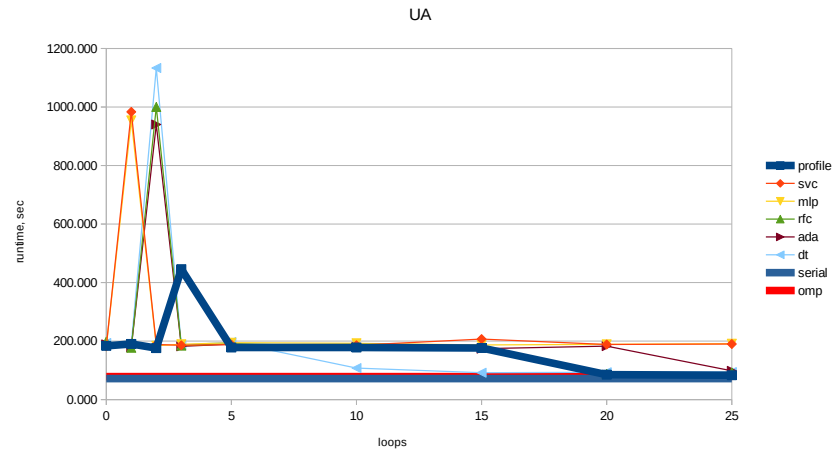
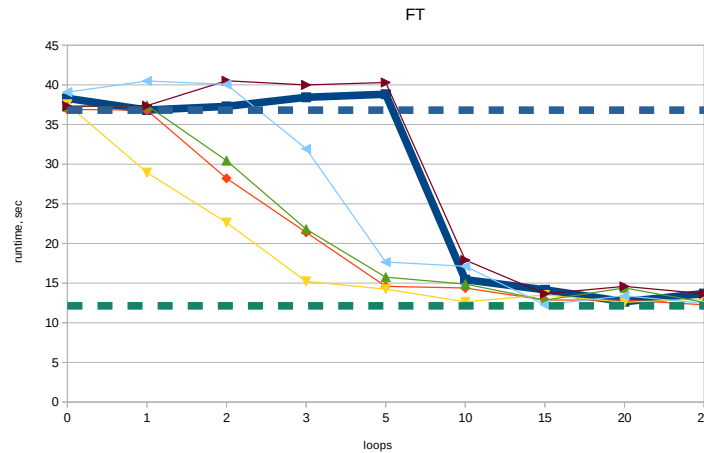
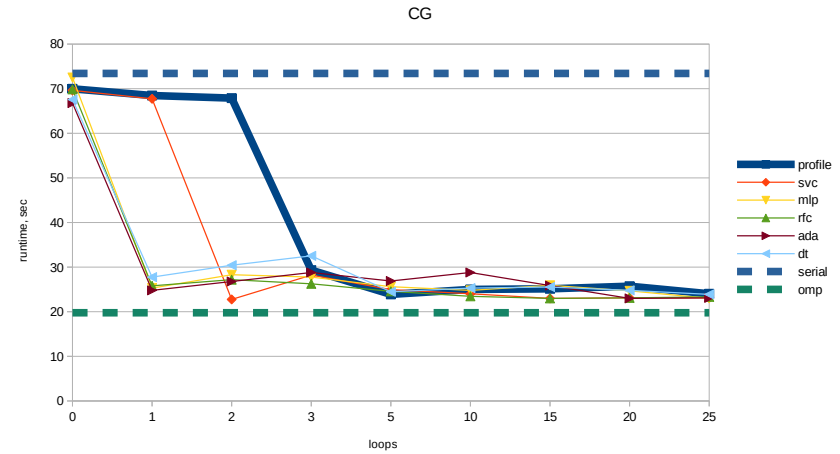
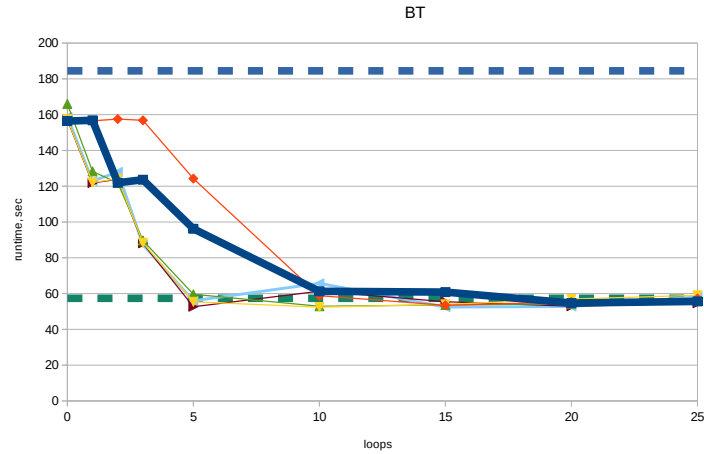
■ parallelizable ■ non-parallelizable

How well do we do?

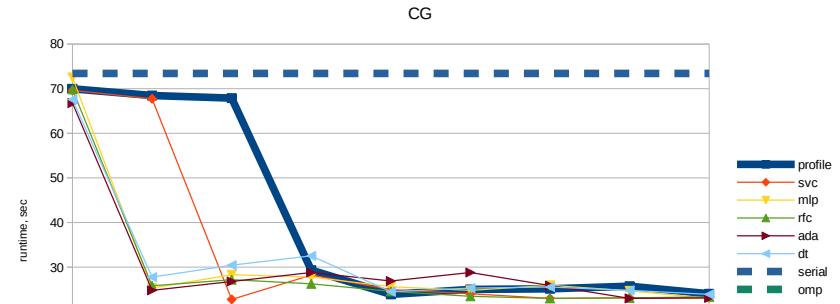
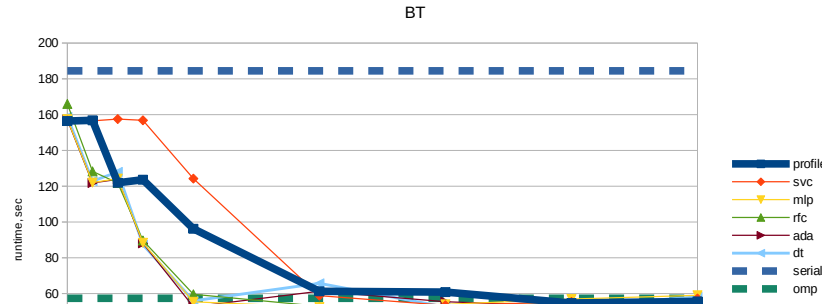
FT



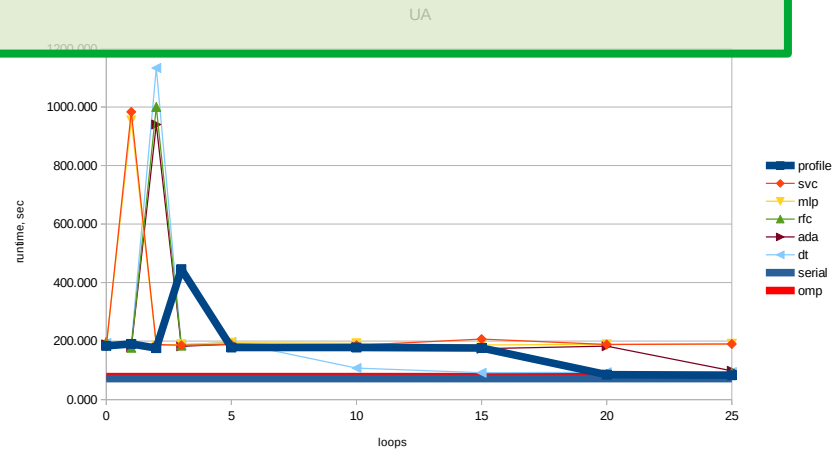
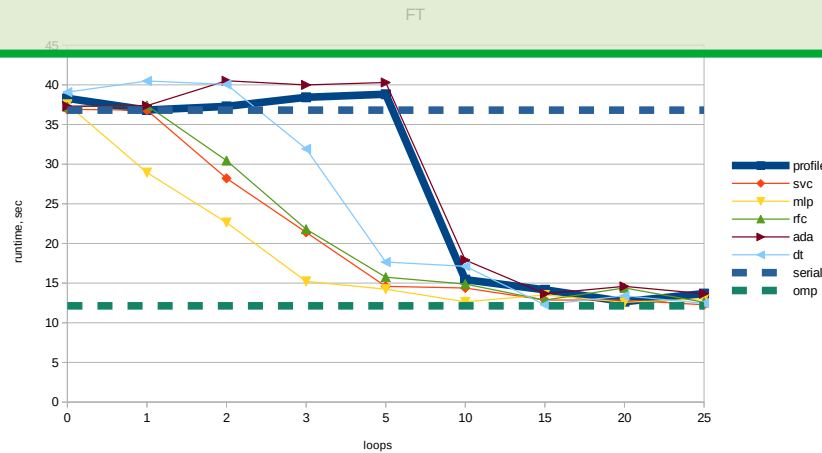
How well do we do?



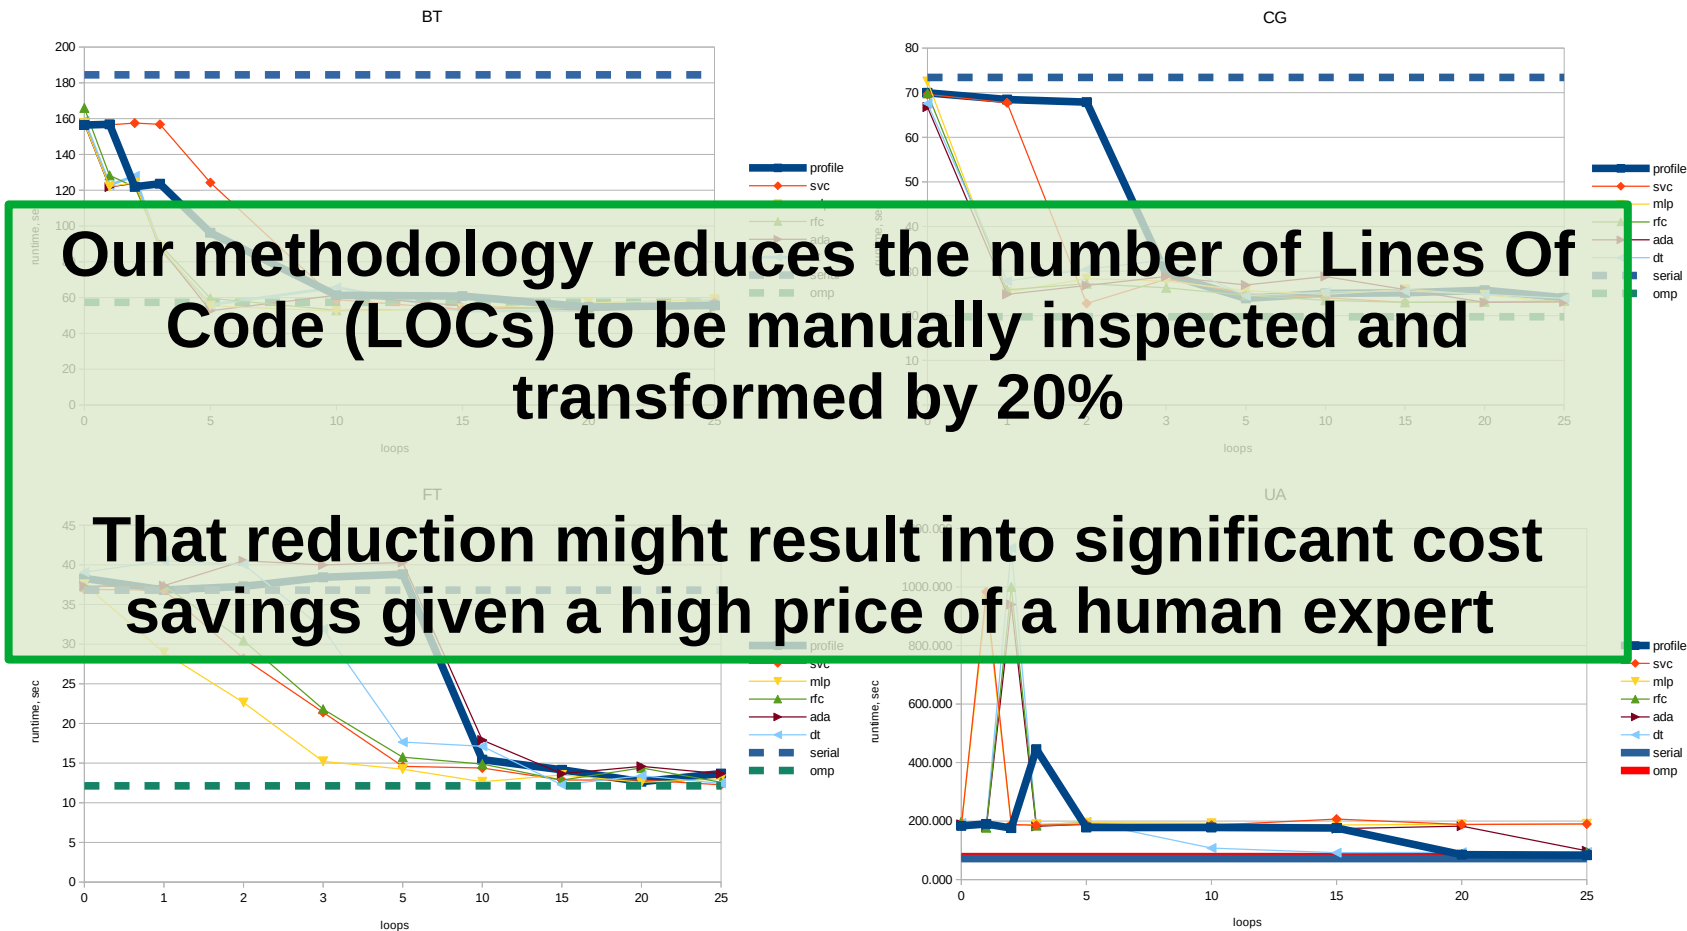
How well do we do?



8 benchmarks in total: 4 wins / 4 draws



How well do we do?



ML Model Predictive Performance

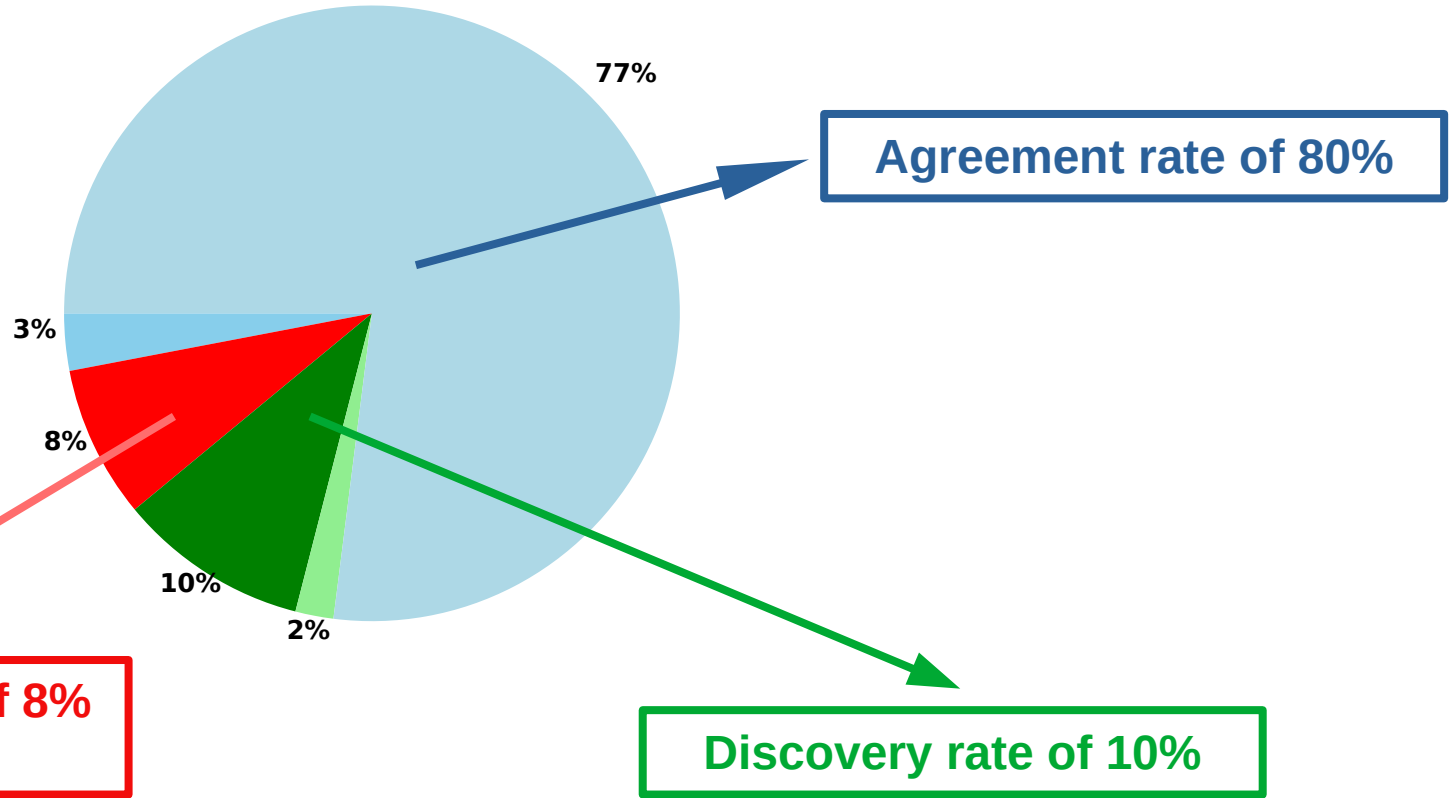
Reference

ML model	accuracy	recall	precision
constant	70.32	100	70.32
uniform	46.27	41.50	69.79
SVC	90.04	95.24	91.06
AdaBoost	86.96	92.92	89.06
DT	84.36	89.57	87.90
RFC	86.65	93.22	88.47
MLP	89.40	93.77	91.39

Around 90% predictive accuracy

ML Model Predictive Performance

ML Based
Loop Parallelizability
Model



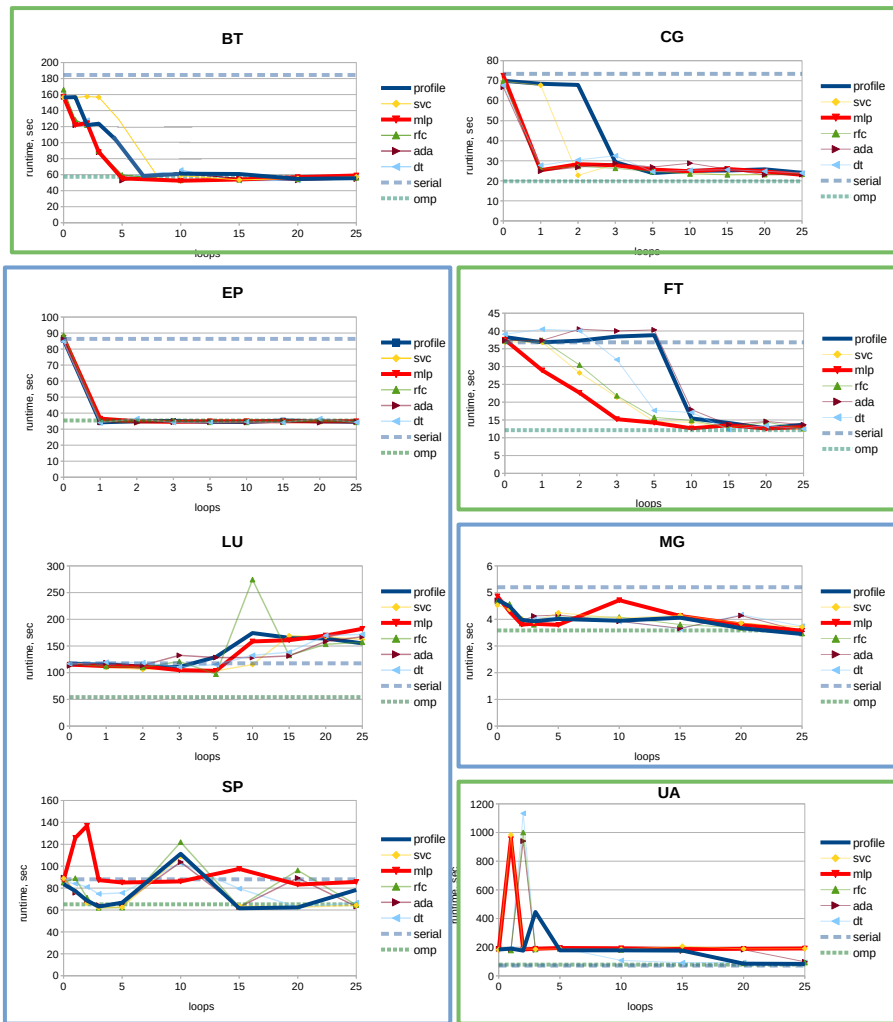
Summary & Conclusions

AI-inspired tool and methodology for manual software parallelization have been designed and implemented as a prototype:

- Loop parallelizability is learnable property
- Designed ML model of loop parallelizability achieves the accuracy of 90%
- The model has been integrated into parallelization assistant
- Deployed against SNU NPB benchmarks our assistant showed a faster parallelization process over traditional profile-guided method: a programmer is required to inspect and transform 20% less Lines Of Code (LOC)



How well do we do?



Motivating example



SNU NPB Conjugate Gradient (CG) benchmark

Profiler's Loop Ranking

- 1 cg.c:326
- 2 cg.c:484
- 3 cg.c:509

VS.

Assistant's Loop Ranking

- 1 cg.c:509
- 2 cg.c:326
- 3 cg.c:484

Assistant's Parallel Loop Probability

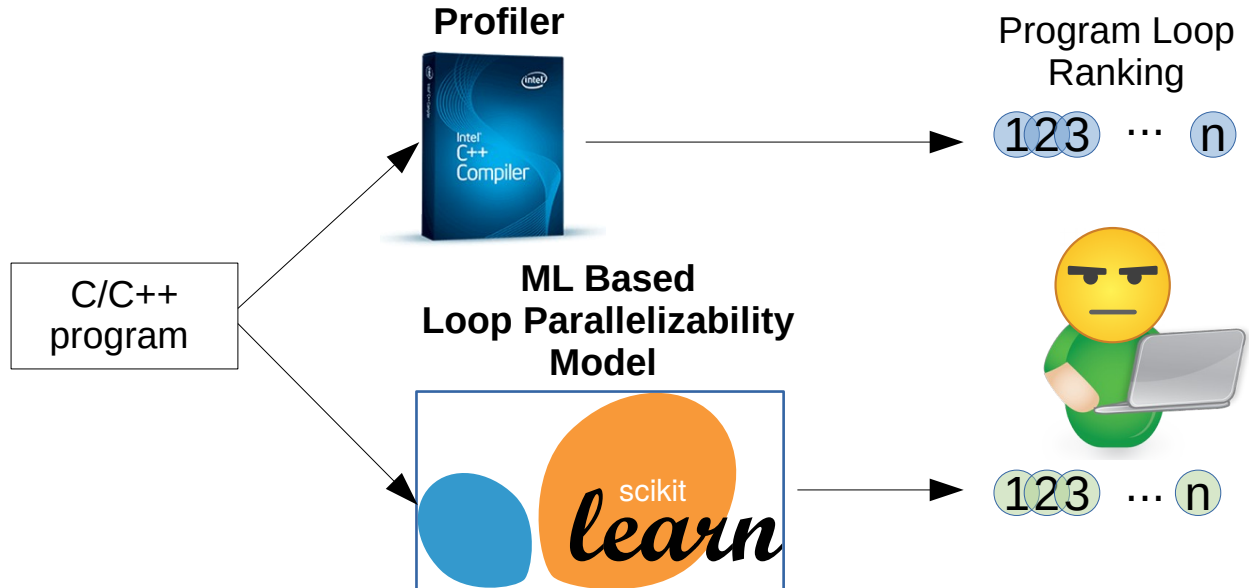
- [85%]
[29%]
[8%]

```
for (j = 0; j < lastrow-firstrow+1; j++) {  
    sum1 = 0.0;  
    for (k = rowstr[j]; k < rowstr[j+1]; k++)  
        sum1 = sum1 + a[k]*p[colidx[k]];  
    q[j] = sum1;  
}
```

cg.c:509


```
for (it = 1; it <= NITER; it++) {  
    ...  
    if (timeron) timer_start(T_conj_grad);  
    conj_grad(colidx,rowstr,x,z,a,p,q,r,&rn timer);  
    if (timeron) timer_stop(T_conj_grad);  
    ...  
    printf("      %5d      %20.14E%20.13f\n", it,  
        rn timer, zeta);  
    ...  
}
```

cg.c:326



Problem Statement 2/2

NAS Parallel Benchmarks (NPB)

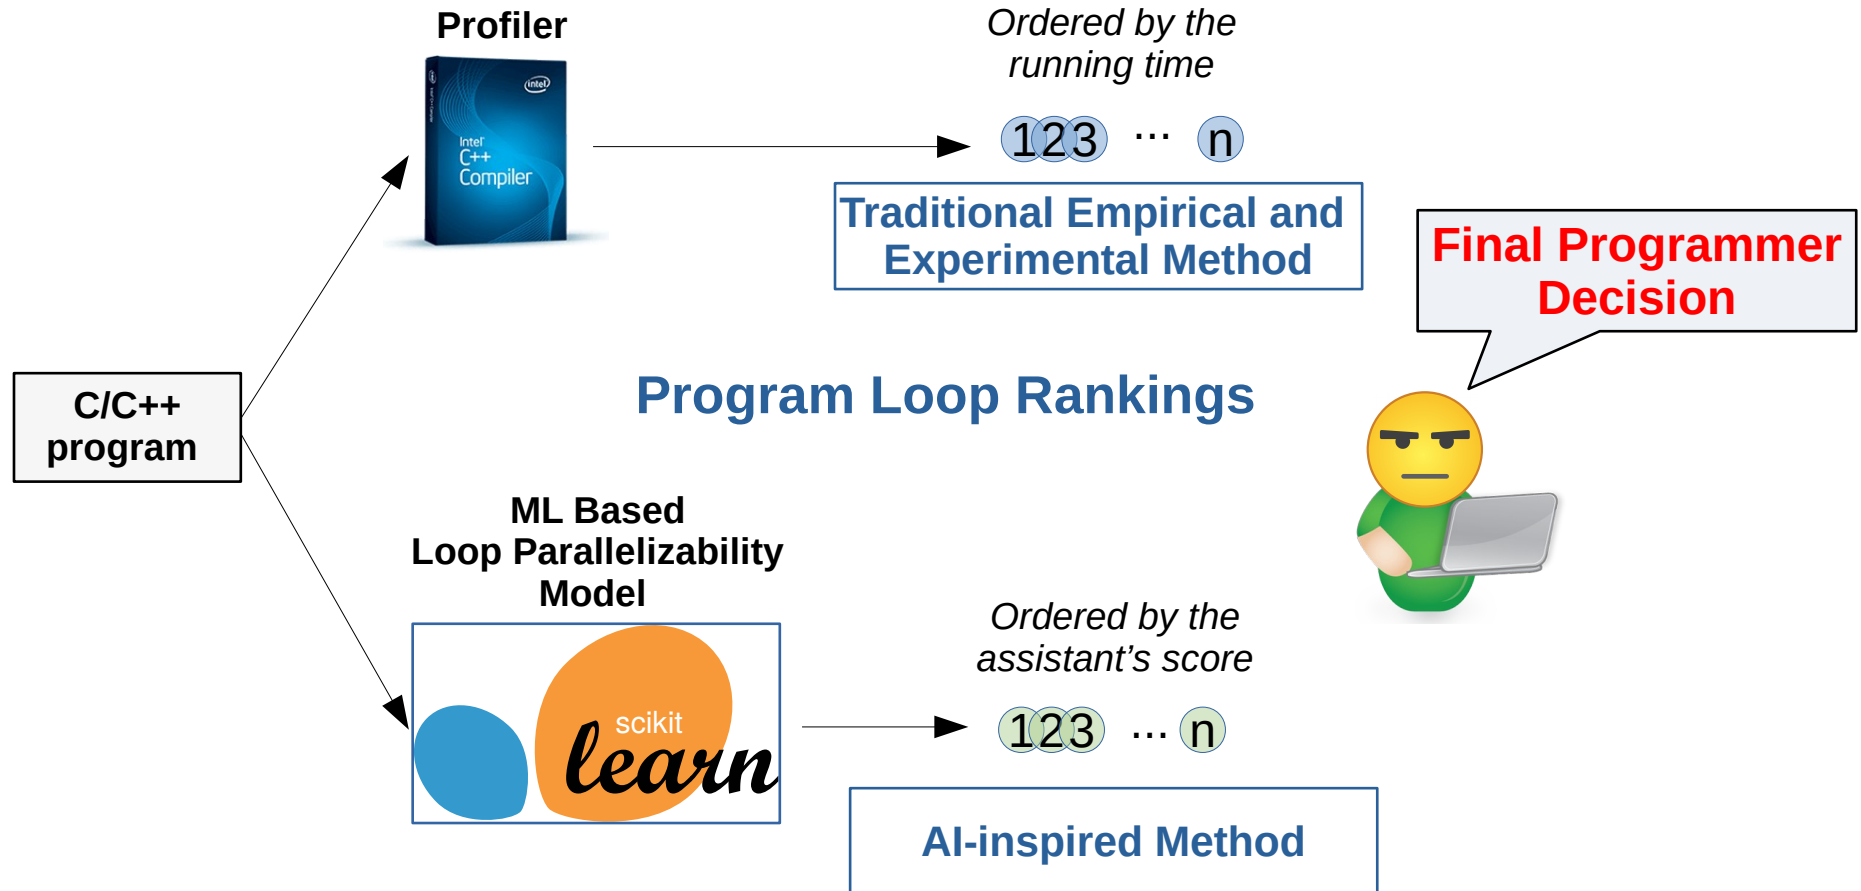


NAS Parallel Benchmarks (NPB)

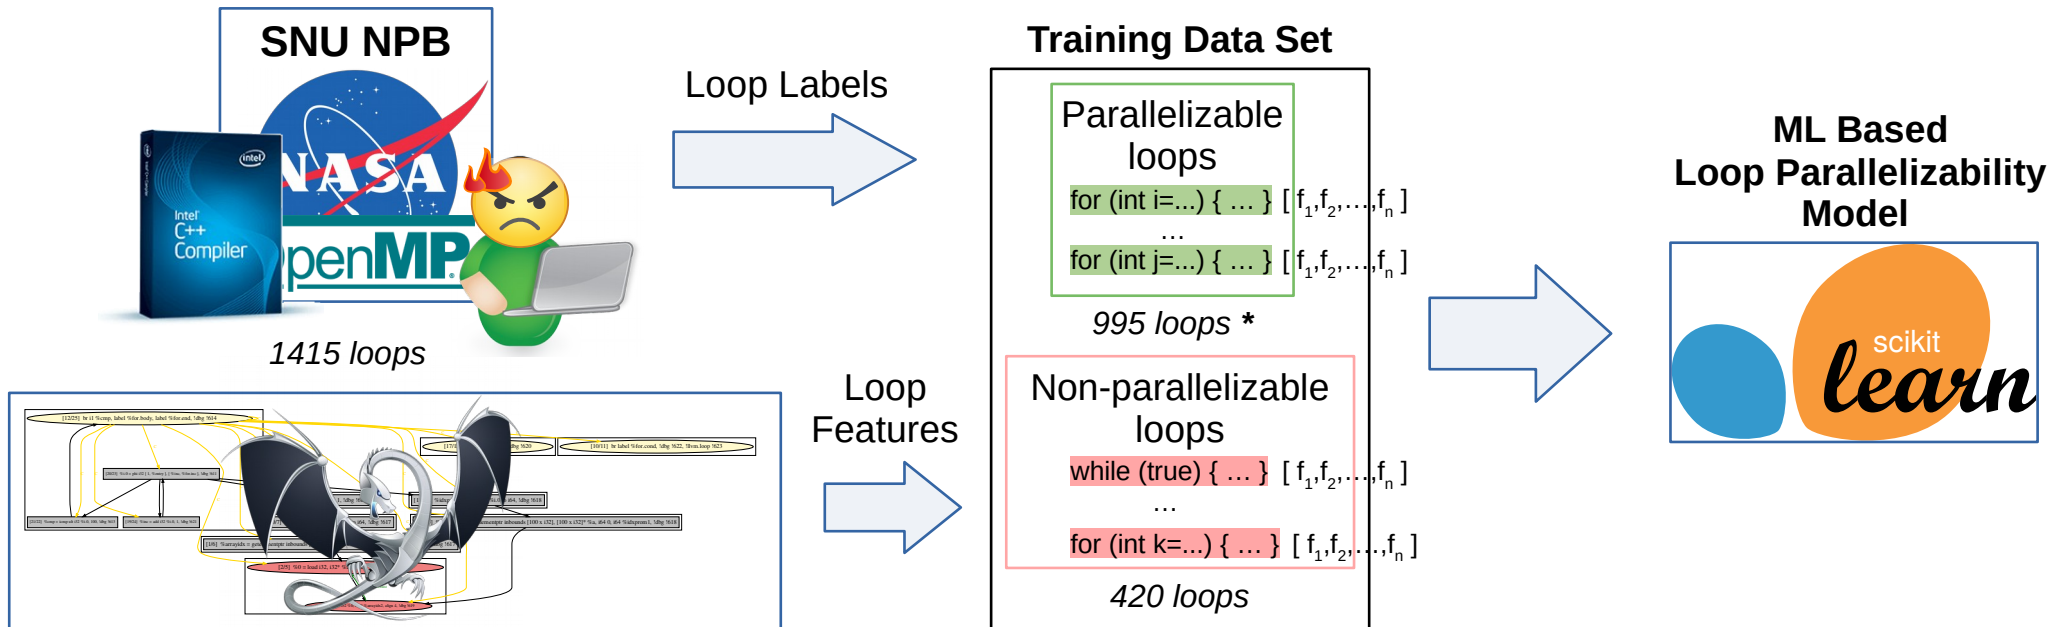


How to use the assistant

(from the programmers perspective)



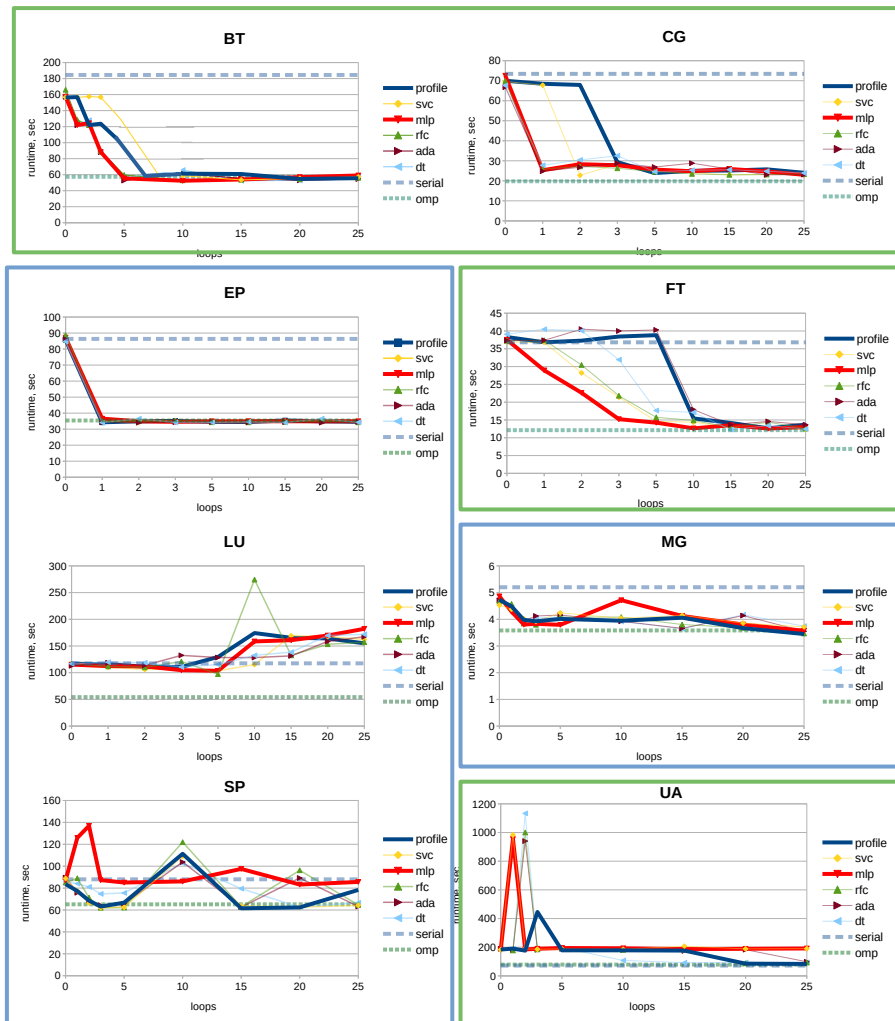
Solution Scheme [1/2]



* Intel Compiler succeeds in parallelizing 812 loops

How well do we do?


Improvement in
4 benchmarks



No change in
4 other
benchmarks

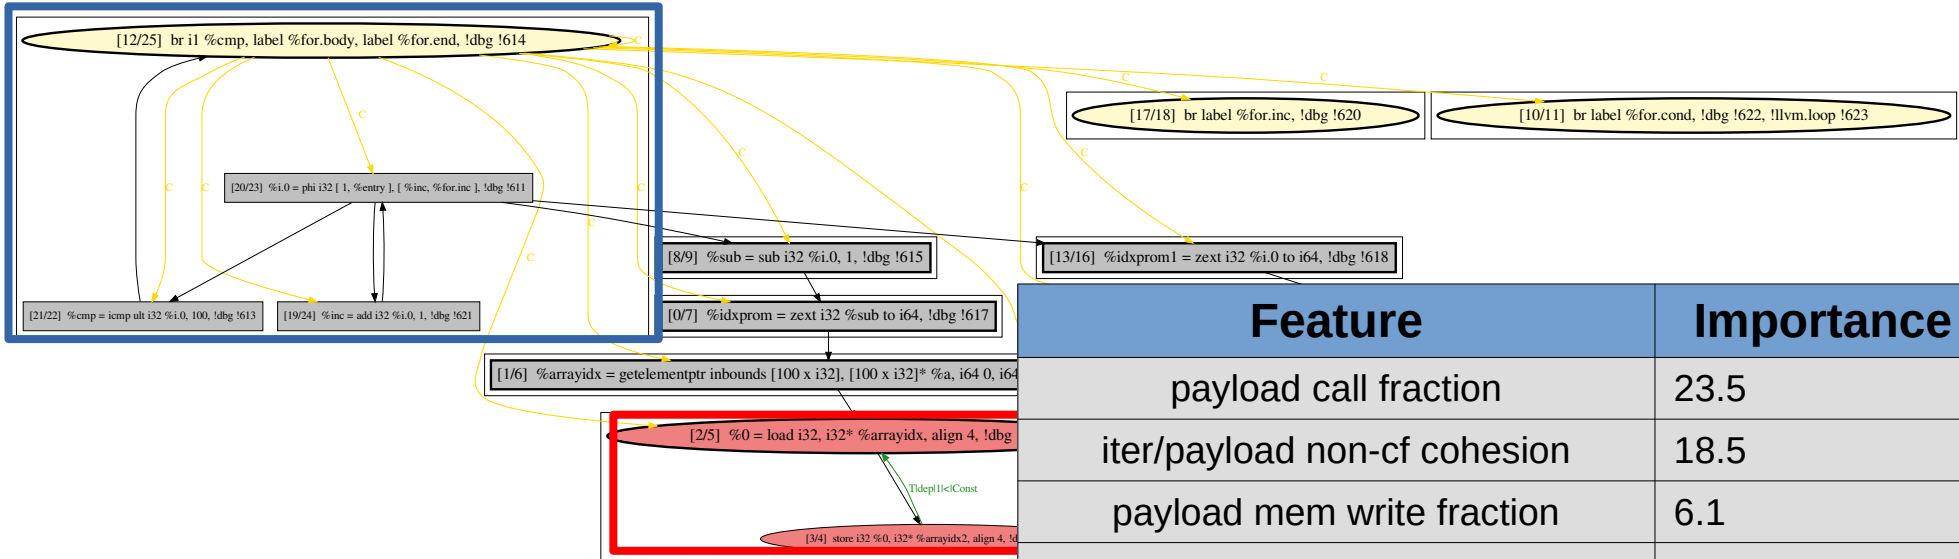
Problem Statement

NAS Parallel Benchmarks (NPB)



Loop Features [74]

Loop Iterator



Feature	Importance
payload call fraction	23.5
iter/payload non-cf cohesion	18.5
payload mem write fraction	6.1
loop absolute size	5.7
critical payload ptr access count	5.3
payload memory dependence count	4.0
critical payload non-cf cohesion	2.9
payload ptr access fraction	2.7



Loop features
of loop program

- Absolute
- Loop Iterator
- Number of dependence edges
- Instruction types (calls, loads/stores, etc.)
- etc.

Assistant Training

