

# A Machine Learning Based Parallelization Assistant

Aleksandr Maramzin, Christos Vasiladiotis, Roberto Castañeda Lozano,  
Björn Franke, Murray Cole

The University of Edinburgh  
United Kingdom



AI-SEPS 2019 workshop



**SPLASH**  
ATHENS 2019

# A Machine Learning Based Parallelization Assistant

Aleksandr Maramzin, Christos Vasiladiotis, Roberto Castañeda Lozano,  
Björn Franke, Murray Cole

“It Looks Like You’re Writing a Parallel Loop”



AI-SEPS 2019 workshop



**SPLASH**  
ATHENS 2019

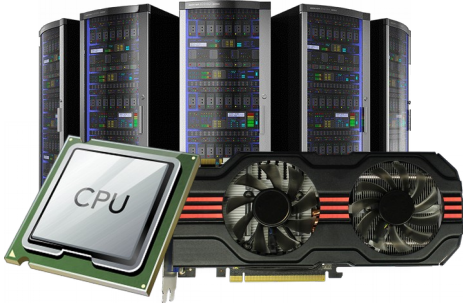
# Presentation Structure

- Problem statement
- How to use the assistant tool
- Assistant internals
- Parallelization of SNU NPB benchmarks

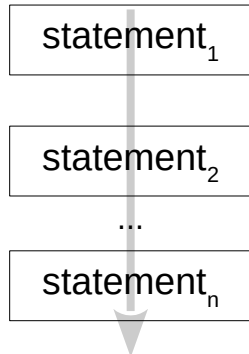


# Problem Statement

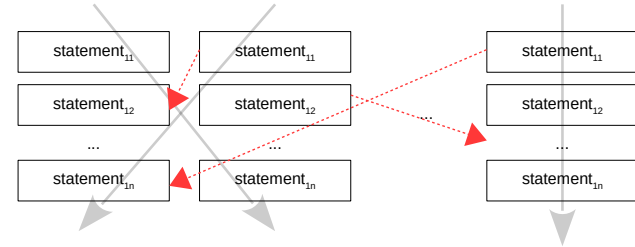
**Parallel Hardware  
is Ubiquitous**



**Software  
is Sequential**



**Automatic Parallelization  
is Limited**



**Manual Parallelization  
is Hard**



# Problem Statement

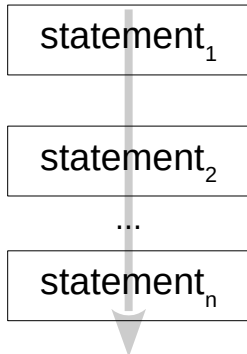
Parallel Hardware  
is Ubiquitous

Automatic Parallelization  
is Limited

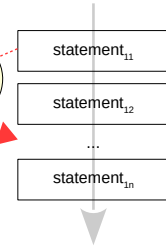
The **assistant solution** we propose alleviates  
the process of manual parallelization



Software  
is Sequential



Manual Parallelization  
is Hard



# Problem Statement

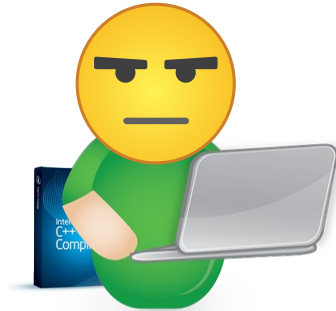
NAS Parallel  
Benchmarks  
(NPB)



**Parallelization Speedup  
(Gmean)**

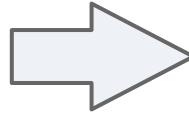
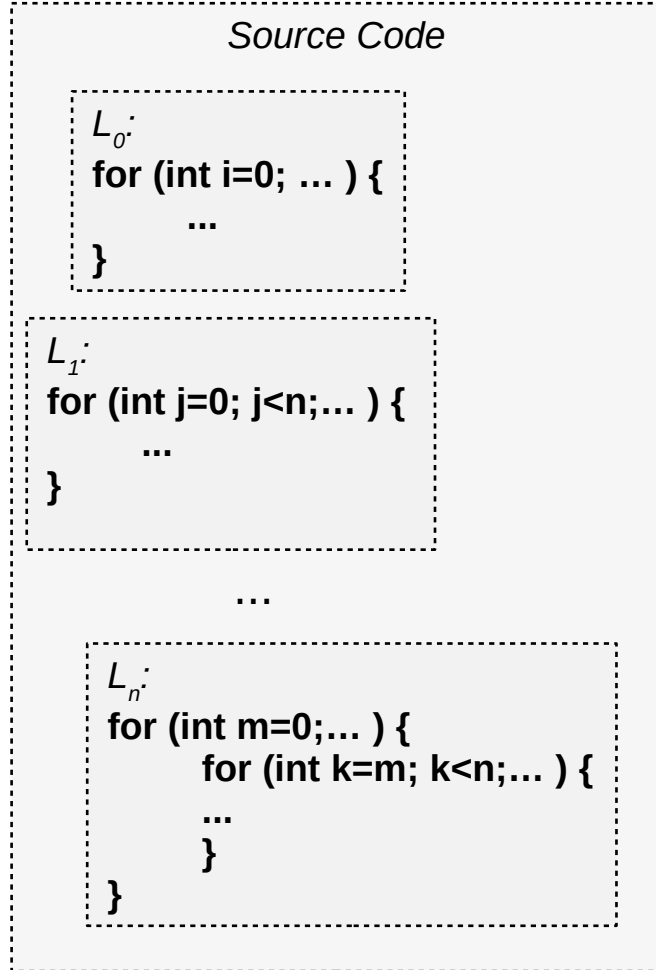
Manual: **1.73x**

Automatic: **0.79x - 1.10x**

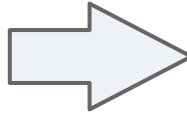
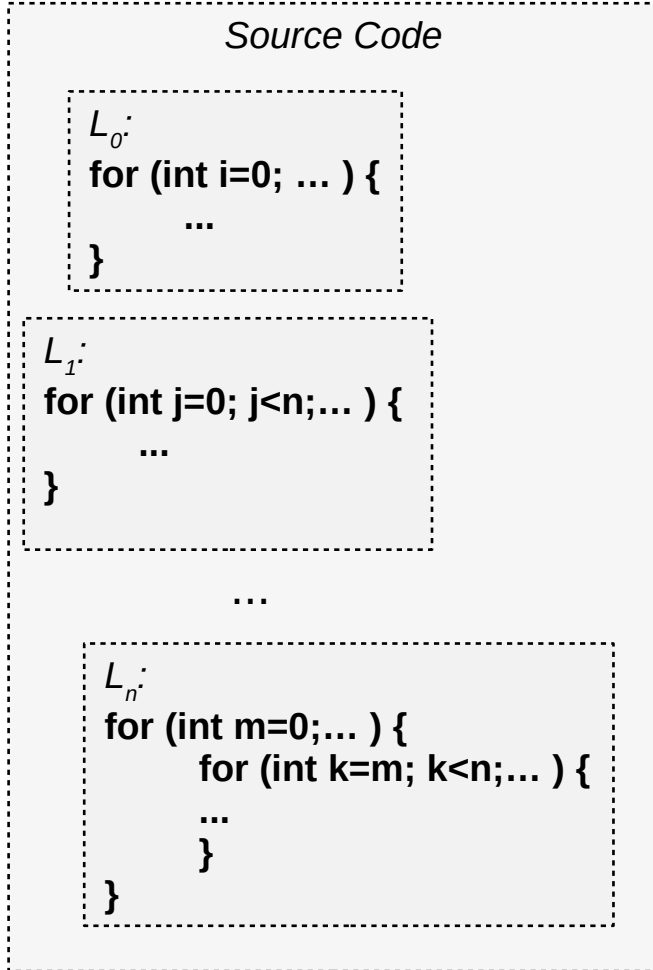


**Programmer still plays the major role  
in the hard task of software  
parallelization**

# Manual Parallelization: Where Should I Start?



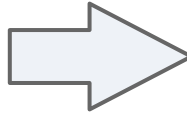
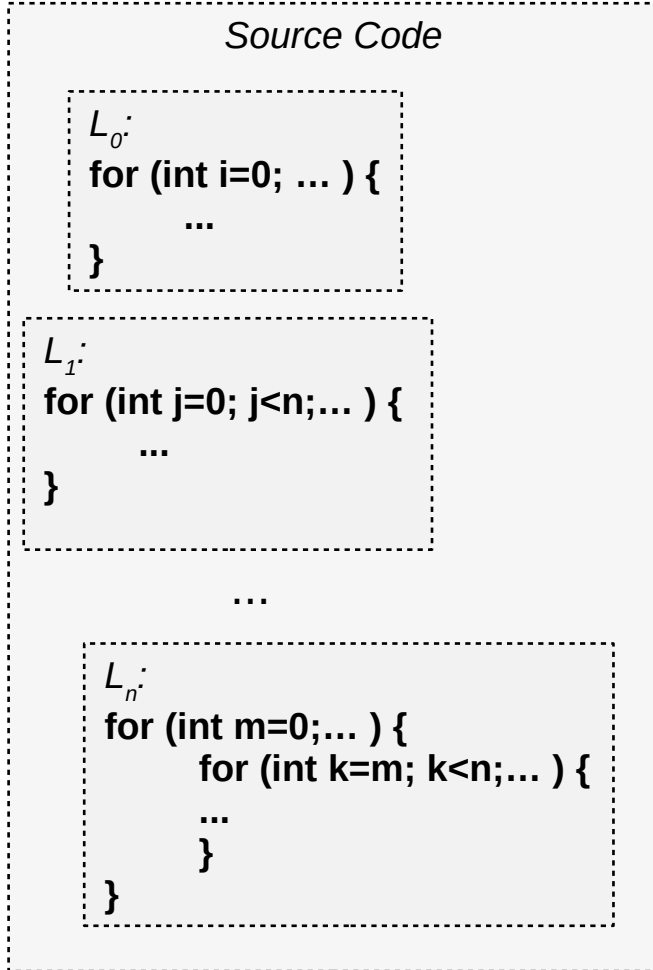
# Manual Parallelization: Where Should I Start?



Where Should I Start?



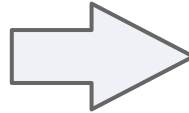
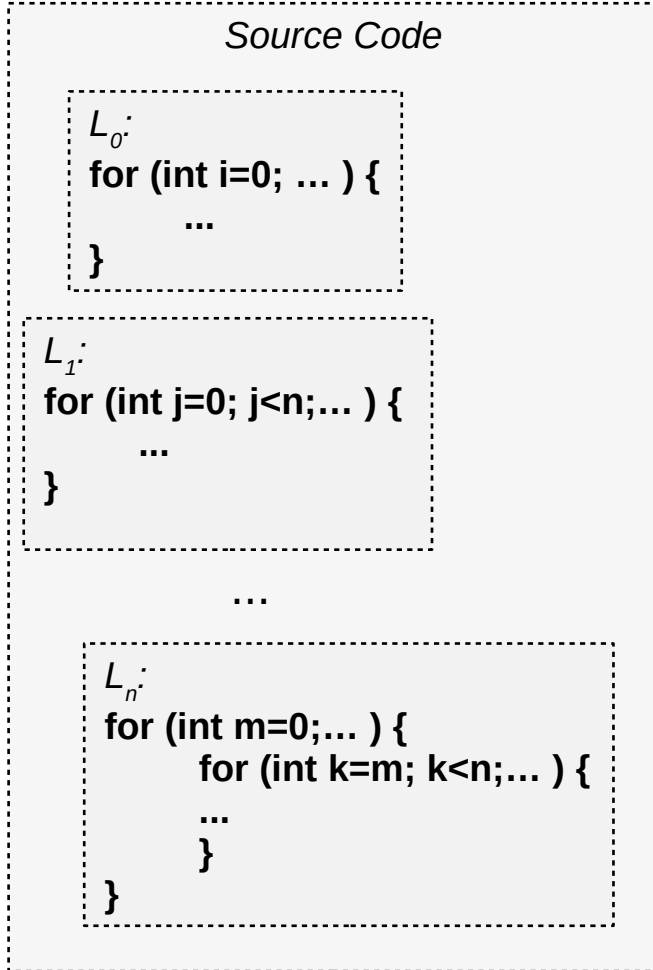
# Manual Parallelization: Where Should I Start?



## Program Loop Rankings

*Source Order*  
 $L_0, L_1, \dots, L_n$

# Manual Parallelization: Where Should I Start?

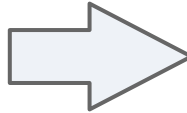
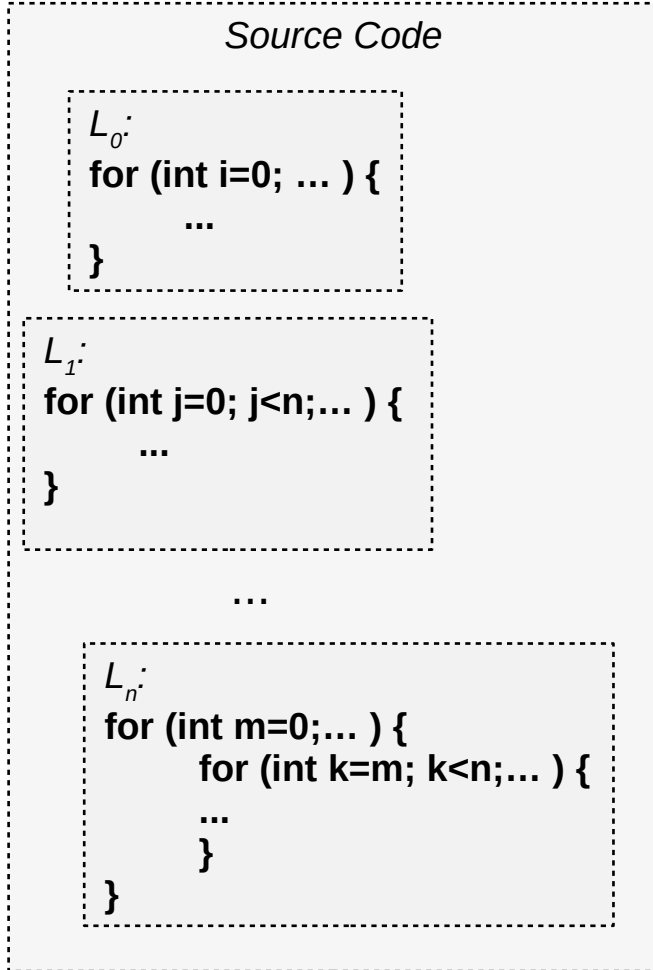


## Program Loop Rankings



Source Order  
 $L_0, L_1, \dots, L_n$

# Manual Parallelization: Where Should I Start?



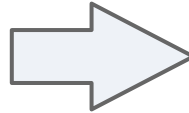
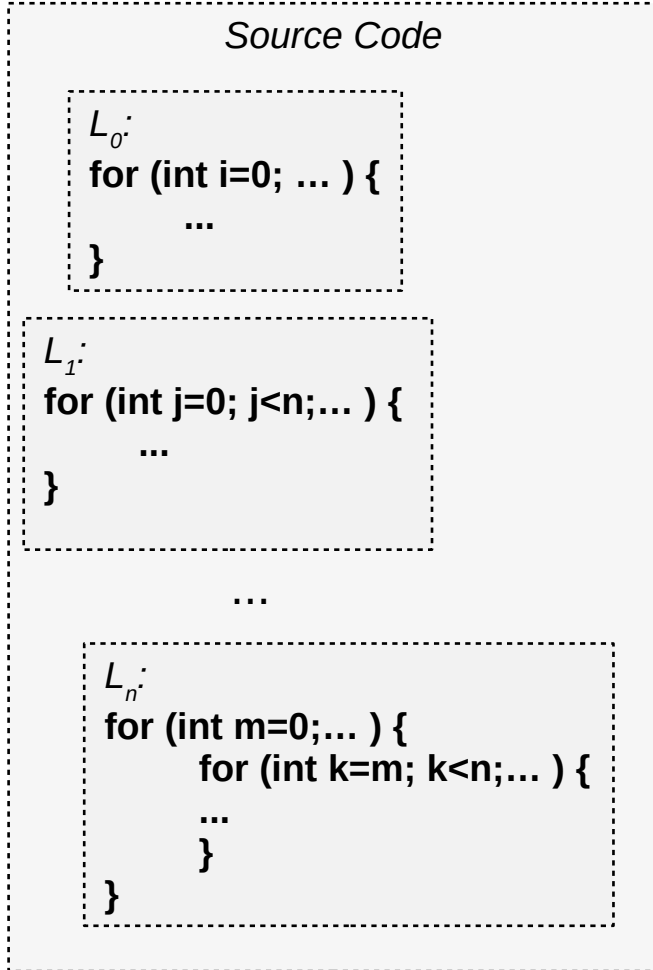
## Program Loop Rankings



Source Order  
 $L_0, L_1, \dots, L_n$



# Manual Parallelization: Where Should I Start?



## Program Loop Rankings



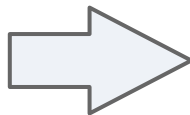
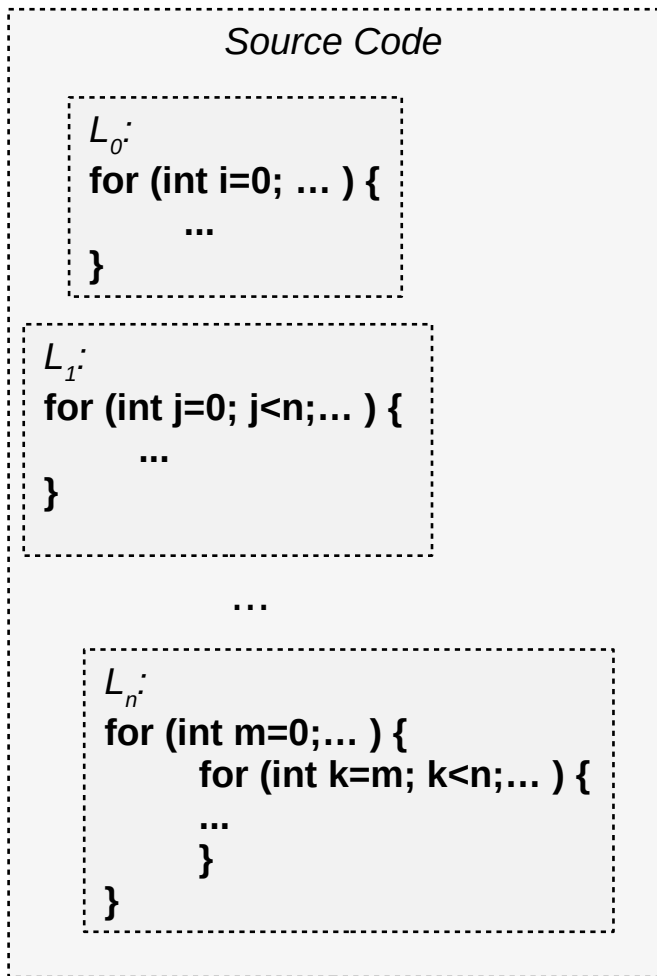
Source Order  
 $L_0, L_1, \dots, L_n$



Traditional Empirical and  
Experimental Method

Profile Order  
 $L_{17}, L_{12}, \dots, L_1$

# Manual Parallelization: Where Should I Start?



Traditional Empirical and  
Experimental Method

## Program Loop Rankings



Source Order

$L_0, L_1, \dots, L_n$

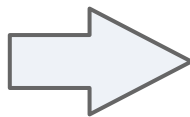
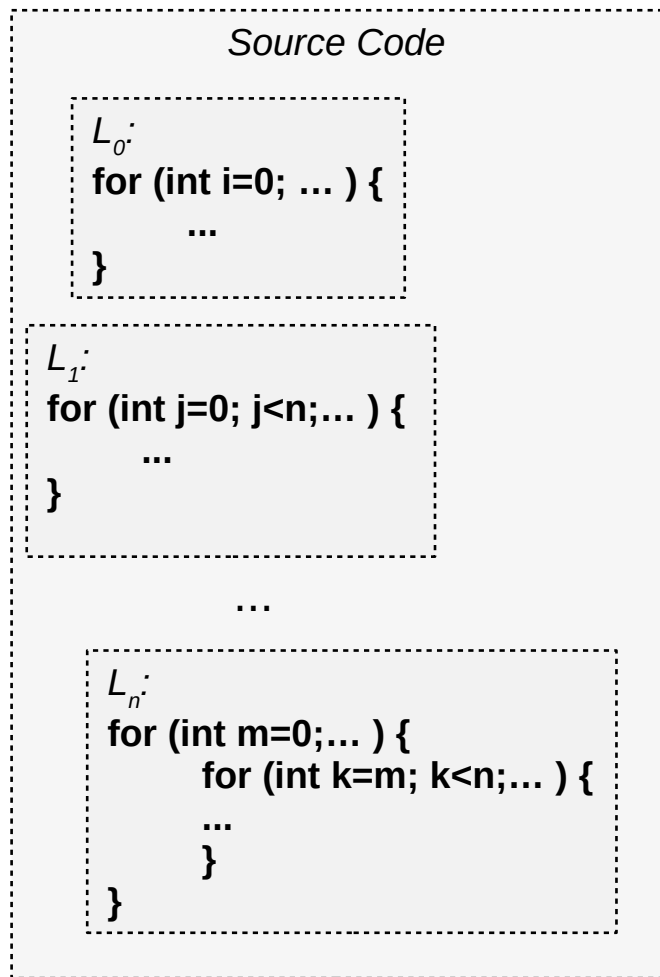


Profile Order

$L_{17}, L_{12}, \dots, L_1$



# Manual Parallelization: Where Should I Start?



Traditional Empirical and Experimental Method

## Program Loop Rankings



Source Order  
 $L_0, L_1, \dots, L_n$



Profile Order  
 $L_{17}, L_{12}, \dots, L_1$

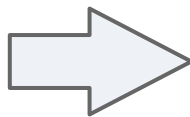
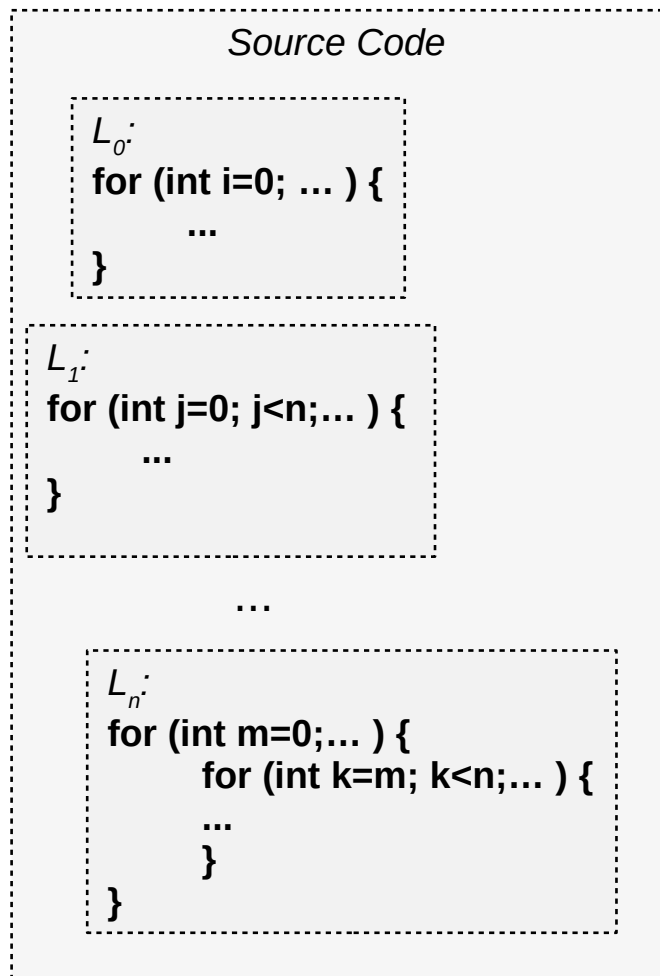


AI-inspired Method

Profitability\* Order  
 $L_{12}, L_7, \dots, L_3$

$$* \text{Profitability} = F\left(\frac{\text{Performance Gain}}{\text{Programmer Effort}}\right)$$

# Manual Parallelization: Where Should I Start?



Traditional Empirical and  
Experimental Method

## Program Loop Rankings



Source Order

$L_0, L_1, \dots, L_n$



Profile Order

$L_{17}, L_{12}, \dots, L_1$



Profitability\* Order

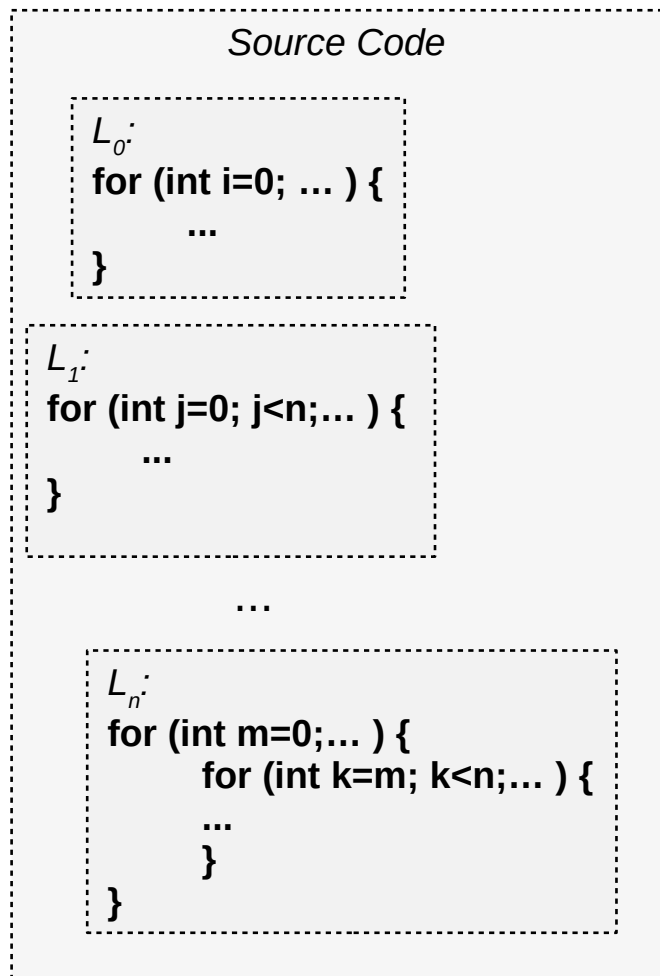
$L_{12}, L_7, \dots, L_3$



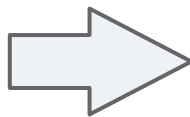
AI-inspired Method

$$* \text{Profitability} = F\left(\frac{\text{Performance Gain}}{\text{Programmer Effort}}\right)$$

# Manual Parallelization: Where Should I Start?



**Final Programmer  
Decision**



**Traditional Empirical and  
Experimental Method**

## Program Loop Rankings



Source Order  
 $L_0, L_1, \dots, L_n$



Profile Order  
 $L_{17}, L_{12}, \dots, L_1$



Profitability\* Order  
 $L_{12}, L_7, \dots, L_3$



**AI-inspired Method**

$$* \text{Profitability} = F\left(\frac{\text{Performance Gain}}{\text{Programmer Effort}}\right)$$



# Proposed Solution

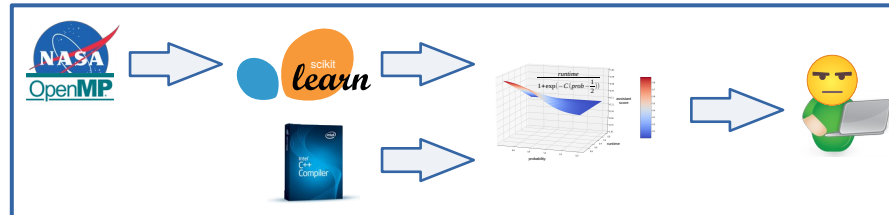
*(ML based software parallelization assistant)*

Tool



&

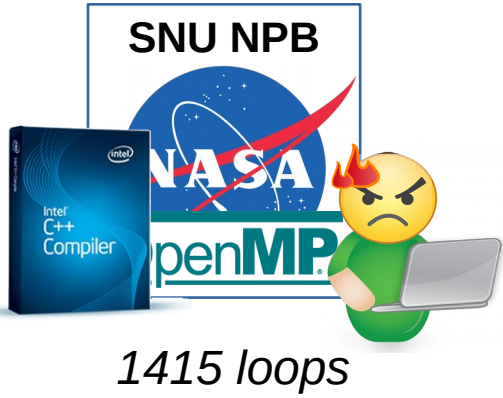
Methodology



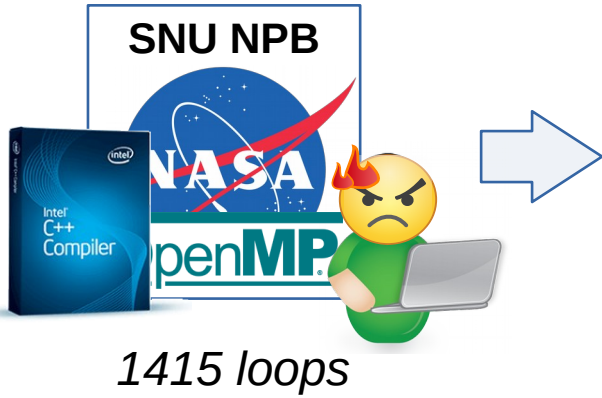
AI-SEPS 2019 workshop

# Assistant Training

# Assistant Training



# Assistant Training



## Training Data Set

### Parallelizable loops

```
for (int i=...) { ... } [ f1, f2, ..., fn ]  
...  
for (int j=...) { ... } [ f1, f2, ..., fn ]
```

995 loops \*

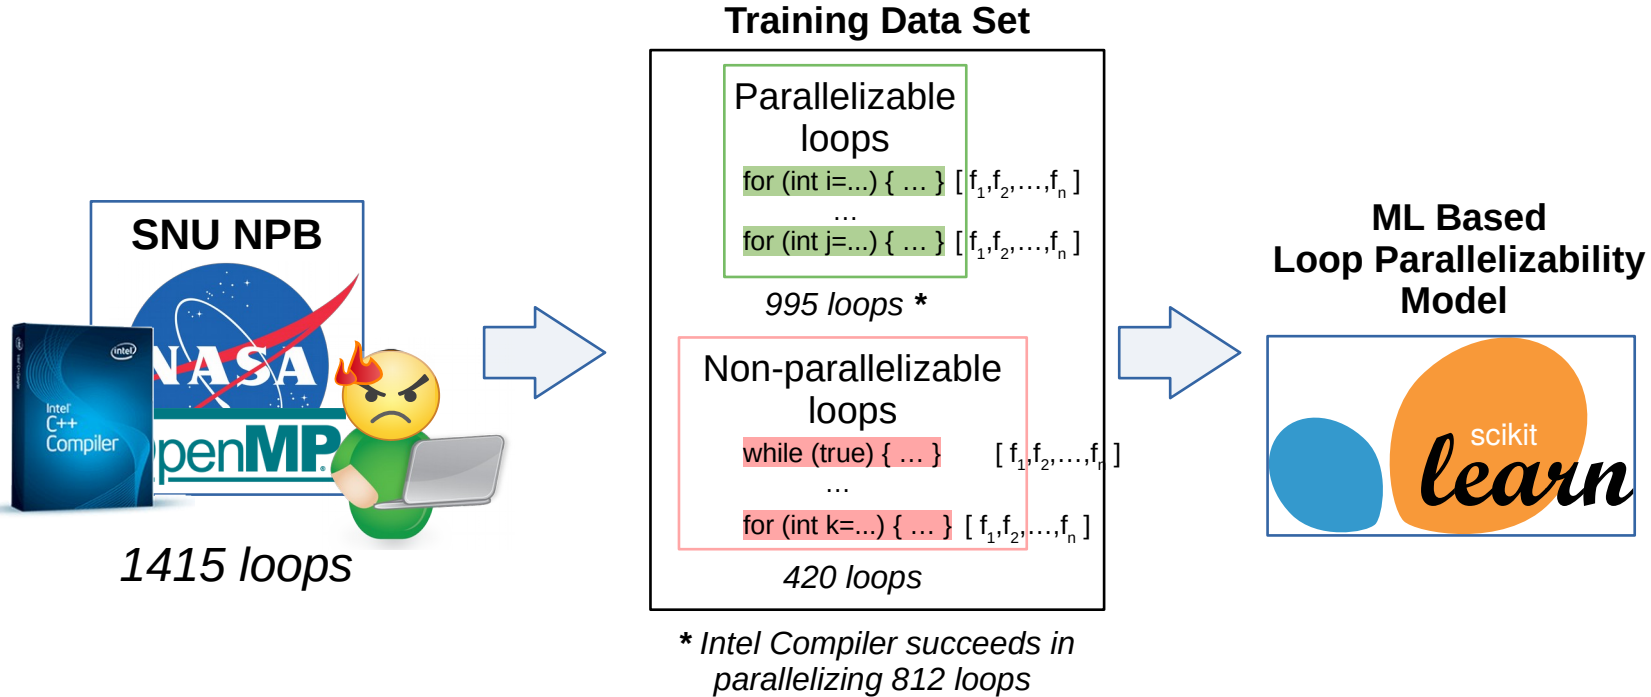
### Non-parallelizable loops

```
while (true) { ... } [ f1, f2, ..., fn ]  
...  
for (int k=...) { ... } [ f1, f2, ..., fn ]
```

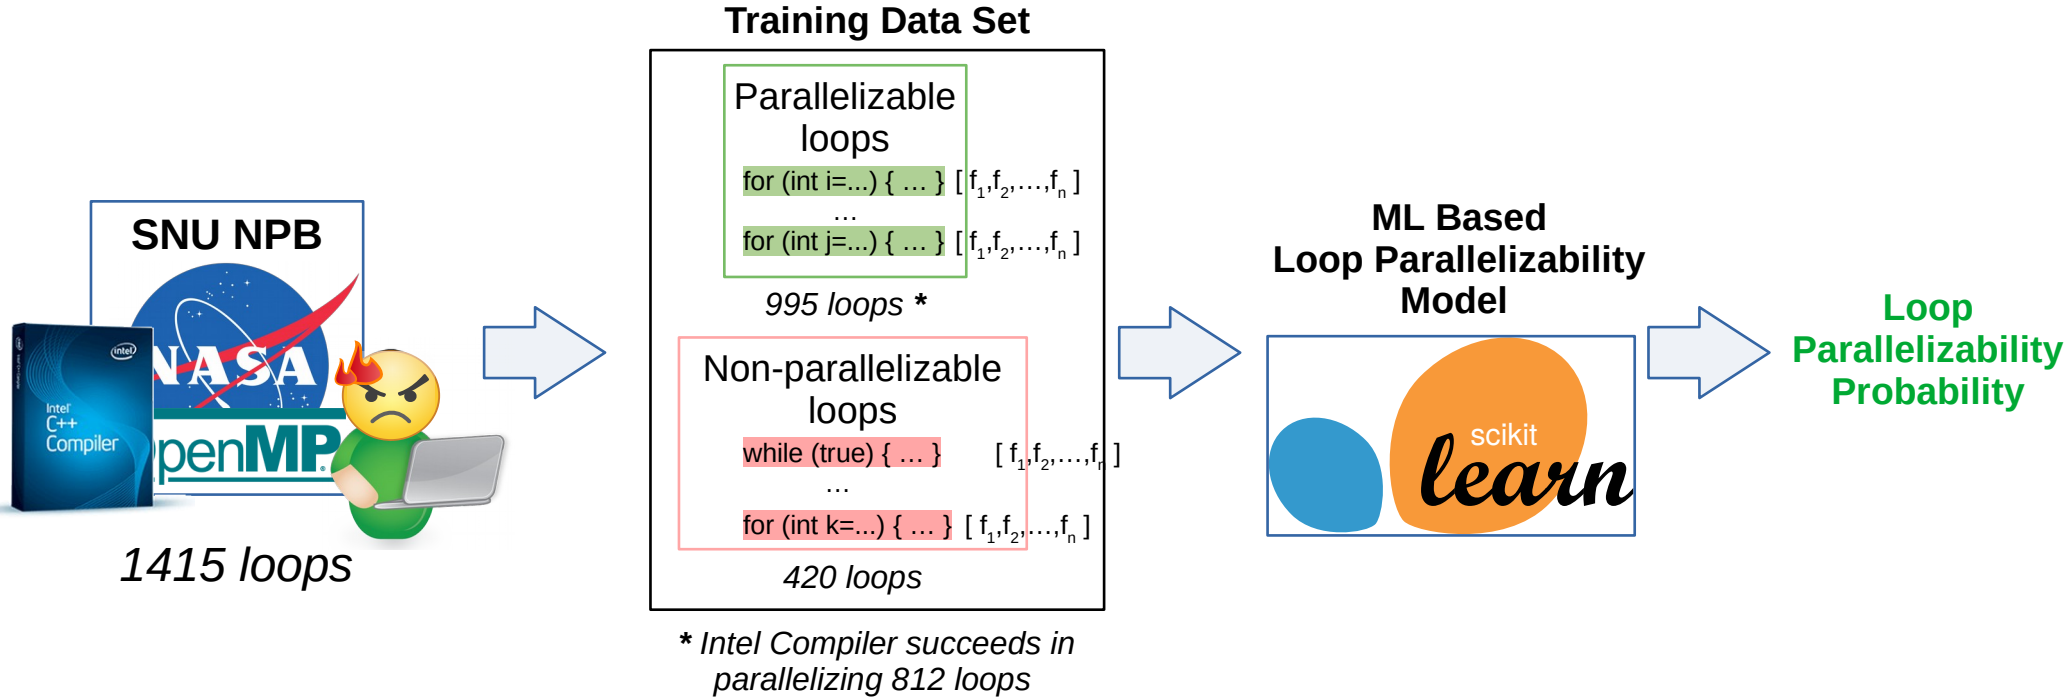
420 loops

\* Intel Compiler succeeds in parallelizing 812 loops

# Assistant Training

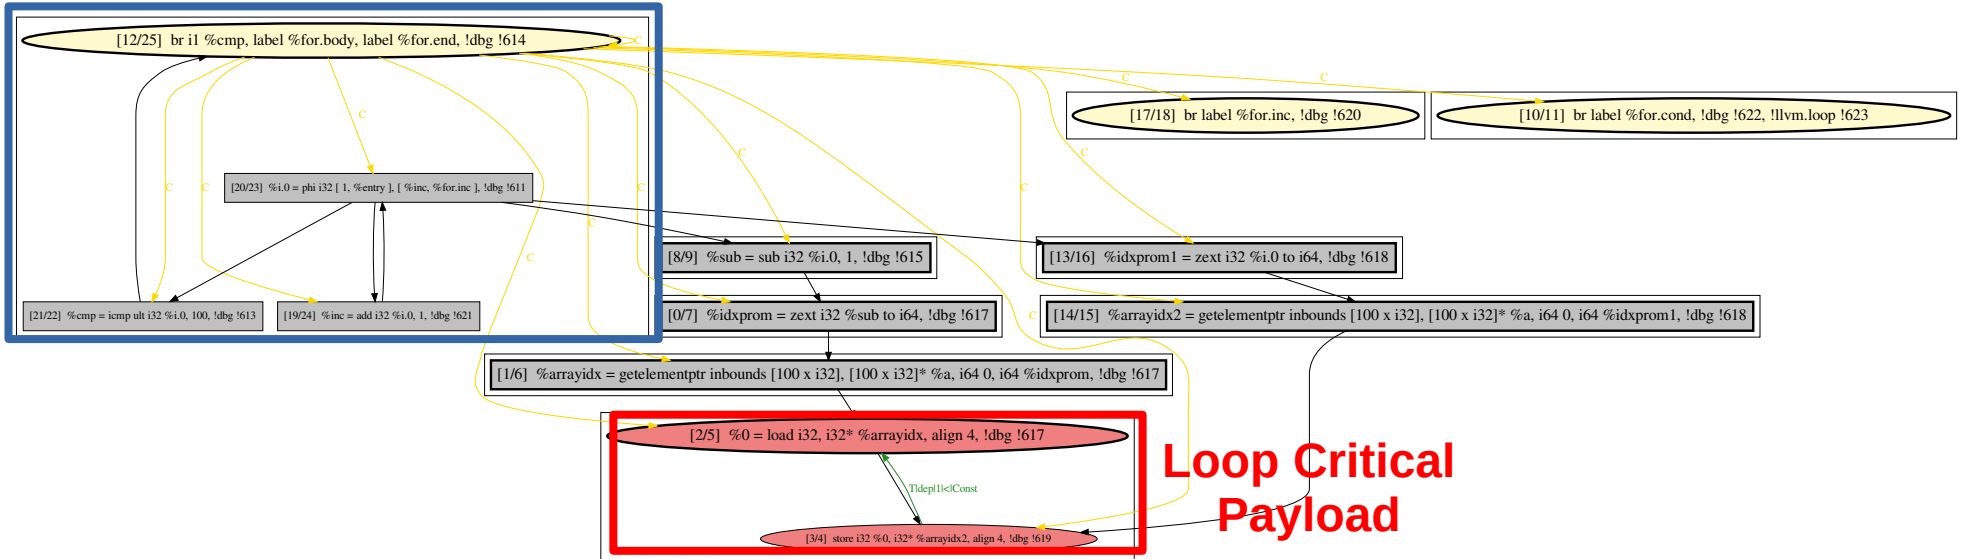


# Assistant Training



# Loop Features [74]

## Loop Iterator

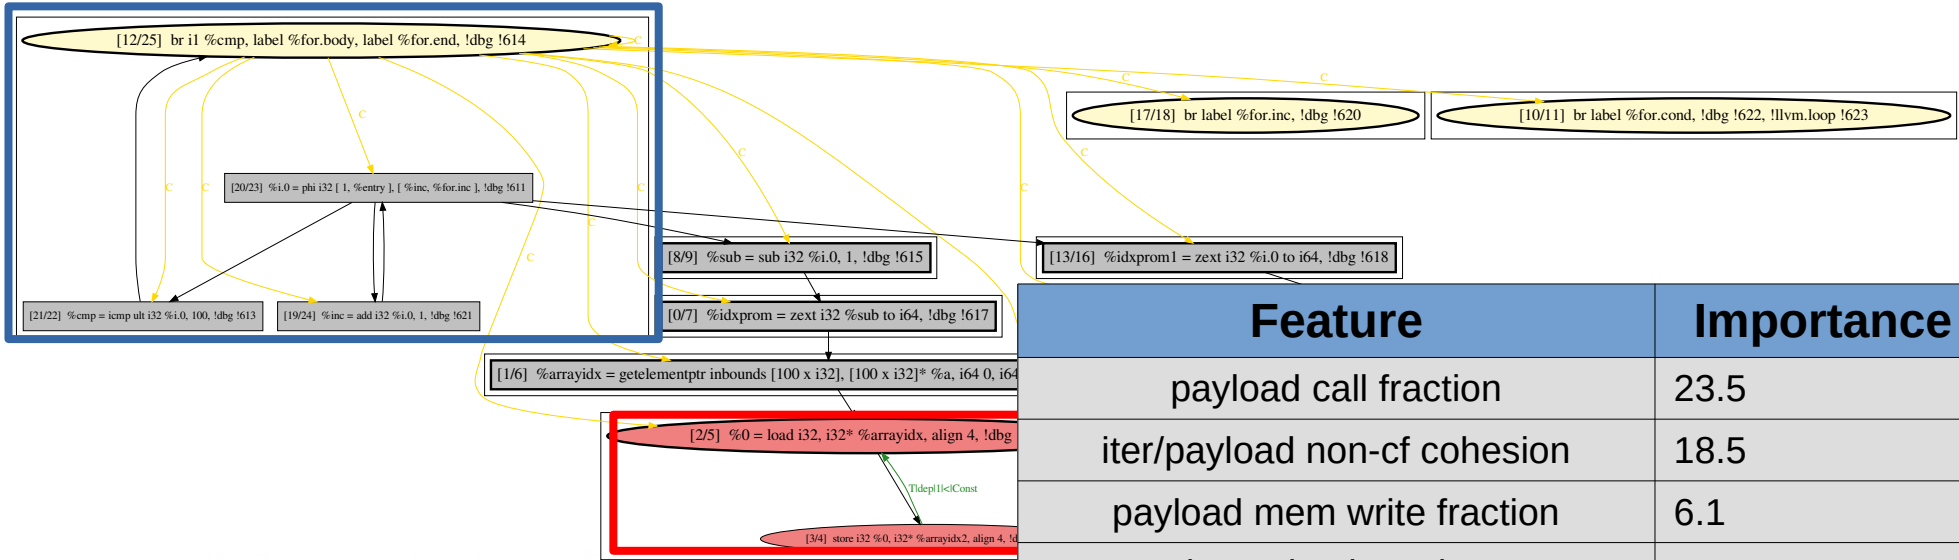


**Loop features** are based on static structural properties of loop program dependence graphs (PDGs):

- Absolute size
- Loop Iterator/Payload cohesion
- Number of dependence edges
- Instruction types (calls, loads/stores, etc.)
- etc.

# Loop Features [74]

## Loop Iterator



Feature	Importance
payload call fraction	23.5
iter/payload non-cf cohesion	18.5
payload mem write fraction	6.1
loop absolute size	5.7
critical payload ptr access count	5.3
payload memory dependence count	4.0
critical payload non-cf cohesion	2.9
payload ptr access fraction	2.7



Loop features  
of loop program

- Absolute
- Loop Iterator
- Number of dependence edges
- Instruction types (calls, loads/stores, etc.)
- etc.



# Loop Ranking Score

ML Based  
Loop Parallelizability  
Model



# Loop Ranking Score

ML Based  
Loop Parallelizability  
Model



Not Aware of Loop  
Running Time

# Loop Ranking Score

**Profiler**



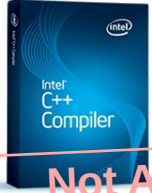
**ML Based  
Loop Parallelizability  
Model**



**Not Aware of Loop  
Running Time**

# Loop Ranking Score

## Profiler



Not Aware of Loop  
Parallelizability

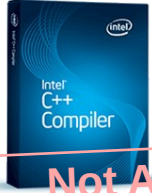
## ML Based Loop Parallelizability Model



Not Aware of Loop  
Running Time

# Loop Ranking Score

**Profiler**



Not Aware of Loop  
Parallelizability

**ML Based  
Loop Parallelizability  
Model**



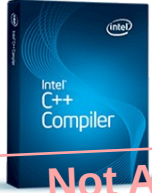
Not Aware of Loop  
Running Time

**Programmer Effort**



# Loop Ranking Score

Profiler



Not Aware of Loop  
Parallelizability

Performance Gain



ML Based  
Loop Parallelizability  
Model



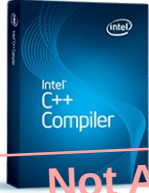
Not Aware of Loop  
Running Time

Programmer Effort



# Loop Ranking Score

Profiler

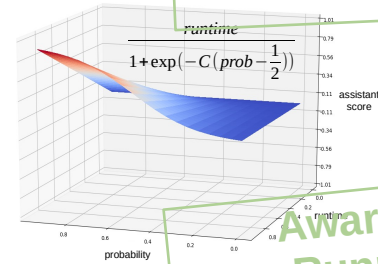


Not Aware of Loop Parallelizability

Performance Gain



Aware of Loop Parallelizability \*



Aware of Loop Running Time

ML Based  
Loop Parallelizability  
Model



Not Aware of Loop Running Time

Programmer Effort



$$* \text{Profitability} = F\left(\frac{\text{Performance Gain}}{\text{Programmer Effort}}\right)$$

# Loop Ranking Score

Profiler



Not Aware of Loop Parallelizability

Performance Gain



ML Based  
Loop Parallelizability  
Model



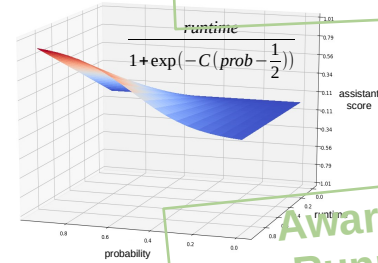
Not Aware of Loop Running Time

Programmer Effort



Aware of Loop Parallelizability

\*



Aware of Loop Running Time

Improved Loop Ranking

$$* \text{Profitability} = F\left(\frac{\text{Performance Gain}}{\text{Programmer Effort}}\right)$$



## Results



Deployment of our assistant on  
SNU NPB benchmarks



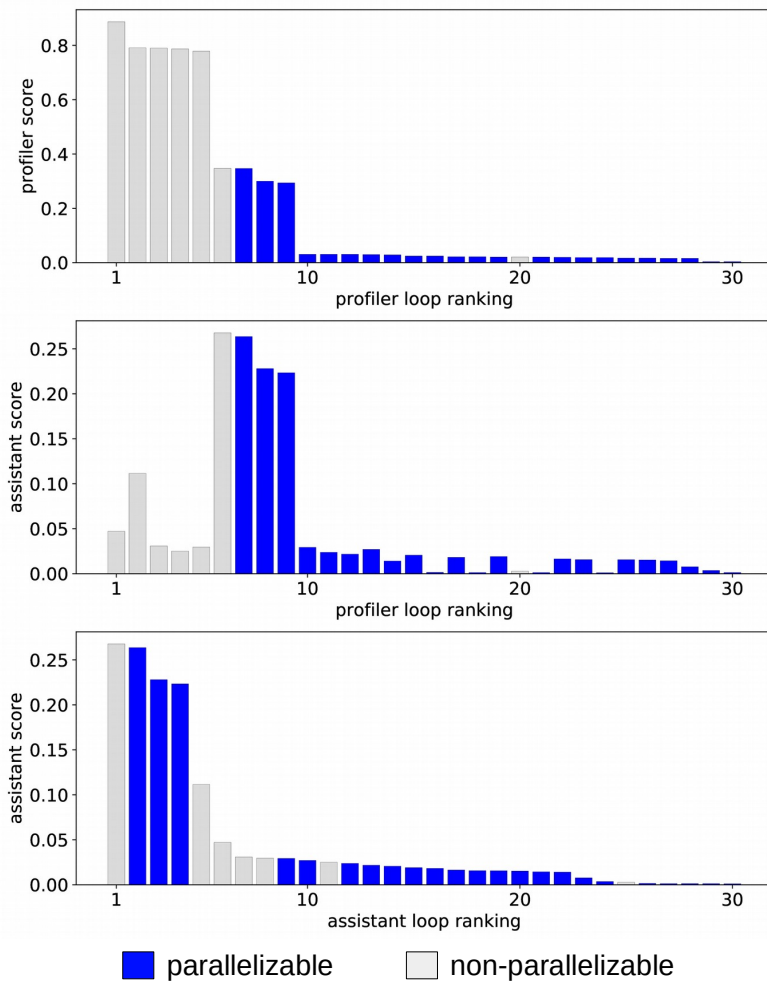
Predictive performance of our ML  
based loop parallelizability model



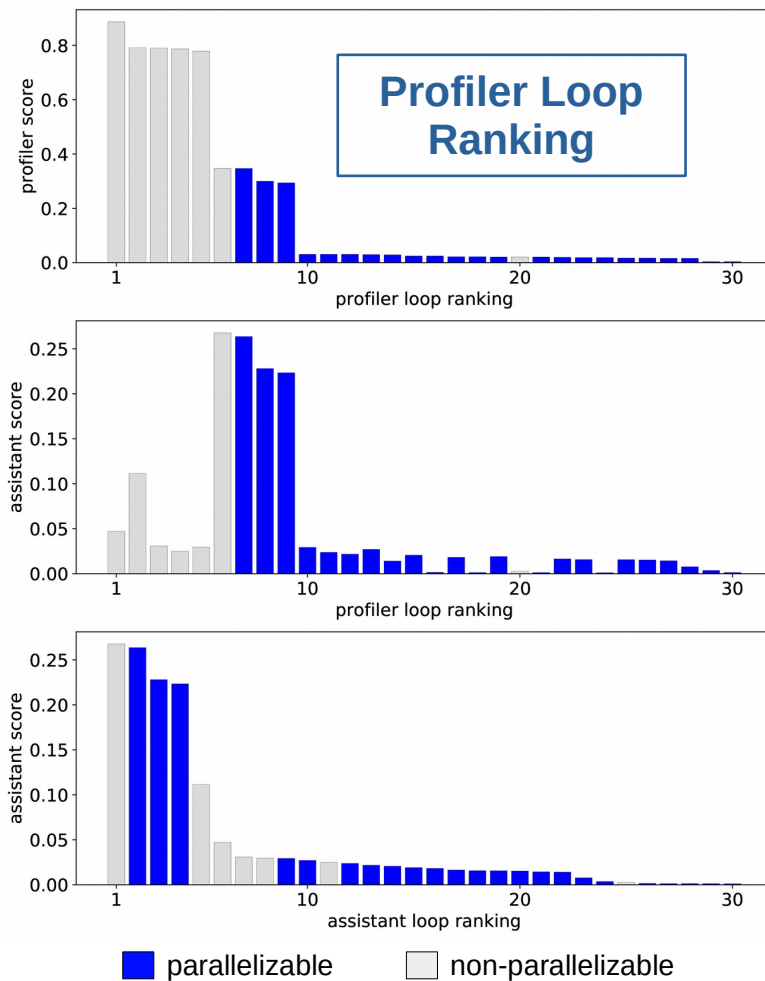
AI-SEPS 2019 workshop



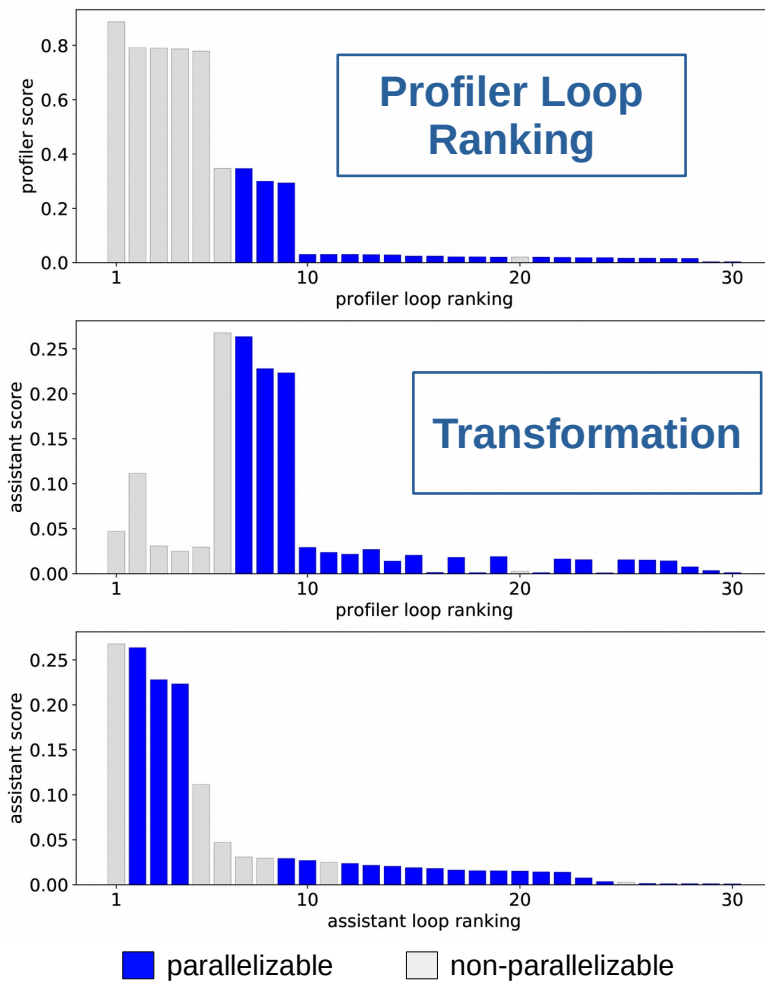
# Loop Rankings: Profiler vs. Assistant



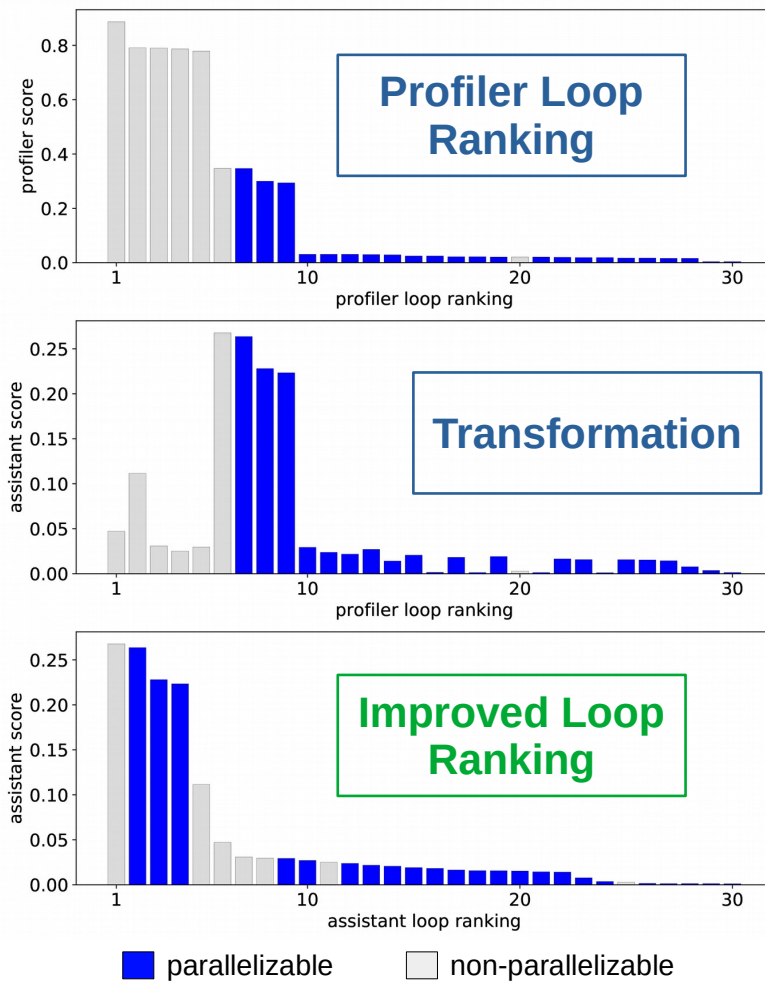
# Loop Rankings: Profiler vs. Assistant



# Loop Rankings: Profiler vs. Assistant

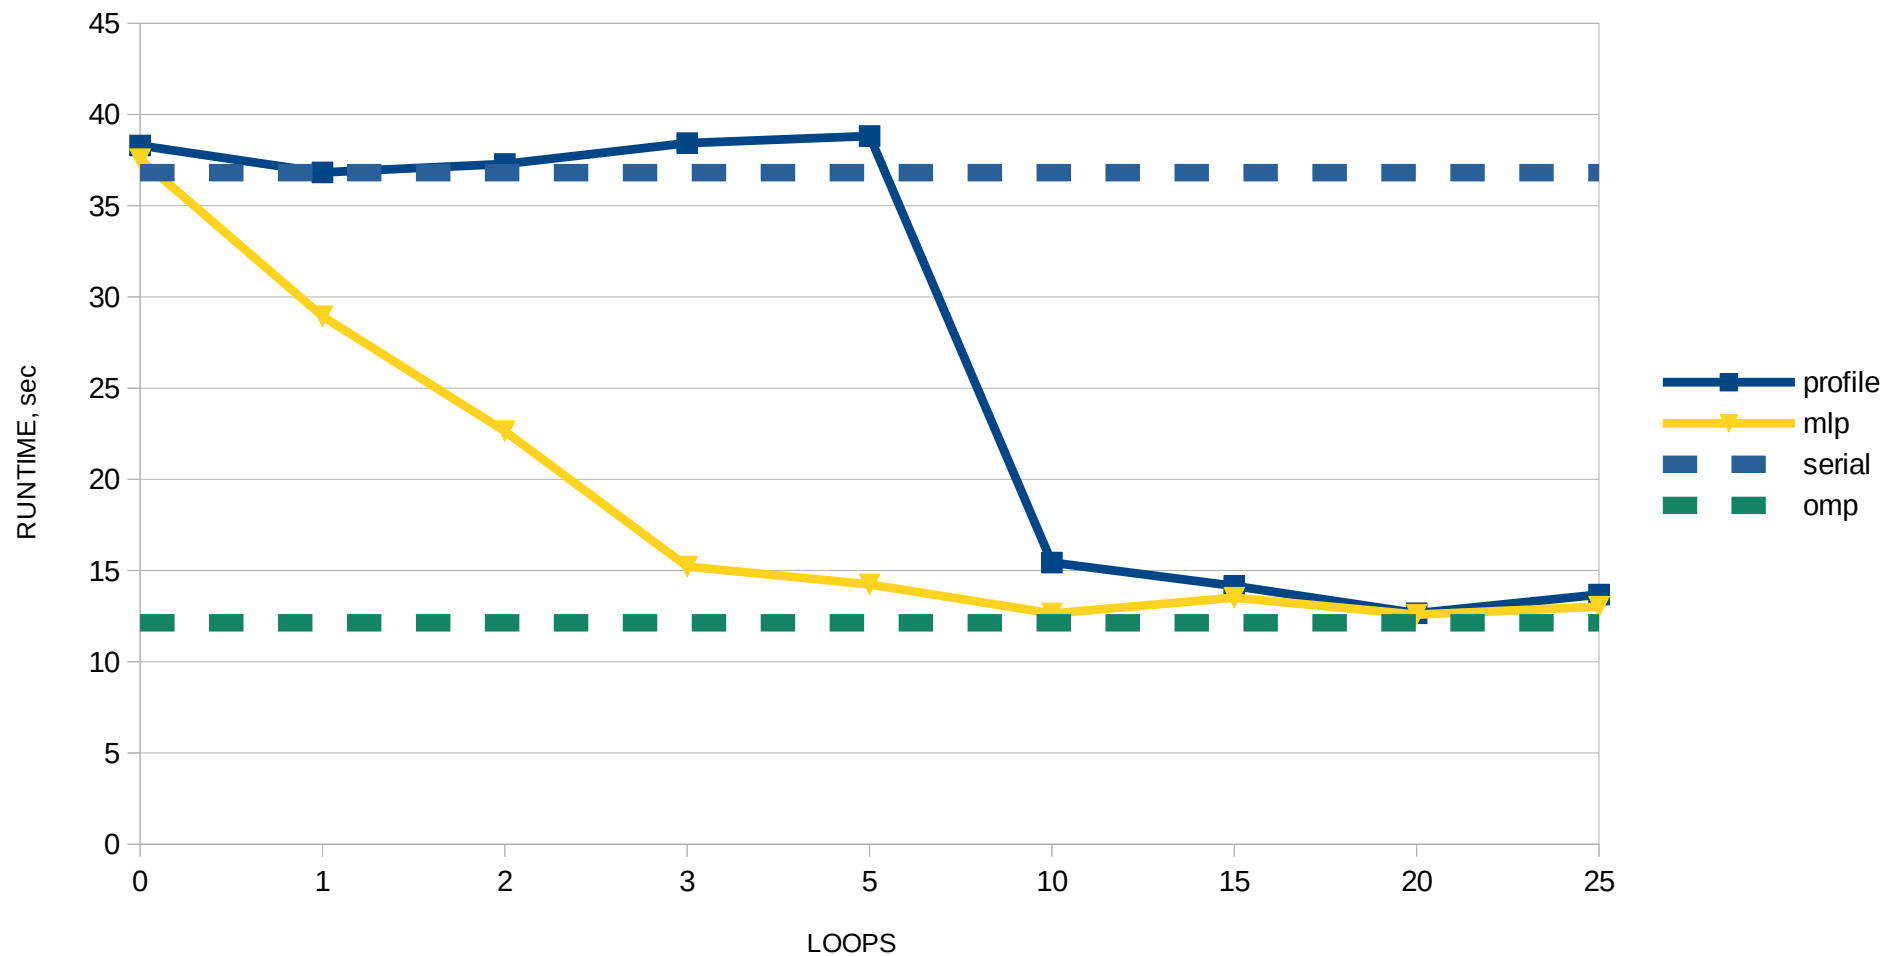


# Loop Rankings: Profiler vs. Assistant

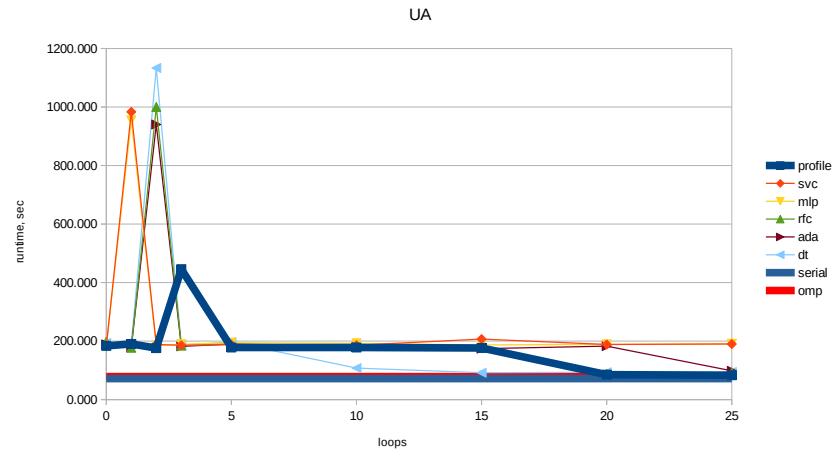
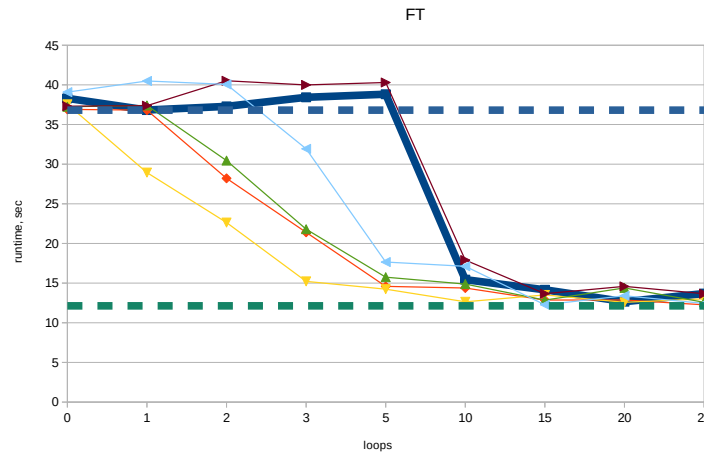
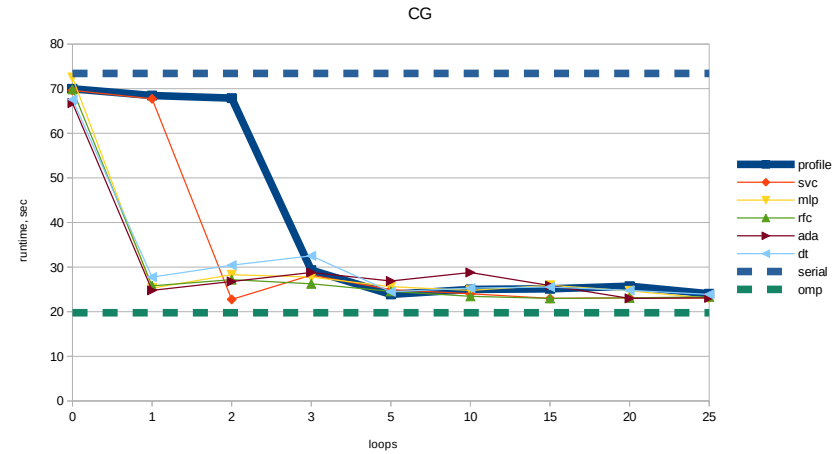
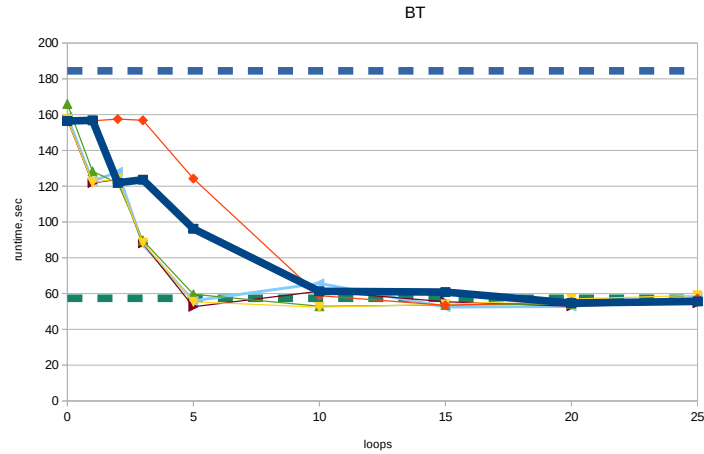


# How well do we do?

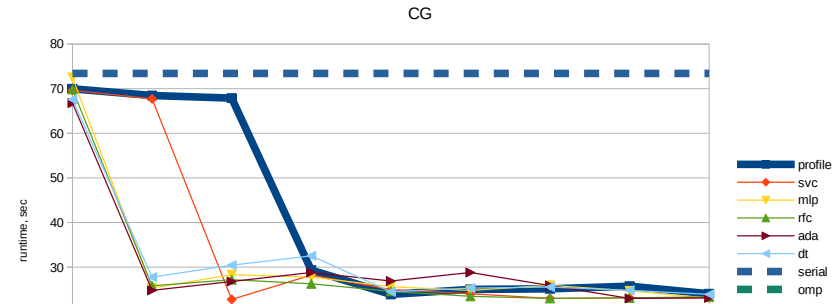
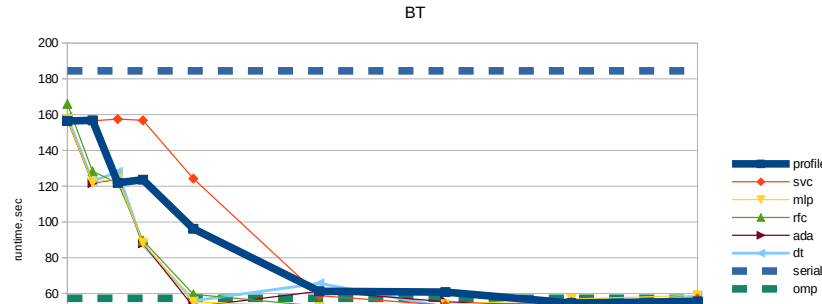
FT



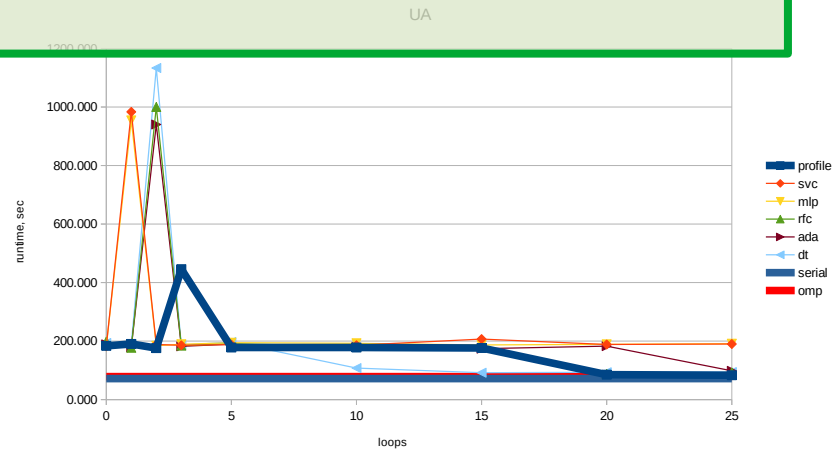
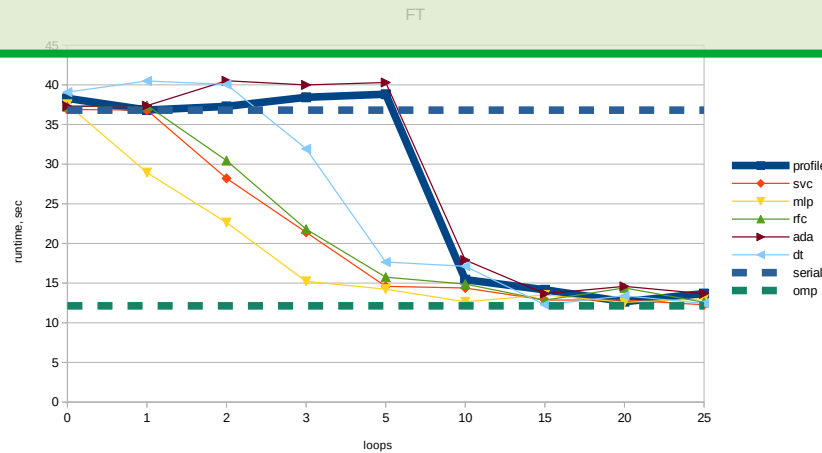
# How well do we do?



# How well do we do?

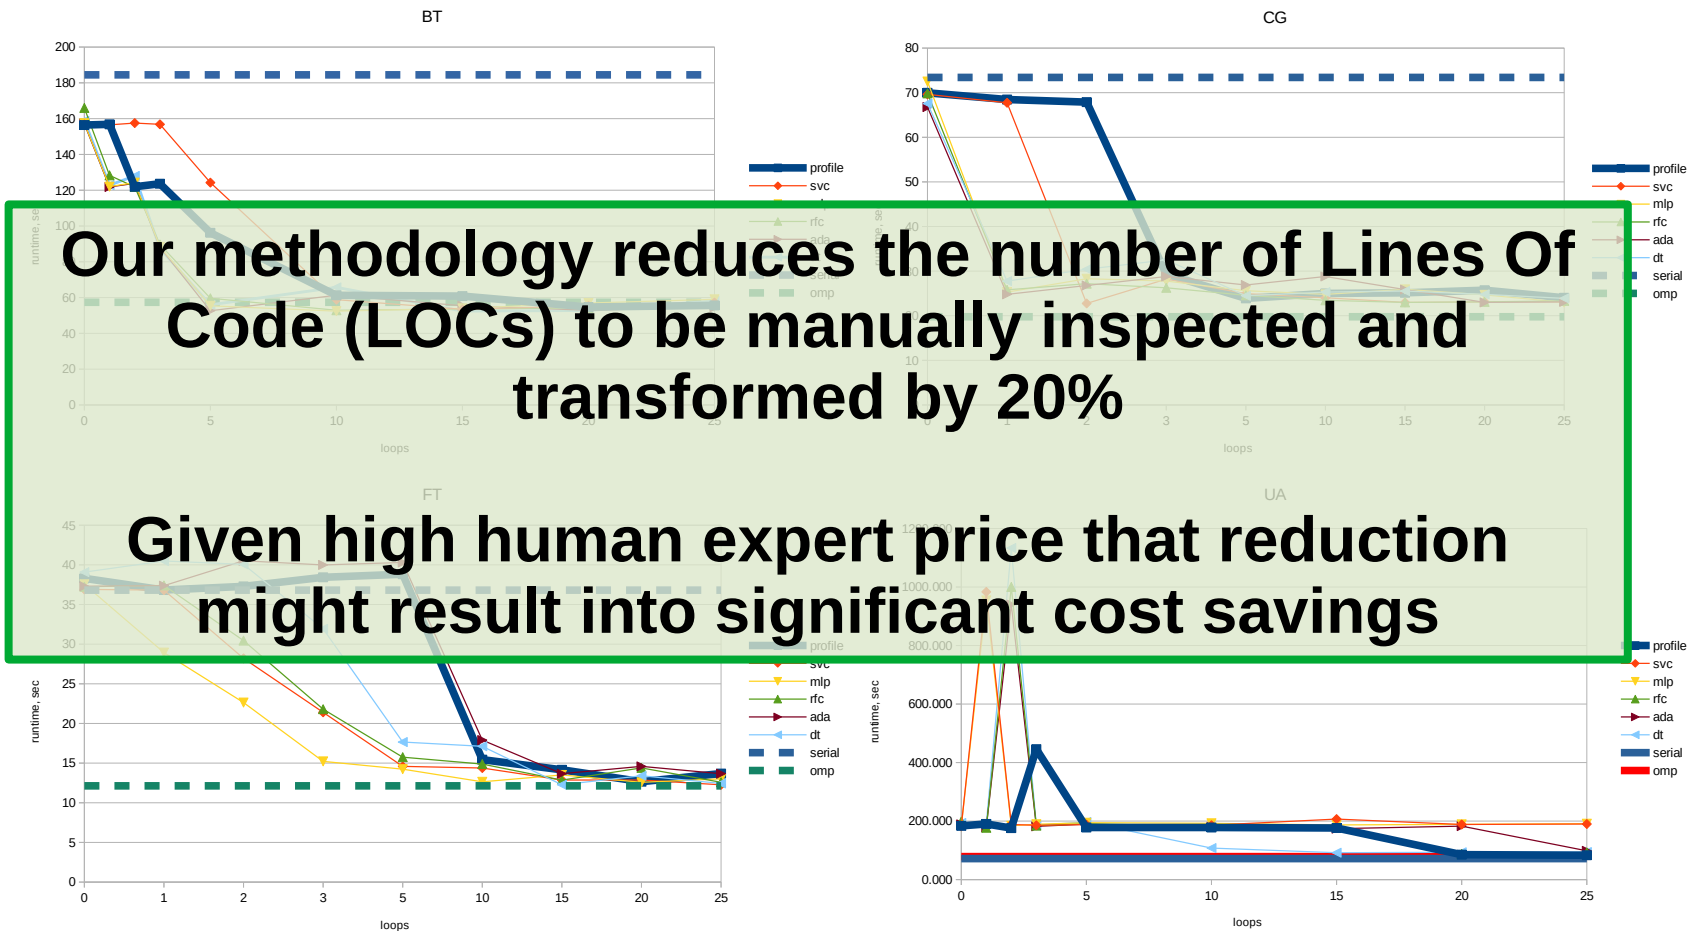


8 benchmarks in total: 4 wins / 4 draws





# How well do we do?



# ML Model Predictive Performance

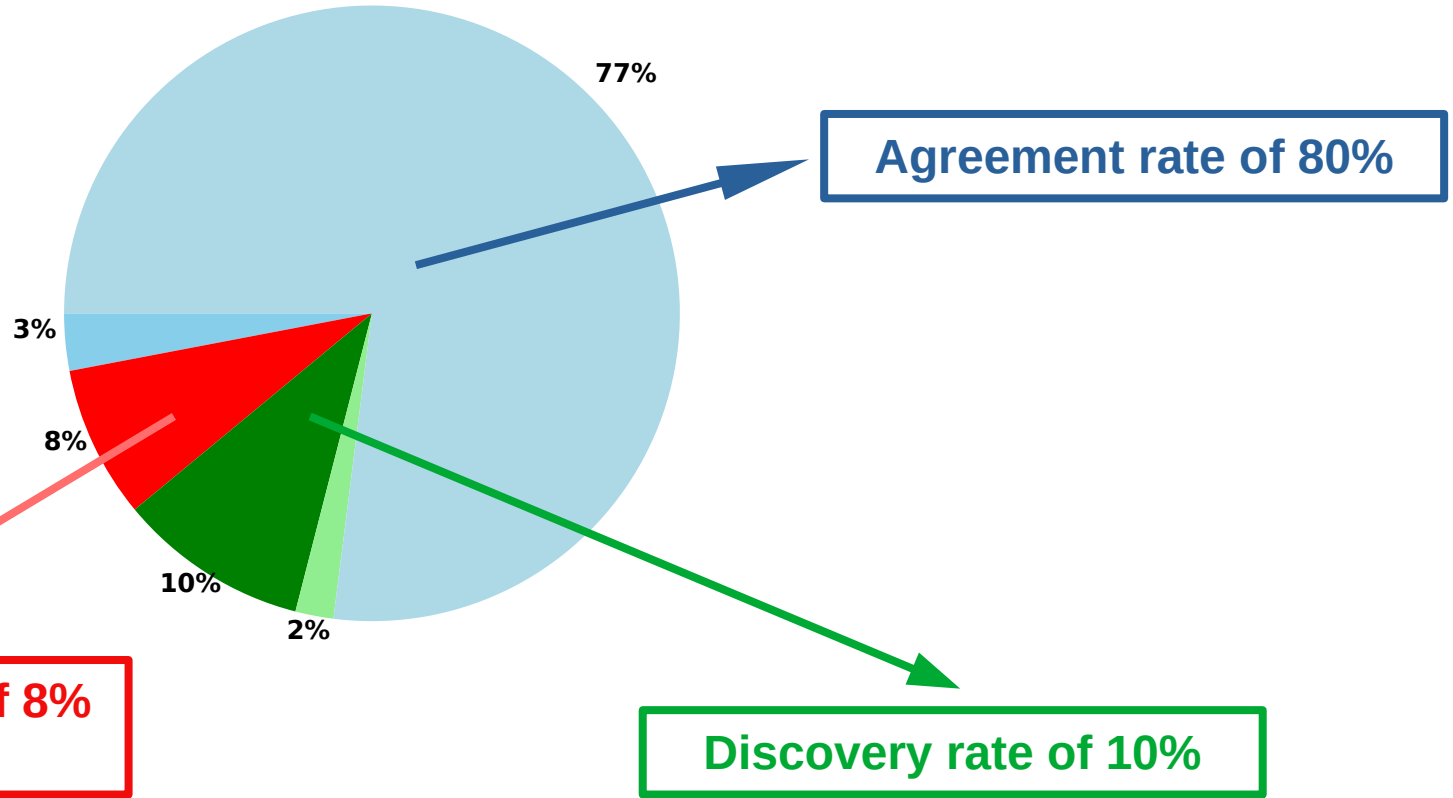
Reference

ML model	accuracy	recall	precision
constant	70.32	100	70.32
uniform	46.27	41.50	69.79
SVC	90.04	95.24	91.06
AdaBoost	86.96	92.92	89.06
DT	84.36	89.57	87.90
RFC	86.65	93.22	88.47
MLP	89.40	93.77	91.39

**Around 90% predictive accuracy**

# ML Model Predictive Performance

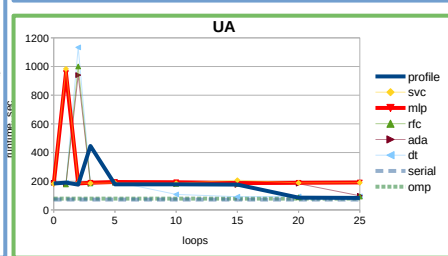
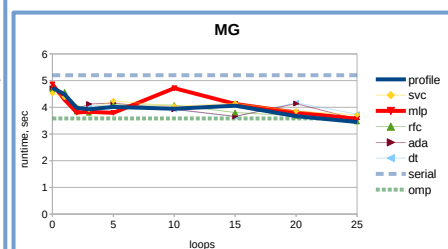
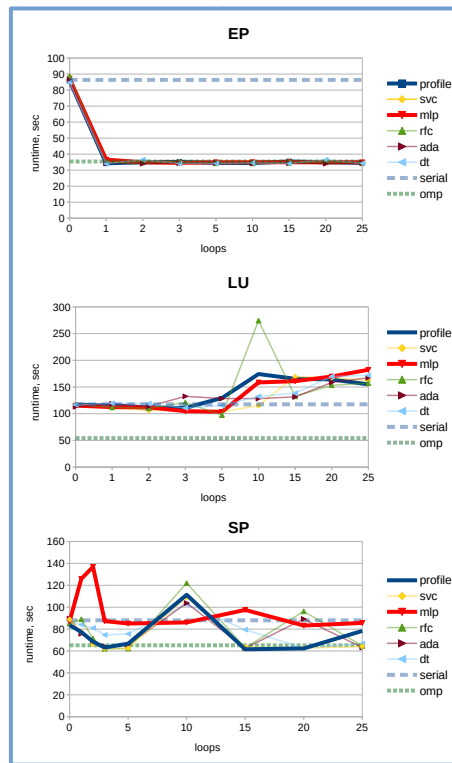
ML Based  
Loop Parallelizability  
Model



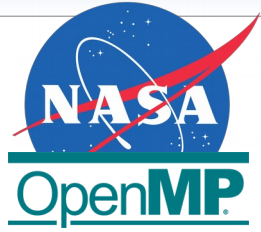
# Summary & Conclusions

**AI-inspired tool and methodology for manual software parallelization have been designed and implemented as a prototype:**

- Loop parallelizability is learnable property (we introduce a machine learning based model and train it to work with the accuracy of above 90%)
- ML model of loop parallelizability has been integrated into parallelization assistant
- Deployed against SNU NPB benchmarks our assistant showed a faster parallelization process over traditional profile-guided method: a programmer is required to inspect and transform 20% less Lines Of Code (LOC)



# Motivating example



## SNU NPB Conjugate Gradient (CG) benchmark

```
for (j = 0; j < lastrow-firstrow+1; j++) {  
    sum1 = 0.0;  
    for (k = rowstr[j]; k < rowstr[j+1]; k++)  
        sum1 = sum1 + a[k]*p[colidx[k]];  
    q[j] = sum1;  
}
```

**cg.c:509**

```
for (it = 1; it <= NITER; it++) {  
    ...  
    if (timeron) timer_start(T_conj_grad);  
    conj_grad(colidx,rowstr,x,z,a,p,q,r,&rn timer);  
    if (timeron) timer_stop(T_conj_grad);  
    ...  
    printf("      %5d      %20.14E%20.13f\n", it,  
        rn timer, zeta);  
    ...  
}
```

**cg.c:326**

### Profiler's Loop Ranking

- 1 **cg.c:326**
- 2 **cg.c:484**
- 3 **cg.c:509**

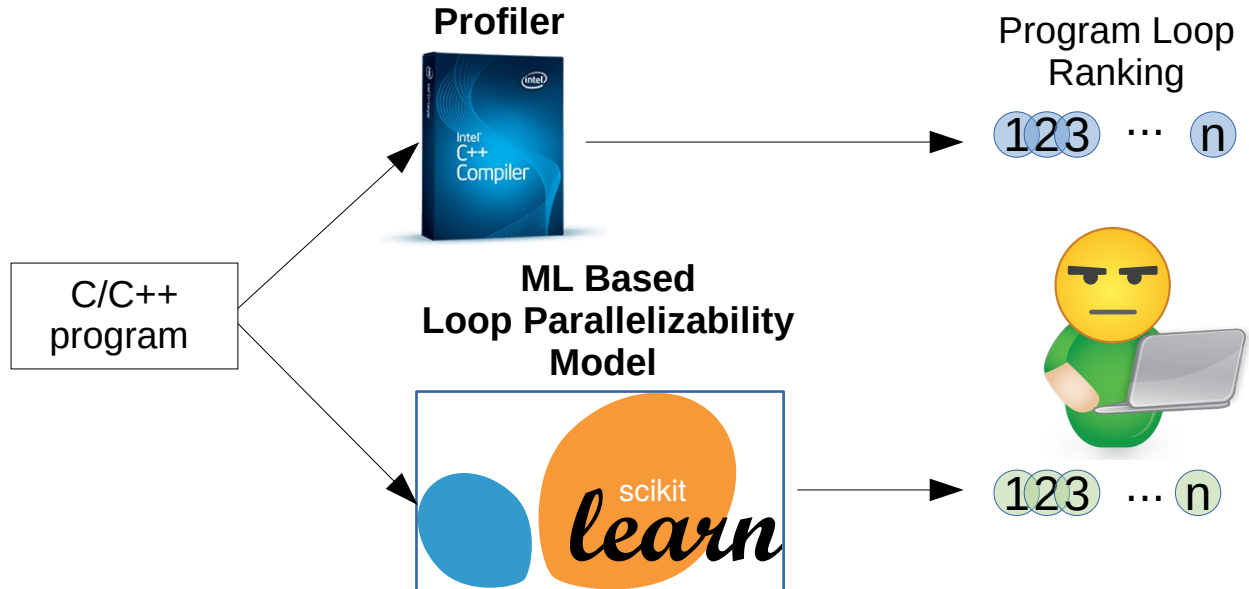
VS.

### Assistant's Loop Ranking

- 1 **cg.c:509**
- 2 **cg.c:326**
- 3 **cg.c:484**

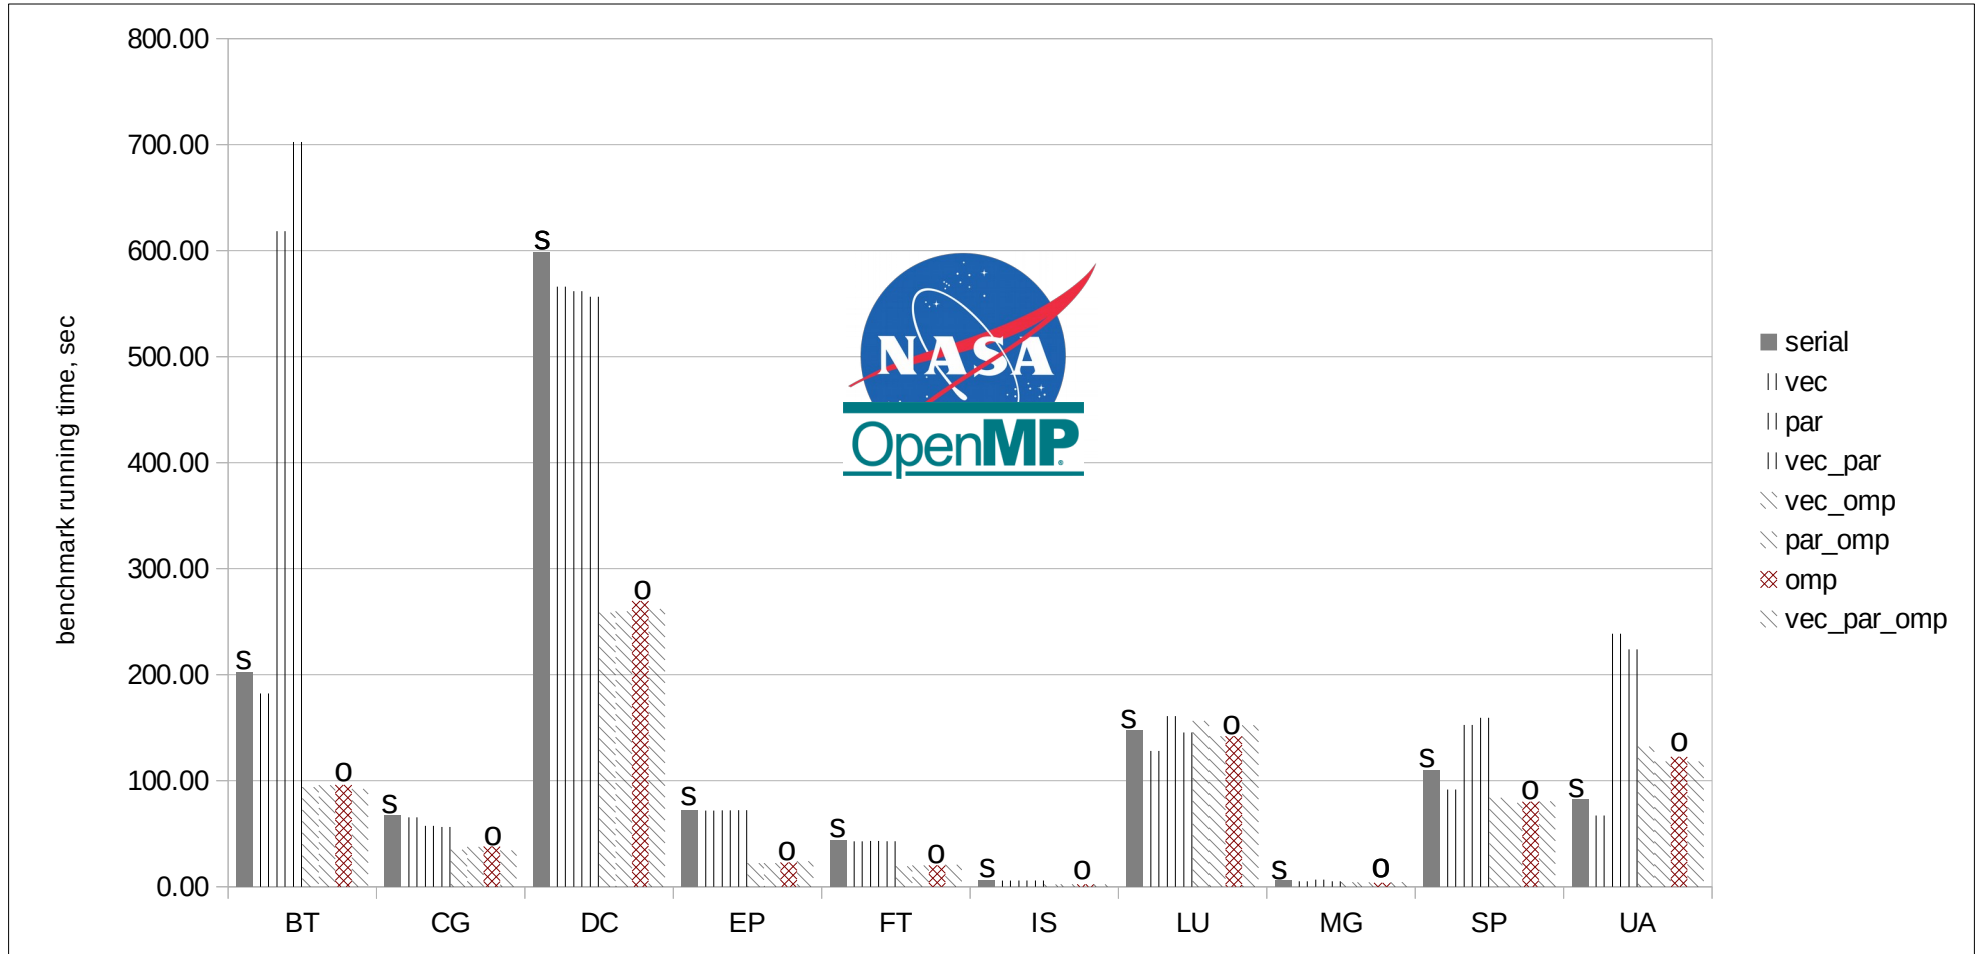
### Assistant's Parallel Loop Probability

- [ 85% ]  
[ 29% ]  
[ 8% ]

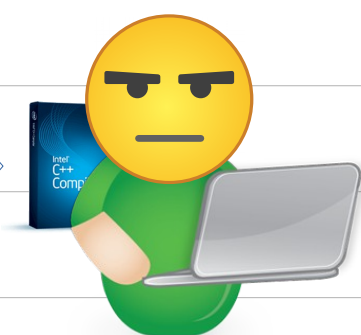
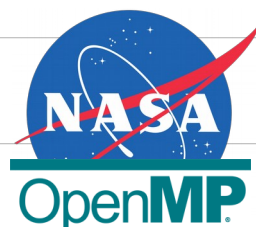


# Problem Statement 2/2

## NAS Parallel Benchmarks (NPB)



## NAS Parallel Benchmarks (NPB)

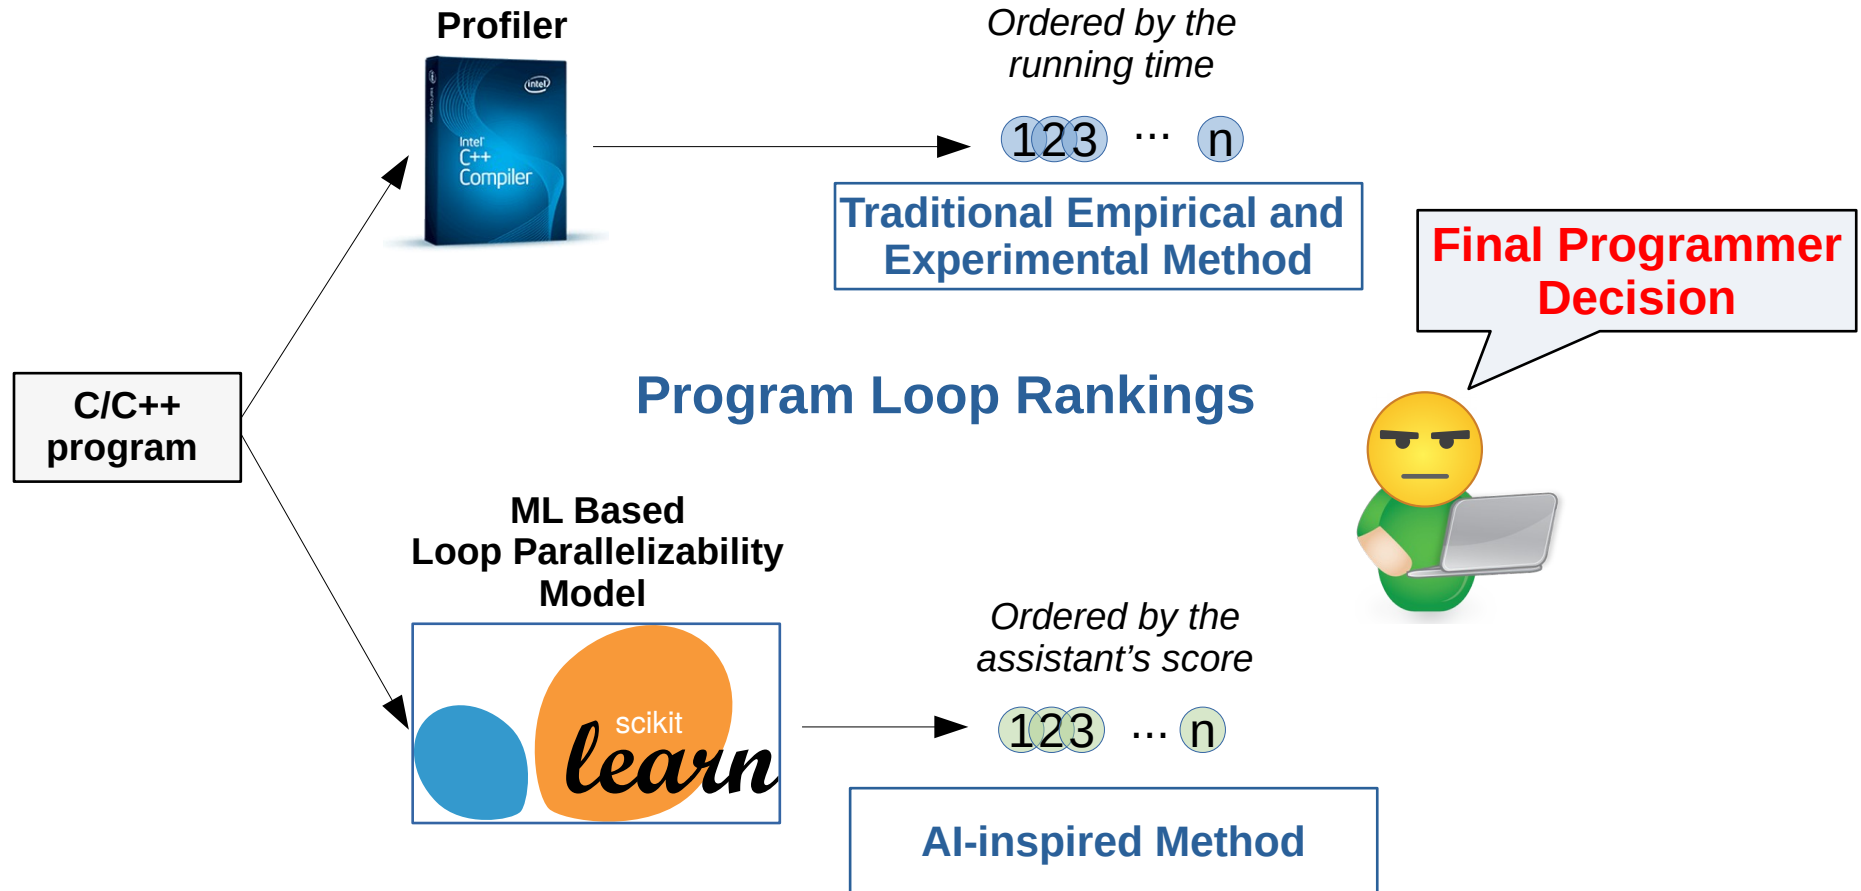


- serial
- || vec
- || par
- || vec\_par
- \\ vec\_omp
- \\ par\_omp
- ⊠ omp
- \\ vec par omp

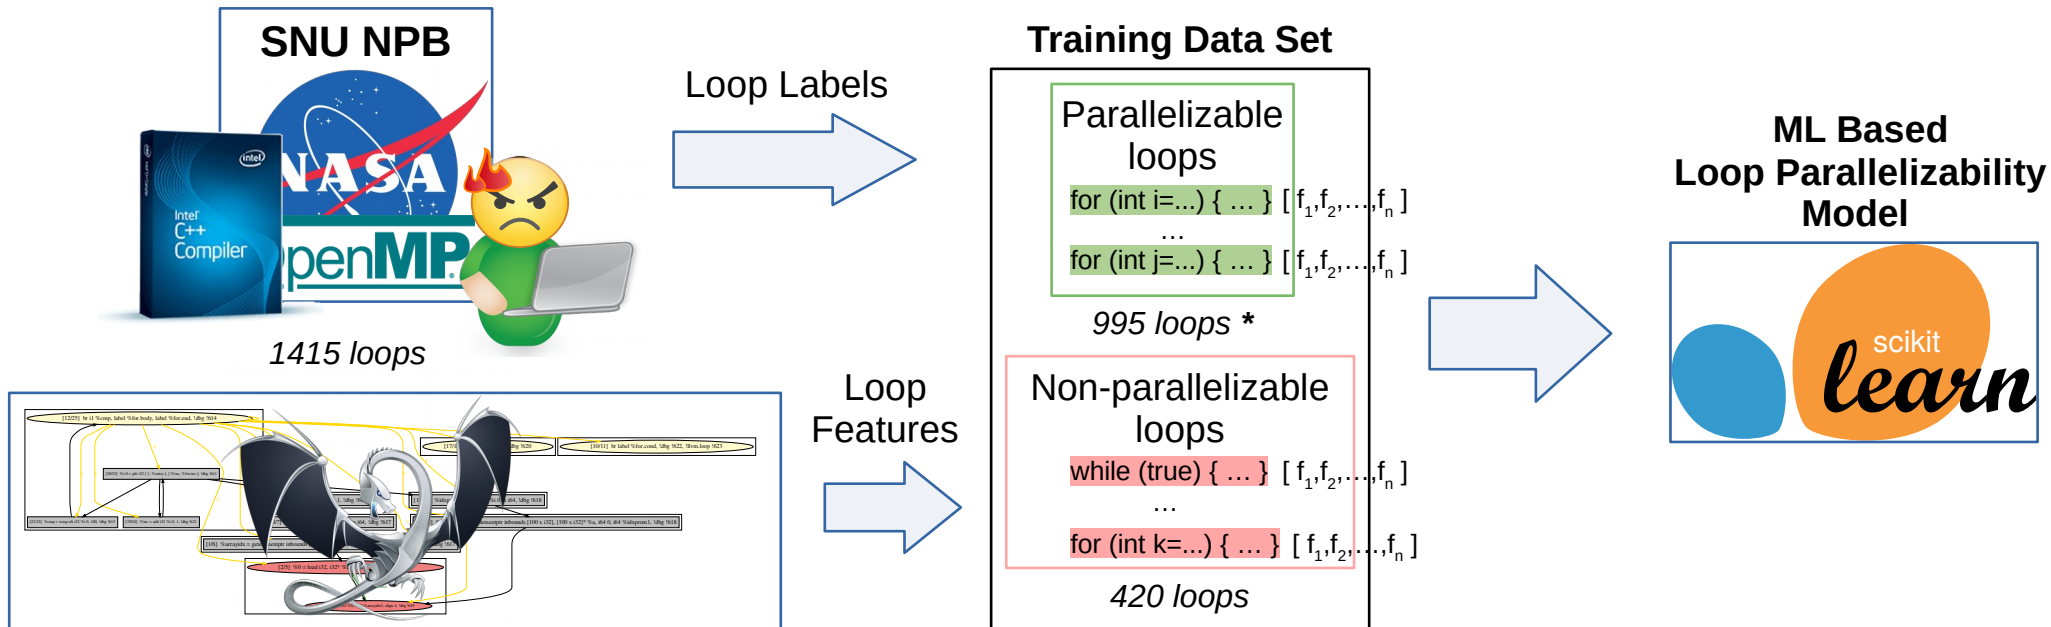


# How to use the assistant

(from the programmers perspective)



# Solution Scheme [1/2]



\* Intel Compiler succeeds in parallelizing 812 loops

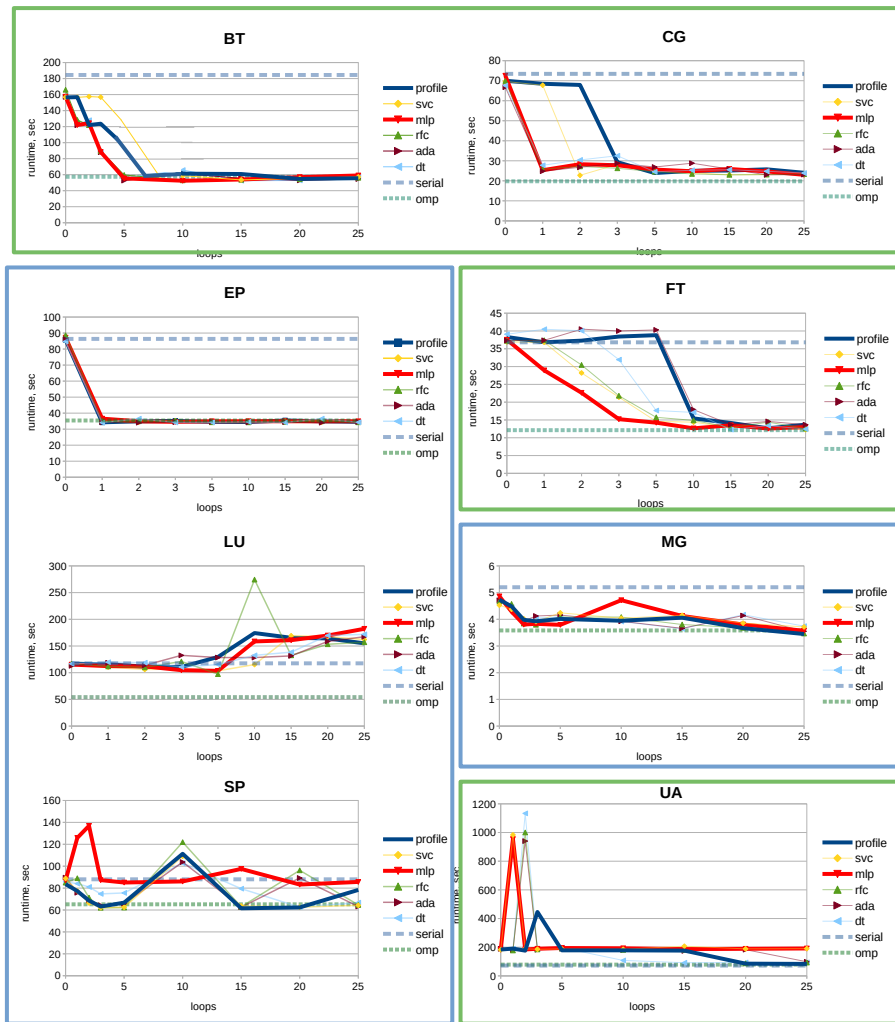
**Loop features** are based on static structural properties of loop program dependence graphs (PDGs):

- Absolute size
- Loop Iterator/Payload cohesion
- Number of dependence edges
- Instruction types (calls, loads/stores, etc.)

**Loop labels** are derived out of OpenMP pragmas present in parallelized SNU NPB versions as well as from the Intel Compiler's parallelization/vectorization reports.

# How well do we do?

Improvement in  
4 benchmarks



No change in  
4 other  
benchmarks

# Problem Statement

## NAS Parallel Benchmarks (NPB)

