

Metric Based Code Parallelisability Analyzer




THE UNIVERSITY *of* EDINBURGH
informatics

EPSRC Centre for Doctoral Training in
Pervasive Parallelism

Basic Idea

Software quality

Cyclomatic Complexity (CC)



Software parallelisability

```
int a[1000];
int b[1000];
int c;

// initialize arrays

c = 0;
for (int i = 0; i < 1000; i++) {
    c += a[i] * b[i];
}
```


Good
Score

```
Node *a;
Node *b;
int c;

// initialize linked lists

0;
while (a && b) {
    c += a->val * b->val;
    a = a->next;
    b = b->next;
}
```

Bad
Score



*.cpp

Compiler

Parallelisable?

YES

NO



optimized.cpp



*.cpp



*.cpp

PPar Tool

88%
parallelisable



Loop Proportion Metric Set:

loop-absolute-size: 22
loop-payload-fraction: 0,55
loop-proper-sccs-number: 1

Loop Cohesion Metric Set:

iterator-payload-total-cohesion: 13 %
iterator-payload-non-cf-cohesion: 2,17 %
iterator-payload-mem-cohesion: 0 %
critical-payload-total-cohesion: 29,4 %
critical-payload-non-cf-cohesion: 11,8 %
critical-payload-mem-cohesion: 0 %

Current Project State (~4-5 months)

- Devised an initial set of parallelisability metrics (*17 metrics*)
 - Developed a tool for metrics collection (*~4300 lines of C++ code*) (*used iterator recognition work as a reference*)
 - Prepared a testing framework (*kindly provided by Christos Vasiladiotis*), allowing to run NAS benchmarks through developed tool
 - Tool and initial metrics are debugged and validated with small hand-written test programs
-
- Tool developed as a dynamic library (.so) to be plugged into LLVM 6.0
 - Uses standard LLVM passes to build dependence graph of the source code
 - Built graphs might be printed and studied by a human reader
 - Tool computes a set of metrics on the built dependence graphs

Things to be done

Minimal plan (has to be done in ~2,5 months)

- Learn NAS benchmarks (11 benchmarks)
- Debug and validate metrics on NAS benchmarks
- Analyze received results
- Write MSc by Research thesis

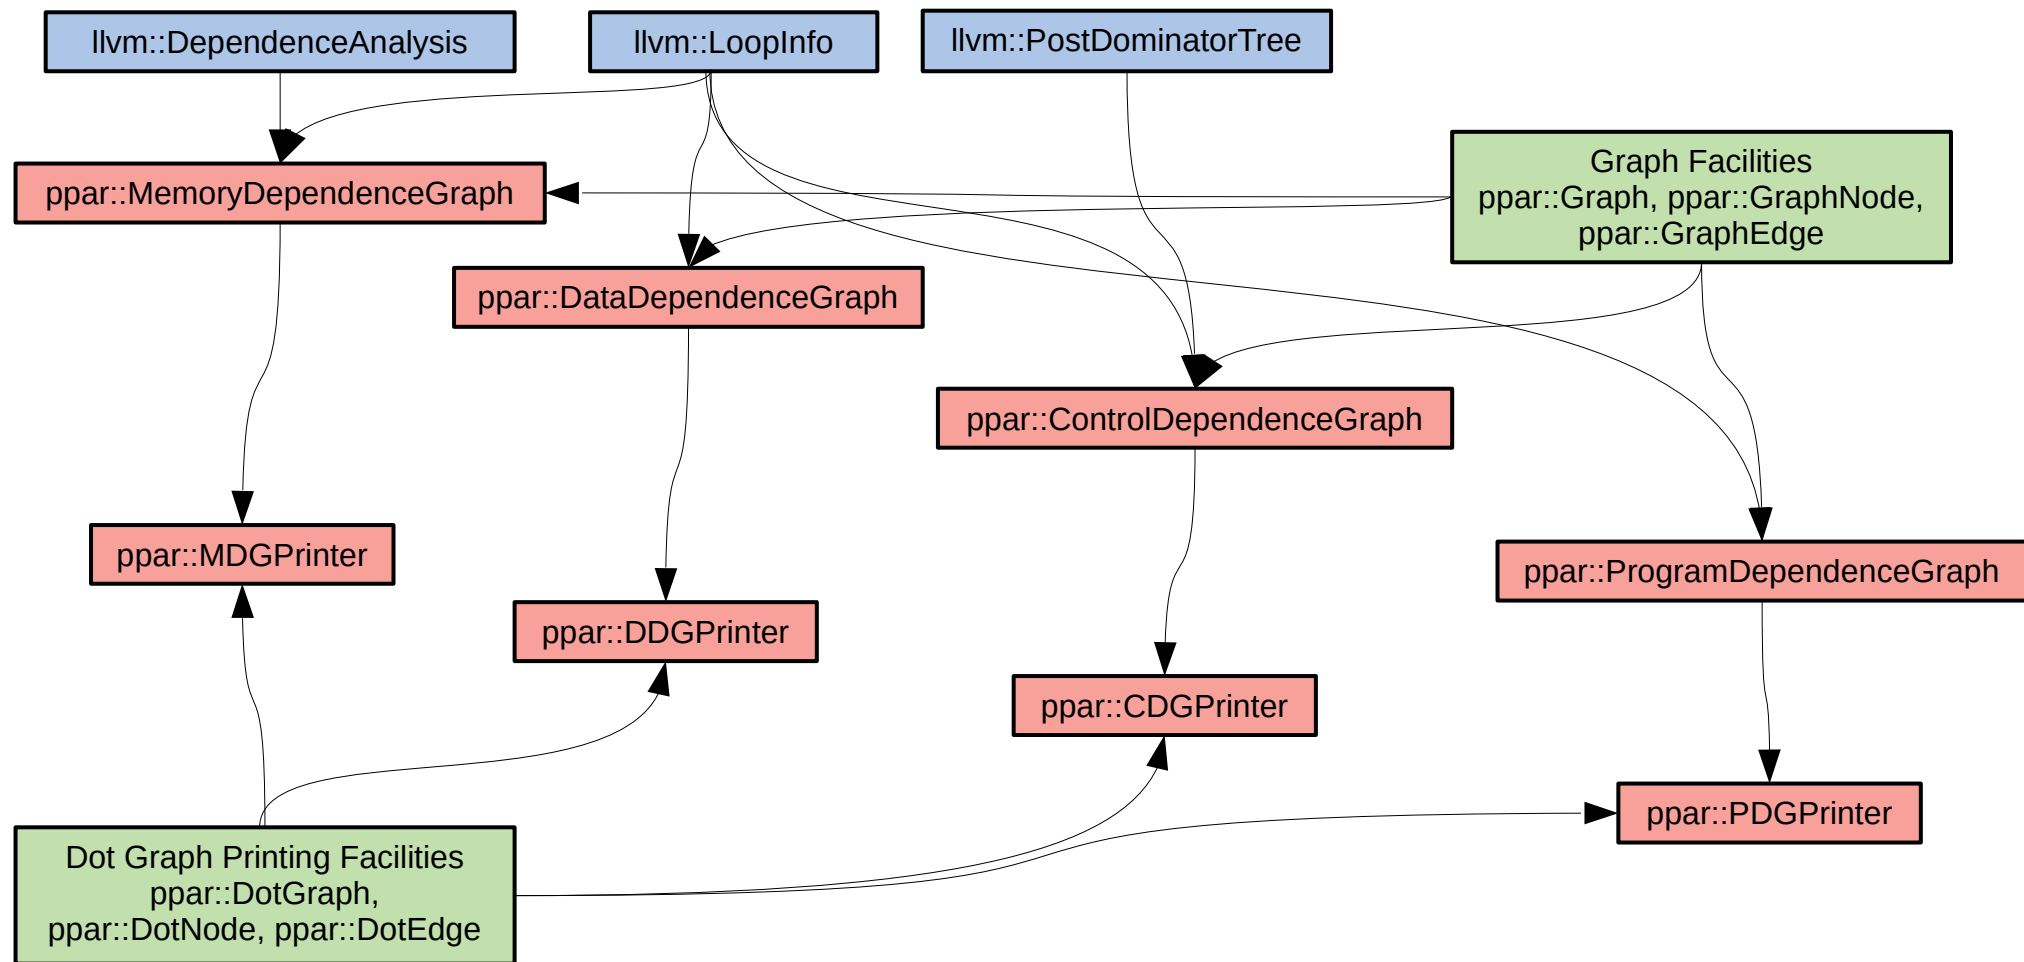
Desired plan (requires more than ~2,5 months)

- Add more benchmarks
- Add more metrics
- Derive more results and correlations
- Add machine-learning part
- Integrate into IDE
- If it all shapes into something promising... Make a publication?

Technical details

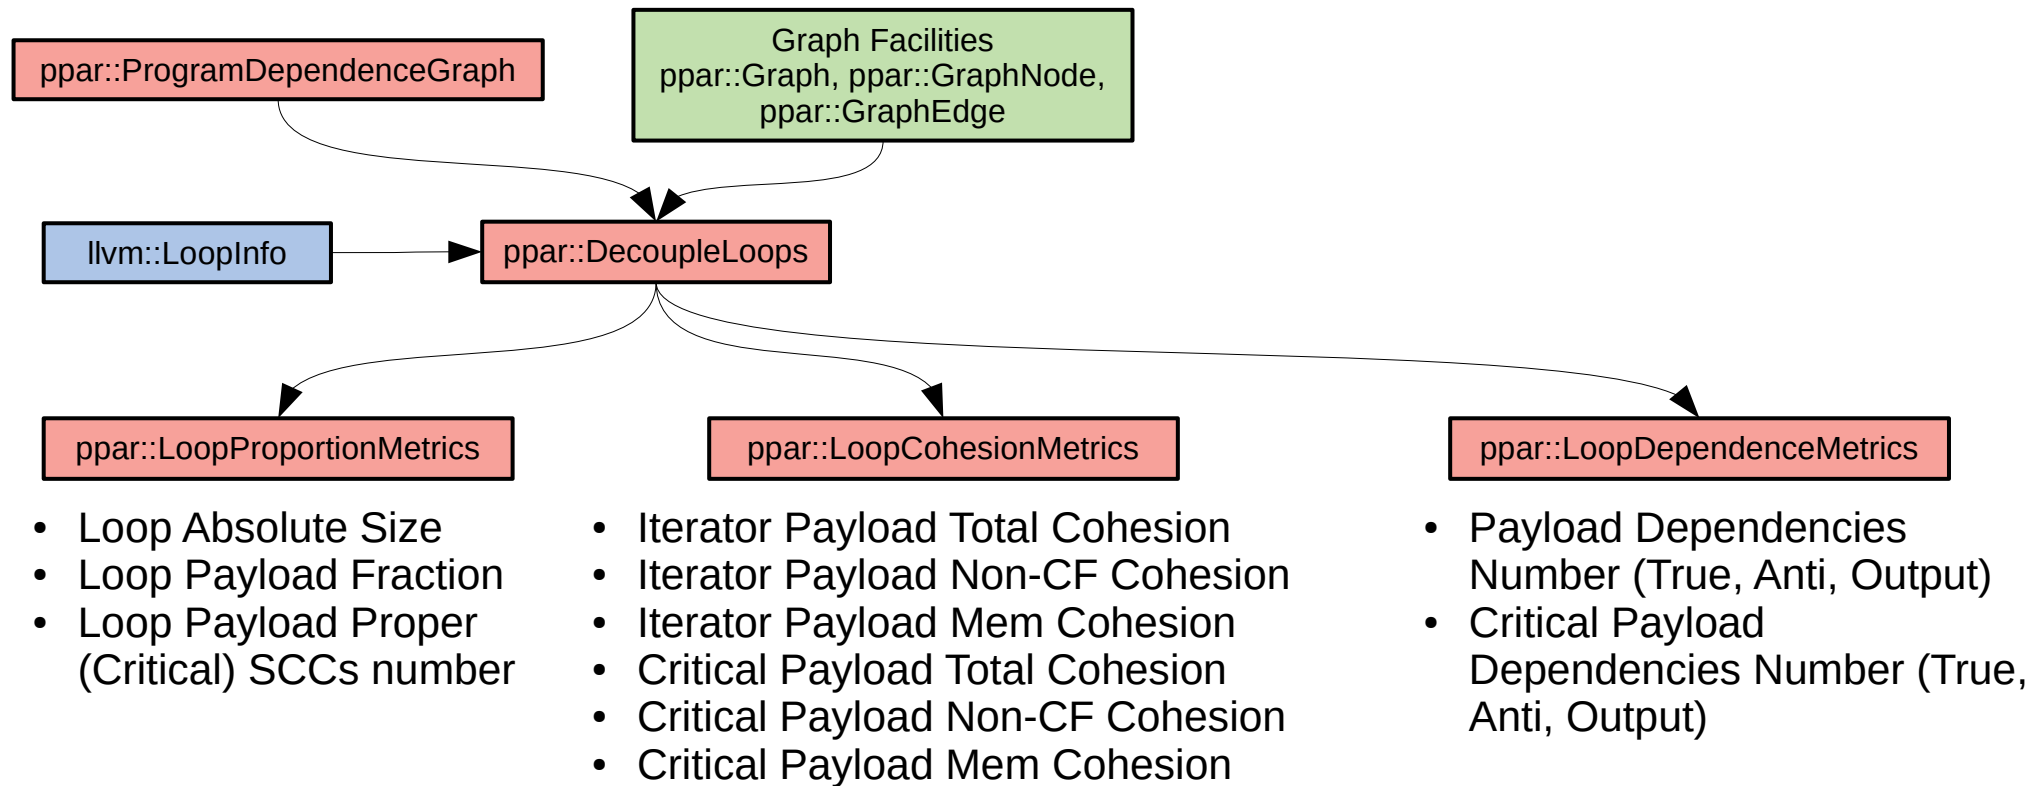
Software Architecture of Graph Building and Printing Facilities

A collection of LLVM passes



Software Architecture of Metrics Collecting Facilities


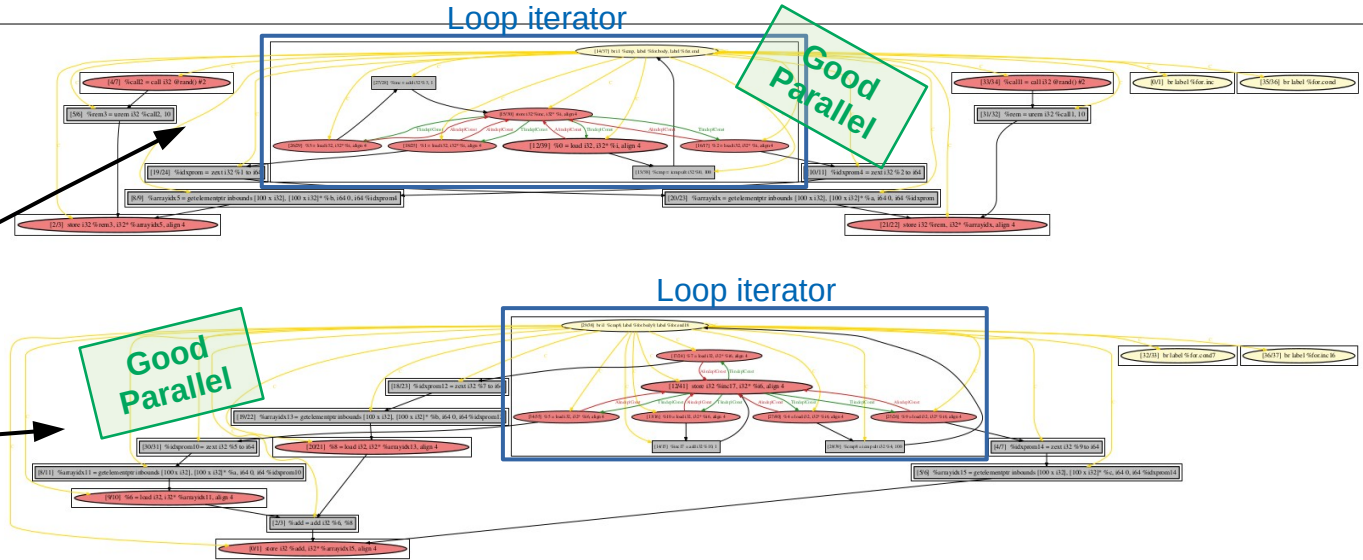
A collection of LLVM passes



Tool workflow

Tool builds Data Dependence Graph (DDG) + Memory Dependence Graph (MDG) + Control Dependence Graph (CDG) and combines them into = Program Dependence Graph (PDG) for all functions and for all loops inside functions. Then it computes a set of metrics on the built graphs.

```
1 #include <cstdlib>
2 #include <ctime>
3 #include <cmath>
4
5 using namespace std;
6
7 static const unsigned int size = 100;
8 static const unsigned int range = 10;
9
10 int main() {
11     unsigned int a[size];
12     unsigned int b[size];
13     unsigned int c[size];
14
15     std::srand(std::time(nullptr));
16
17     // initialization loop
18     for (unsigned int i = 0; i < size; i++) {
19         a[i] = rand() % range;
20         b[i] = rand() % range;
21     }
22
23     // vector sum computation
24     for (unsigned int i = 0; i < size; i++) {
25         c[i] = a[i] + b[i];
26     }
27
28     return 0;
29 }
```



Case study [1] : parallelisable loop

```
// vector sum computation
for (unsigned int i = 0; i < size; i++) {
    c[i] = a[i] + b[i];
}
```

[1] Loop at depth 1 containing: %for.cond7<header><exiting>,
%for.body9,%for.inc16<latch>

Loop Proportion Metric Set:

loop-absolute-size: 21

Loop-payload-fraction: 57,14 %

loop-proper-sccs-number: 0

Loop Cohesion Metric Set:

Iterator-payload-total-cohesion: 31,91 %

Iterator-payload-non-cf-cohesion: 6.38 %

Loop Dependence Metric Set:

payload-total-dependencies-number: 9

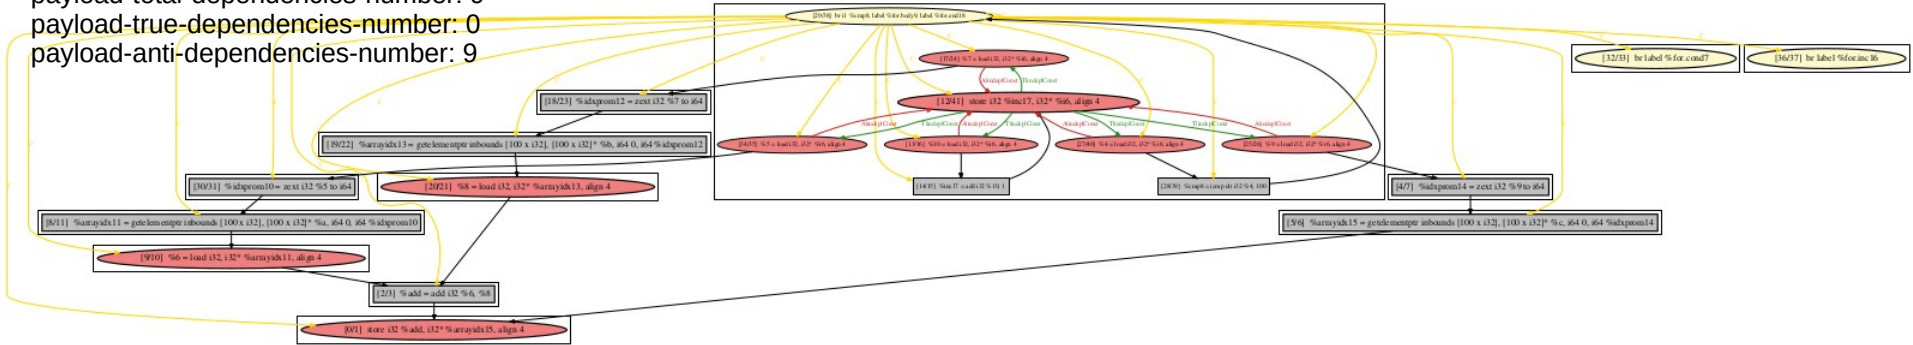
payload-true-dependencies-number: 0

payload-anti-dependencies-number: 9

```
for.cond7:                                     ; preds = %for.inc16, %for.end
    %4 = load i32, @i32, align 4
    %cmp8 = icmp ult i32, %4, 100
    br i1 %cmp8, label %for.body9, label %for.end18

for.body9:                                     ; preds = %for.cond7
    %5 = load i32, @i32, align 4
    %idxprom10 = zext i32, %5 to i64
    %arrayidx11 = getelementptr @inbounds, [100 x i32], [100 x i32]* %a, @i64_0, @i64 %idxprom10
    %6 = load i32, @i32, align 4
    %7 = load i32, @i32, align 4
    %idxprom12 = zext i32, %7 to i64
    %arrayidx13 = getelementptr @inbounds, [100 x i32], [100 x i32]* %b, @i64_0, @i64 %idxprom12
    %8 = load i32, @i32, align 4
    %add = add i32, %6, %8
    %9 = load i32, @i32, align 4
    %idxprom14 = zext i32, %9 to i64
    %arrayidx15 = getelementptr @inbounds, [100 x i32], [100 x i32]* %c, @i64_0, @i64 %idxprom14
    store i32, @i32, align 4, %arrayidx15, align 4
    br label %for.inc16

for.inc16:                                     ; preds = %for.body9
    %10 = load i32, @i32, align 4
    %inc17 = add i32, %10, 1
    store i32, @i32, align 4, %10, align 4
    br label %for.cond7
```



Case study [2] : non-parallelisable loop

```
for.cond:                                ; preds = %for.inc, %entry
%0 = load i32, i32* %i, align 4
%cmp = icmp ult i32 %0, 100
br i1 %cmp, label %for.body, label %for.end
```

```
for.body:                                ; preds = %for.cond
%call = call i8* @Znwm(i64 16) #8
%1 = bitcast i8* %call to %struct.list_node*
call void @ZN9list_nodeC2Ev(%struct.list_node* %1)
br label %invoke.cont
```

```
invoke.cont:                             ; preds = %for.body
%2 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%next = getelementptr inbounds %struct.list_node, %struct.list_node* %2, i32 0, i32 1
store %struct.list_node* %1, %struct.list_node** %next, align 8
%3 = load i32, i32* %i, align 4
%4 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%value = getelementptr inbounds %struct.list_node, %struct.list_node* %4, i32 0, i32 0
store i32 %3, i32* %value, align 8
%5 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%next1 = getelementptr inbounds %struct.list_node, %struct.list_node* %5, i32 0, i32 1
%6 = load %struct.list_node*, %struct.list_node** %next1, align 8
store %struct.list_node* %6, %struct.list_node** %list_it, align 8
br label %for.inc
```

```
for.inc:                                 ; preds = %invoke.cont
%7 = load i32, i32* %i, align 4
%inc = add i32 %7, 1
store i32 %inc, i32* %i, align 4
br label %for.cond
```

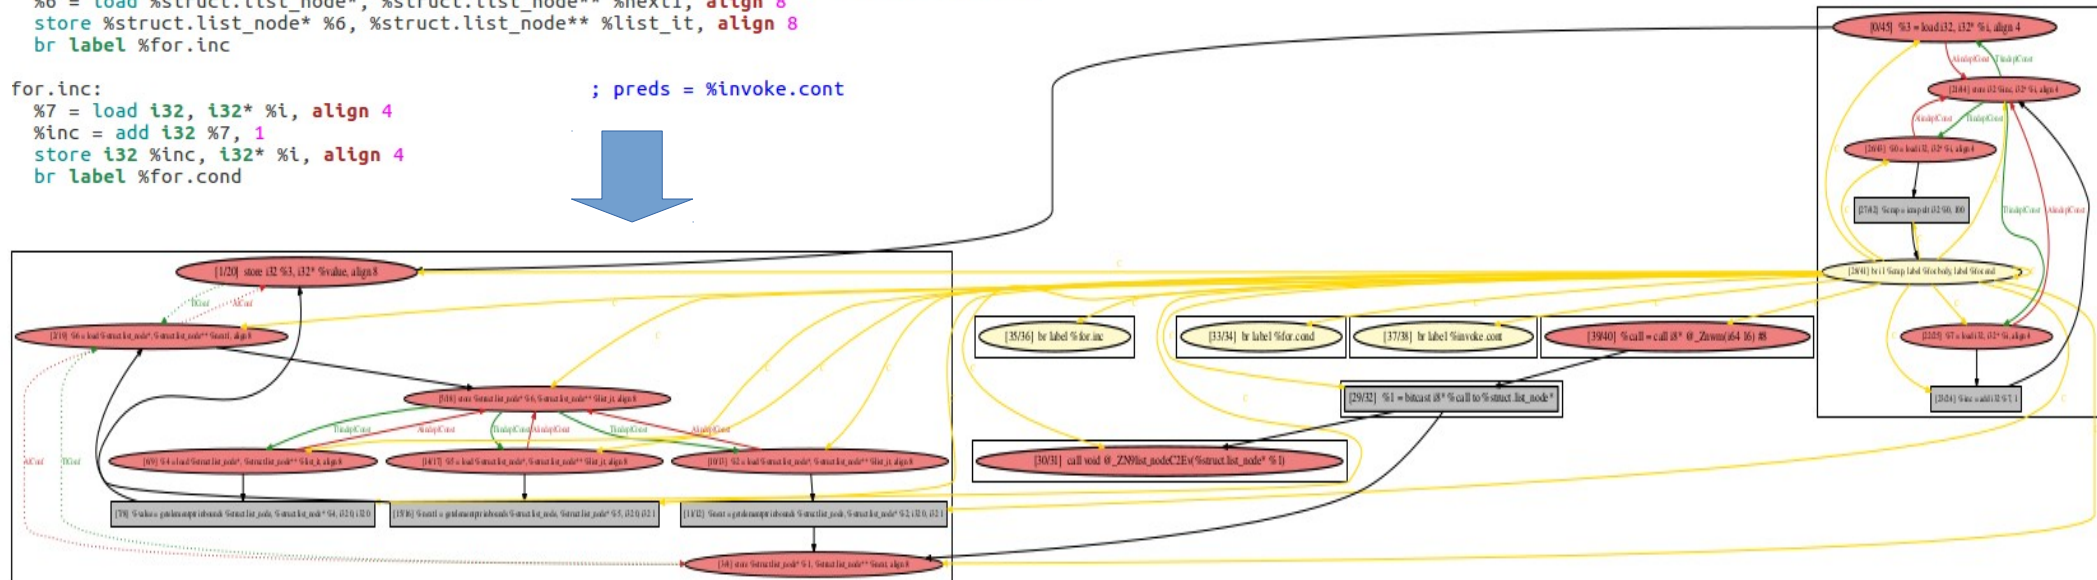
```
list_node_t begin;
list_node_t* list_it;
```

```
list_it = &begin;
for (unsigned int i = 1; i < size; i++) {

    list_it->next = new list_node_t;

    list_it->value = i;

    list_it = list_it->next;
}
```



Case study [2] : non-parallelisable loop

[3] Loop at depth 1 containing: %for.cond<header><exiting>,%for.body,
%invoke.cont,%for.inc<latch>

Loop Proportion Metric Set:

loop-absolute-size: 23
Loop-payload-fraction: 69,57 %
loop-proper-sccs-number: 1

Loop Cohesion Metric Set:

Iterator-payload-total-cohesion: 31,48 %
Iterator-payload-non-cf-cohesion: 1.85%
Iterator-payload-mem-cohesion: 0 %
Critical-payload-total-cohesion: 5 %
Critical-payload-non-cf-cohesion: 5 %
critical-payload-mem-cohesion: 0

Loop Dependence Metric Set:

payload-total-dependencies-number: 3
payload-true-dependencies-number: 0
payload-anti-dependencies-number: 3
critical-payload-total-dependencies-number: 17
critical-payload-true-dependencies-number: 5
critical-payload-anti-dependencies-number: 12

```
list_node_t begin;  
list_node_t* list_it;  
  
list_it = &begin;  
for (unsigned int i = 1; i < size; i++) {  
  
    list_it->next = new list_node_t;  
  
    list_it->value = i;  
  
    list_it = list_it->next;  
}
```




Backup slides


Basic idea

```
1 #include <cstdlib>
2 #include <ctime>
3 #include <cmath>
4
5 using namespace std;
6
7 static const unsigned int size = 100;
8 static const unsigned int range = 10;
9
10 int main() {
11     unsigned int a[size];
12     unsigned int b[size];
13     unsigned int c[size];
14
15     std::srand(std::time(nullptr));
16
17     // initialization loop
18     for (unsigned int i = 0; i < size; i++) {
19         a[i] = rand() % range;
20         b[i] = rand() % range;
21     }
22
23     // vector sum computation
24     for (unsigned int i = 0; i < size; i++) {
25         c[i] = a[i] + b[i];
26     }
27
28     return 0;
29 }
```


Function PDG



Loop0 PDG



Loop1 PDG



llvm::DependenceAnalysis

- LLVM pass which analyses dependencies between memory accesses. As of LLVM 6.0 it is an implementation of the approach described in [1].
- llvm::DependenceAnalysis pass exists to support llvm::DependenceGraph pass (*does not exist in LLVM yet*).
- There are two separate passes, because it is a useful separation of concerns. A dependence exists if two conditions are met:
 - 1) Two instructions reference the same memory location
 - 2) There is a flow of control, leading from one instruction to the other.
- llvm::DependenceAnalysis addresses the first condition, llvm::DependenceGraph (*not available as a standard LLVM pass yet, as of LLVM 6.0*) addresses the second.
- llvm::DependenceAnalysis is *Work In Progress* in LLVM

[1] Practical Dependence Testing. Gina Goff, Ken Kennedy, Chau-Wen Tseng. Department of Computer Science, Rice University, Houston, TX. Proceedings of the ACM SIGPLAN '91 Conference on Programming Language Design and Implementation. Toronto, Ontario, Canada, June 26-28, 1991.

llvm::DependenceAnalysis

```
class Dependence {  
    Dependence(Instruction *Source,  
               Instruction *Destination);
```

```
    bool isInput() const;  
    bool isOutput() const;  
    bool isFlow() const;  
    bool isAnti() const;  
    bool isOrdered() const { return isOutput() || isFlow() || isAnti(); }  
    bool isUnordered() const { return isInput(); }
```

```
    // isLoopIndependent - should be set by the caller if it appears that control flow can reach  
    // from Src to Dst without traversing a loop back edge.  
    virtual bool isLoopIndependent() const { return true; }  
    // isConfused - Returns true if this dependence is confused  
    // (the compiler understands nothing and makes worst-case assumptions).  
    virtual bool isConfused() const { return true; }  
    // isConsistent - Returns true if this dependence is consistent  
    // (occurs every time the source and destination are executed).  
    virtual bool isConsistent() const { return false; }  
    // getLevels - Returns the number of common loops surrounding the  
    // source and destination of the dependence.  
    virtual unsigned getLevels() const { return 0; }  
    // getDirection - Returns the direction associated with a particular level.  
    virtual unsigned getDirection(unsigned Level) const { return DVEntry::ALL; }  
    // getDistance - Returns the distance (or NULL) associated with a particular level.  
    virtual const SCEV *getDistance(unsigned Level) const { return nullptr; }  
    // isPeelFirst - Returns true if peeling the first iteration from this loop will break this dependence.  
    virtual bool isPeelFirst(unsigned Level) const { return false; }  
    // isPeelLast - Returns true if peeling the last iteration from this loop will break this dependence.  
    virtual bool isPeelLast(unsigned Level) const { return false; }  
    // isSplittable - Returns true if splitting this loop will break the dependence.  
    virtual bool isSplittable(unsigned Level) const { return false; }  
    // isScalar - Returns true if a particular level is scalar; that is, if no subscript in the source or  
    // destination mention the induction variable associated with the loop at this level.  
    virtual bool isScalar(unsigned Level) const;  
};
```

- LLVM class Dependence represents a dependence between two memory references in a function.

llvm::DependenceGraph

*Not yet ready in LLVM!
Work In Progress!*

- Generally, the dependence analyzer will be used to build a dependence graph for a function. Looking for cycles in the graph shows us loops, that cannot be trivially vectorized/parallelized
- We can try to improve the situation by examining all the dependences that make up the cycle, looking for ones we can break. Sometimes, peeling the first or last iteration of a loop will break dependences, and there are flags for those possibilities. Sometimes, splitting a loop at some other iteration will do the trick, and we've got a flag for that case. Rather than waste the space to record the exact iteration (since we rarely know), we provide a method that calculates the iteration. It's a drag that it must work from scratch, but wonderful in that it's possible.

Here's an example:

```
for (i = 0; i < 10; i++)  
    A[i] = ...  
    ... = A[11 - i]
```

There's a loop-carried flow dependence from the store to the load, found by the weak-crossing SIV test. The dependence will have a flag, indicating that the dependence can be broken by splitting the loop. Calling `getSplitIteration` will return 5. Splitting the loop breaks the dependence, like so:

```
for (i = 0; i <= 5; i++)  
    A[i] = ...  
    ... = A[11 - i]  
for (i = 6; i < 10; i++)  
    A[i] = ...  
    ... = A[11 - i]
```

breaks the dependence and allows us to vectorize/parallelize both loops.

Case study [2] : non-parallelisable loop

```
list_node_t begin;  
list_node_t* list_it;
```

```
list_it = &begin; Iterator
```

```
for (unsigned int i = 1; i < size; i++) {
```

```
    list_it->next = new list_node_t;
```

```
    list_it->value = i;
```

```
    list_it = list_it->next;
```

```
}
```

Payload

```
list_node_t begin;  
list_node_t* list_it;
```

```
list_it = &begin;
```

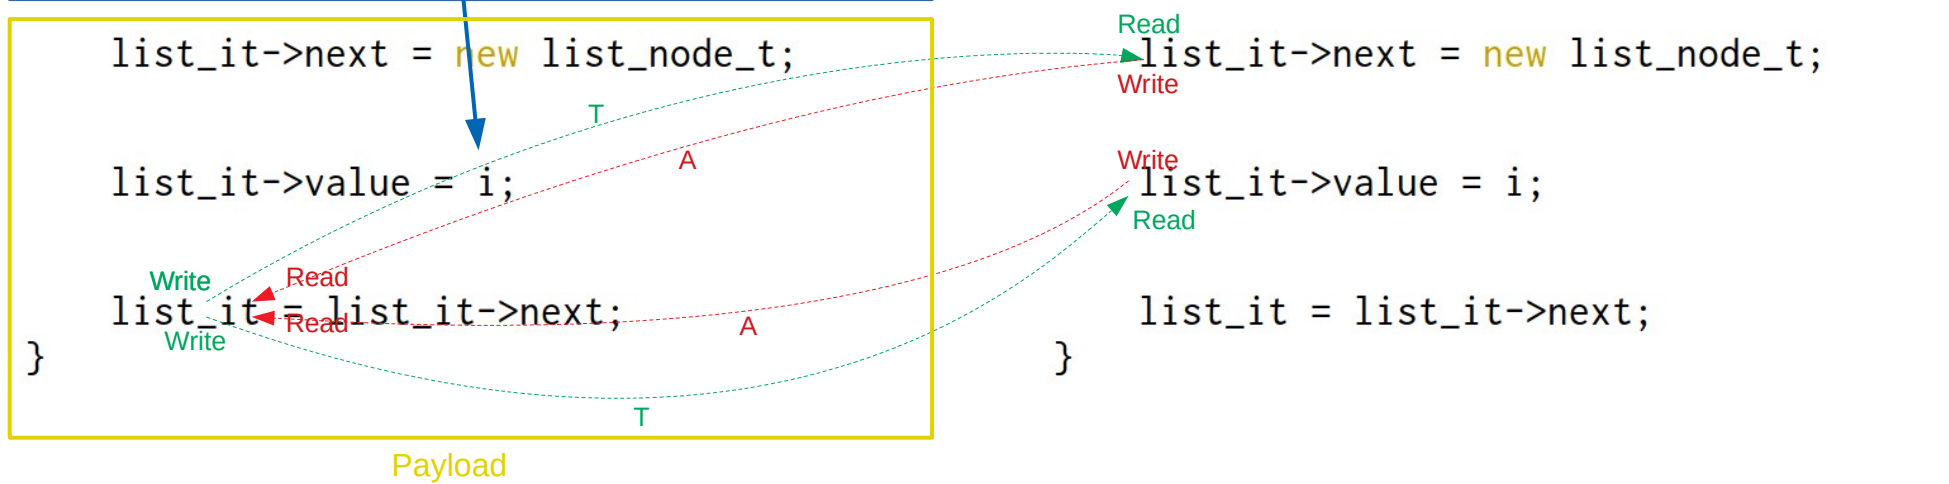
```
for (unsigned int i = 1; i < size; i++) {
```

```
    list_it->next = new list_node_t;
```

```
    list_it->value = i;
```

```
    list_it = list_it->next;
```

```
}
```



General work observation

- To judge about true program parallelisability, we have to disassemble the program (its program dependence graph) to the smallest finest-grain pieces
- As it appears at the moment, the true indicator of the loop parallelisability is the absence of strongly connected components (besides the iterator one) of the size, greater than 1 instruction: there are cycles present in such strongly connected components, which imply dependence between instructions tying up/entangling loops and preventing them from parallelisation
- According to code comments, there is a work going on in LLVM on building `llvm::DependenceGraph` and breaking its edges with loop splitting and peeling.