


# Machine Learning Based Code Parallelisability Analyzer


# Basic idea

```
1 #include <cstdlib>
2 #include <ctime>
3 #include <cmath>
4
5 using namespace std;
6
7 static const unsigned int size = 100;
8 static const unsigned int range = 10;
9
10 int main() {
11     unsigned int a[size];
12     unsigned int b[size];
13     unsigned int c[size];
14
15     std::srand(std::time(nullptr));
16
17     // initialization loop
18     for (unsigned int i = 0; i < size; i++) {
19         a[i] = rand() % range;
20         b[i] = rand() % range;
21     }
22
23     // vector sum computation
24     for (unsigned int i = 0; i < size; i++) {
25         c[i] = a[i] + b[i];
26     }
27
28     return 0;
29 }
```


Function PDG



Loop0 PDG



Loop1 PDG




# Pervasive Parallelism (PPar) Code Parallelisability Analyzer

- A tool to be used for software parallelisability assessment
- Developed as a dynamic library (.so) to be plugged into LLVM 6.0
- Uses standard LLVM passes to build dependence graph of the source code
- Built graphs might be printed and studied
- Tool computes a set of metrics on the built dependence graphs

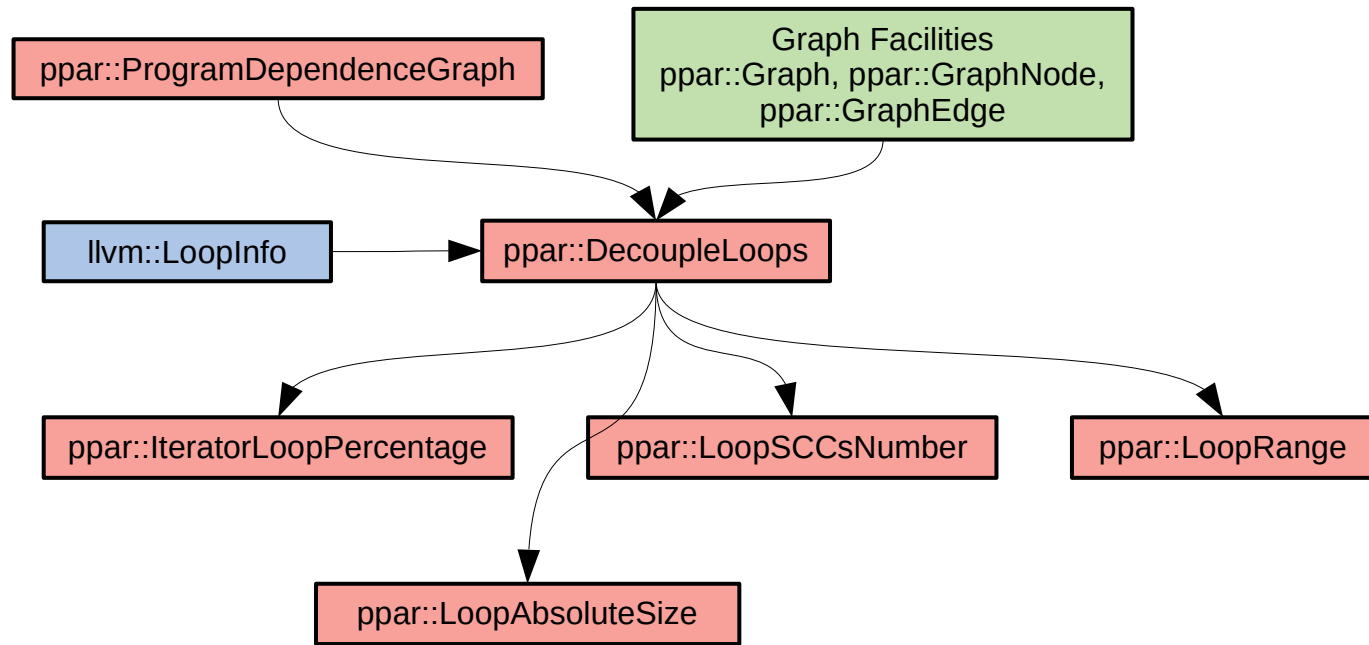
# Software Architecture of Graph Building and Printing Facilities

A collection of LLVM passes



# Software Architecture of Metrics Collecting Facilities

A collection of LLVM passes



## llvm::DependenceAnalysis

- LLVM pass which analyses dependencies between memory accesses. As of LLVM 6.0 it is an implementation of the approach described in [1].
- llvm::DependenceAnalysis pass exists to support llvm::DependenceGraph pass (*does not exist in LLVM yet*).
- There are two separate passes, because it is a useful separation of concerns. A dependence exists if two conditions are met:
  - 1) Two instructions reference the same memory location
  - 2) There is a flow of control, leading from one instruction to the other.
- llvm::DependenceAnalysis addresses the first condition, llvm::DependenceGraph (*not available as a standard LLVM pass yet, as of LLVM 6.0*) addresses the second.
- llvm::DependenceAnalysis is *Work In Progress* in LLVM

[1] Practical Dependence Testing. Gina Goff, Ken Kennedy, Chau-Wen Tseng. Department of Computer Science, Rice University, Houston, TX. Proceedings of the ACM SIGPLAN '91 Conference on Programming Language Design and Implementation. Toronto, Ontario, Canada, June 26-28, 1991.

## llvm::DependenceAnalysis

```
class Dependence {  
    Dependence(Instruction *Source,  
               Instruction *Destination);
```

```
    bool isInput() const;  
    bool isOutput() const;  
    bool isFlow() const;  
    bool isAnti() const;  
    bool isOrdered() const { return isOutput() || isFlow() || isAnti(); }  
    bool isUnordered() const { return isInput(); }
```

```
    // isLoopIndependent - should be set by the caller if it appears that control flow can reach  
    // from Src to Dst without traversing a loop back edge.  
    virtual bool isLoopIndependent() const { return true; }  
    // isConfused - Returns true if this dependence is confused  
    // (the compiler understands nothing and makes worst-case assumptions).  
    virtual bool isConfused() const { return true; }  
    // isConsistent - Returns true if this dependence is consistent  
    // (occurs every time the source and destination are executed).  
    virtual bool isConsistent() const { return false; }  
    // getLevels - Returns the number of common loops surrounding the  
    // source and destination of the dependence.  
    virtual unsigned getLevels() const { return 0; }  
    // getDirection - Returns the direction associated with a particular level.  
    virtual unsigned getDirection(unsigned Level) const { return DVEntry::ALL; }  
    // getDistance - Returns the distance (or NULL) associated with a particular level.  
    virtual const SCEV *getDistance(unsigned Level) const { return nullptr; }  
    // isPeelFirst - Returns true if peeling the first iteration from this loop will break this dependence.  
    virtual bool isPeelFirst(unsigned Level) const { return false; }  
    // isPeelLast - Returns true if peeling the last iteration from this loop will break this dependence.  
    virtual bool isPeelLast(unsigned Level) const { return false; }  
    // isSplittable - Returns true if splitting this loop will break the dependence.  
    virtual bool isSplittable(unsigned Level) const { return false; }  
    // isScalar - Returns true if a particular level is scalar; that is, if no subscript in the source or  
    // destination mention the induction variable associated with the loop at this level.  
    virtual bool isScalar(unsigned Level) const;  
};
```

- LLVM class Dependence represents a dependence between two memory references in a function.

## llvm::DependenceGraph

*Not yet ready in LLVM!  
Work In Progress!*

- Generally, the dependence analyzer will be used to build a dependence graph for a function. Looking for cycles in the graph shows us loops, that cannot be trivially vectorized/parallelized
- We can try to improve the situation by examining all the dependences that make up the cycle, looking for ones we can break. Sometimes, peeling the first or last iteration of a loop will break dependences, and there are flags for those possibilities. Sometimes, splitting a loop at some other iteration will do the trick, and we've got a flag for that case. Rather than waste the space to record the exact iteration (since we rarely know), we provide a method that calculates the iteration. It's a drag that it must work from scratch, but wonderful in that it's possible.

Here's an example:

```
for (i = 0; i < 10; i++)  
    A[i] = ...  
    ... = A[11 - i]
```

There's a loop-carried flow dependence from the store to the load, found by the weak-crossing SIV test. The dependence will have a flag, indicating that the dependence can be broken by splitting the loop. Calling `getSplitIteration` will return 5. Splitting the loop breaks the dependence, like so:

```
for (i = 0; i <= 5; i++)  
    A[i] = ...  
    ... = A[11 - i]  
for (i = 6; i < 10; i++)  
    A[i] = ...  
    ... = A[11 - i]
```

breaks the dependence and allows us to vectorize/parallelize both loops.



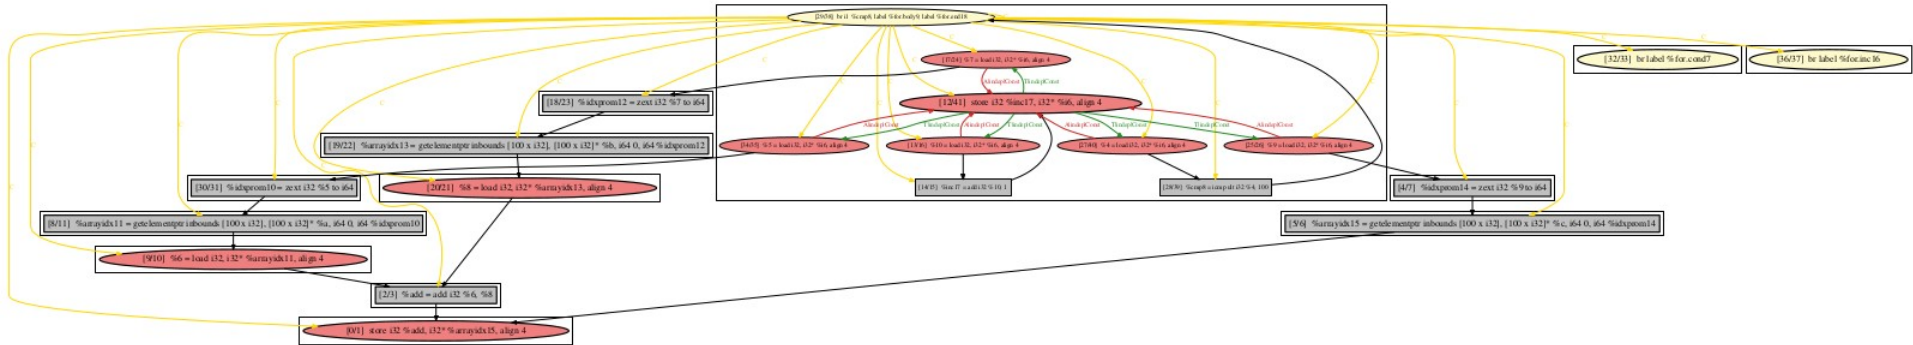
# Case study [1] : parallelisable loop

```
// vector sum computation
for (unsigned int i = 0; i < size; i++) {
    c[i] = a[i] + b[i];
}
```

```
for.cond7:
    %4 = load i32, i32* %i6, align 4          ; preds = %for.inc16, %for.end
    %cmp8 = icmp ult i32 %4, 100
    br i1 %cmp8, label %for.body9, label %for.end18

for.body9:
    %5 = load i32, i32* %i6, align 4          ; preds = %for.cond7
    %idxprom10 = zext i32 %5 to i64
    %arrayidx11 = getelementptr @inbounds [100 x i32], [100 x i32]* %a, i64 0, i64 %idxprom10
    %6 = load i32, i32* %arrayidx11, align 4
    %7 = load i32, i32* %i6, align 4
    %idxprom12 = zext i32 %7 to i64
    %arrayidx13 = getelementptr @inbounds [100 x i32], [100 x i32]* %b, i64 0, i64 %idxprom12
    %8 = load i32, i32* %arrayidx13, align 4
    %add = add i32 %6, %8
    %9 = load i32, i32* %i6, align 4
    %idxprom14 = zext i32 %9 to i64
    %arrayidx15 = getelementptr @inbounds [100 x i32], [100 x i32]* %c, i64 0, i64 %idxprom14
    store i32 %add, i32* %arrayidx15, align 4
    br label %for.inc16

for.inc16:
    %10 = load i32, i32* %i6, align 4
    %inc17 = add i32 %10, 1
    store i32 %inc17, i32* %i6, align 4
    br label %for.cond7
```



# Case study [2] : non-parallelisable loop

```
for.cond:                                ; preds = %for.inc, %entry
%0 = load i32, i32* %i, align 4
%cmp = icmp ult i32 %0, 100
br i1 %cmp, label %for.body, label %for.end
```

```
for.body:                                ; preds = %for.cond
%call = call i8* @Znwm(i64 16) #8
%1 = bitcast i8* %call to %struct.list_node*
call void @ZN9list_nodeC2Ev(%struct.list_node* %1)
br label %invoke.cont
```

```
invoke.cont:                             ; preds = %for.body
%2 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%next = getelementptr inbounds %struct.list_node, %struct.list_node* %2, i32 0, i32 1
store %struct.list_node* %1, %struct.list_node** %next, align 8
%3 = load i32, i32* %i, align 4
%4 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%value = getelementptr inbounds %struct.list_node, %struct.list_node* %4, i32 0, i32 0
store i32 %3, i32* %value, align 8
%5 = load %struct.list_node*, %struct.list_node** %list_it, align 8
%next1 = getelementptr inbounds %struct.list_node, %struct.list_node* %5, i32 0, i32 1
%6 = load %struct.list_node*, %struct.list_node** %next1, align 8
store %struct.list_node* %6, %struct.list_node** %list_it, align 8
br label %for.inc
```

```
for.inc:                                 ; preds = %invoke.cont
%7 = load i32, i32* %i, align 4
%inc = add i32 %7, 1
store i32 %inc, i32* %i, align 4
br label %for.cond
```

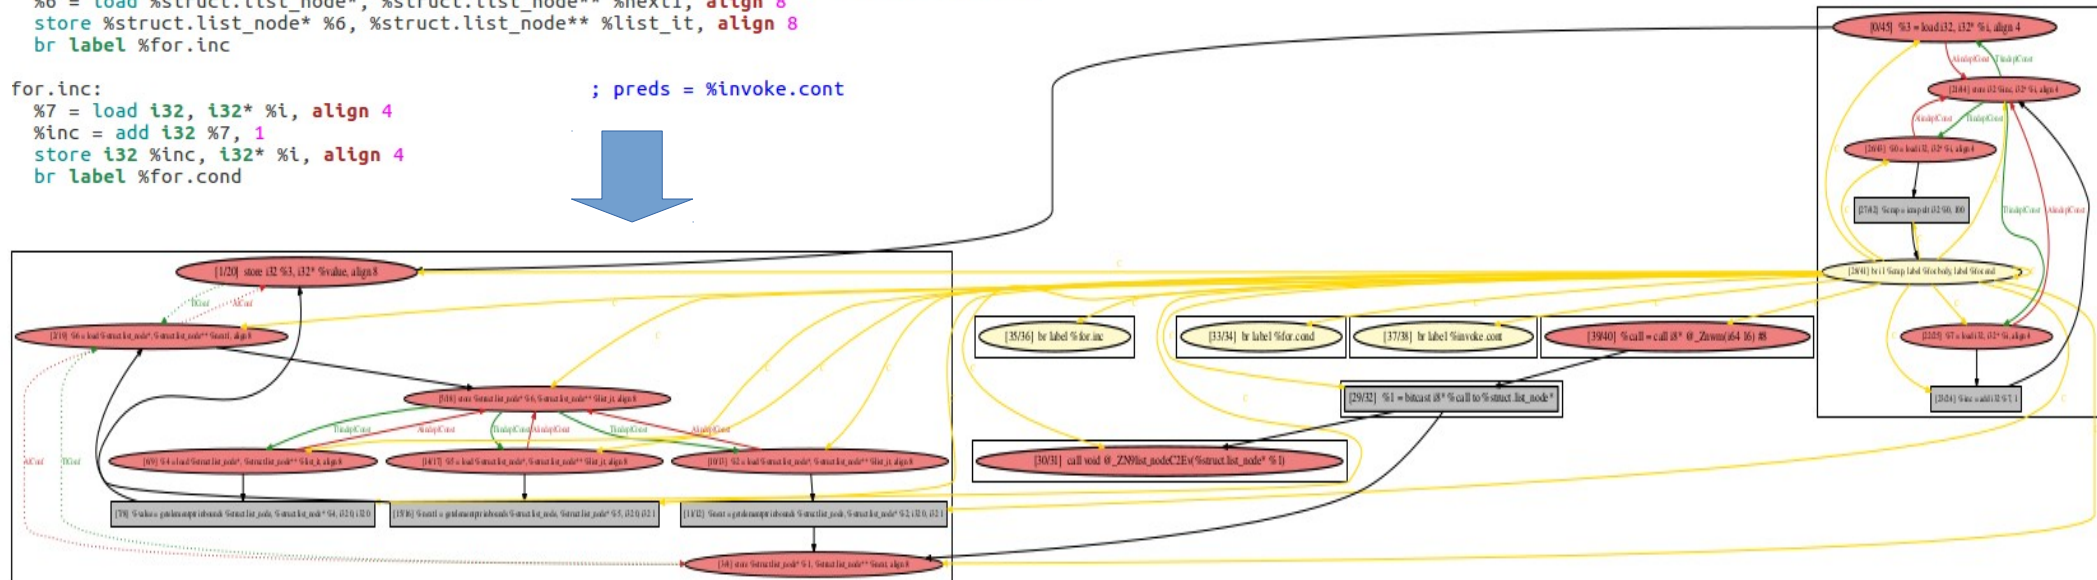
```
list_node_t begin;
list_node_t* list_it;
```

```
list_it = &begin;
for (unsigned int i = 1; i < size; i++) {

    list_it->next = new list_node_t;

    list_it->value = i;

    list_it = list_it->next;
}
```



# Case study [2] : non-parallelisable loop

```
list_node_t begin;  
list_node_t* list_it;
```

```
list_it = &begin; Iterator
```

```
for (unsigned int i = 1; i < size; i++) {
```

```
    list_it->next = new list_node_t;
```

```
    list_it->value = i;
```

```
    list_it = list_it->next;
```

```
}
```

Payload

```
list_node_t begin;  
list_node_t* list_it;
```

```
list_it = &begin;
```

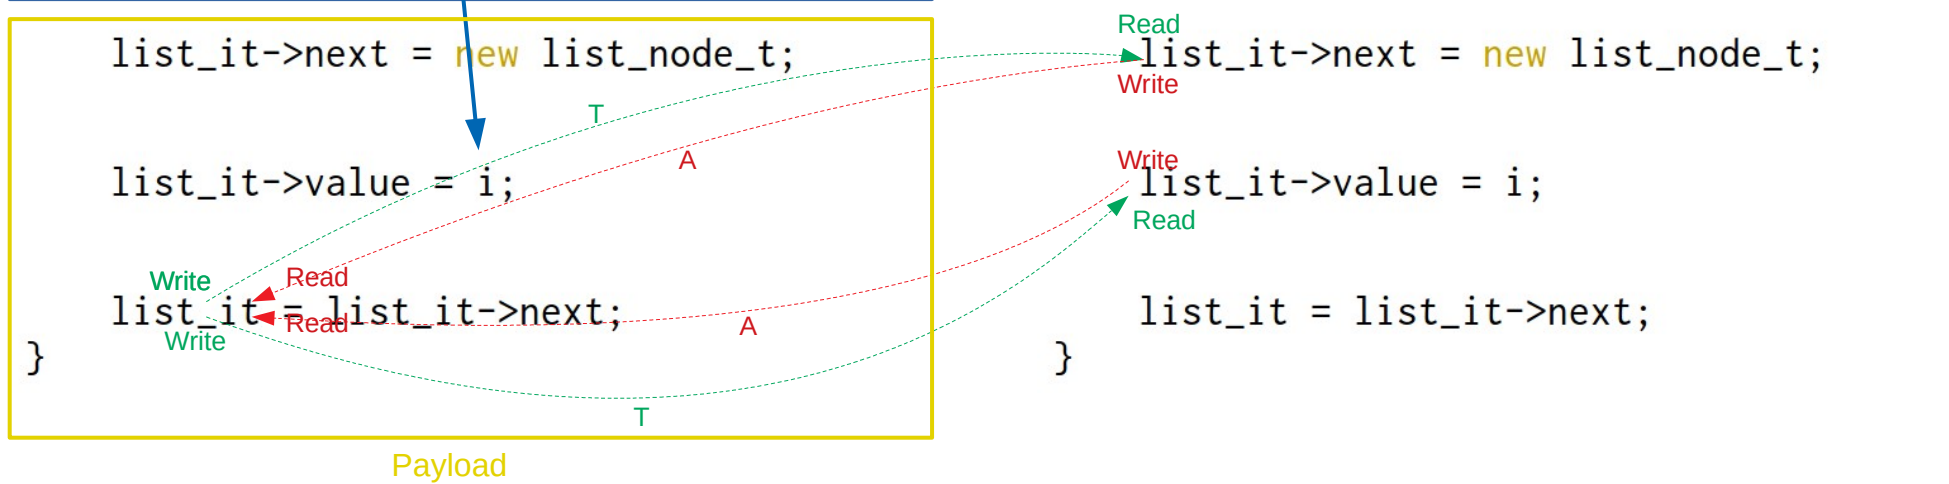
```
for (unsigned int i = 1; i < size; i++) {
```

```
    list_it->next = new list_node_t;
```

```
    list_it->value = i;
```

```
    list_it = list_it->next;
```

```
}
```

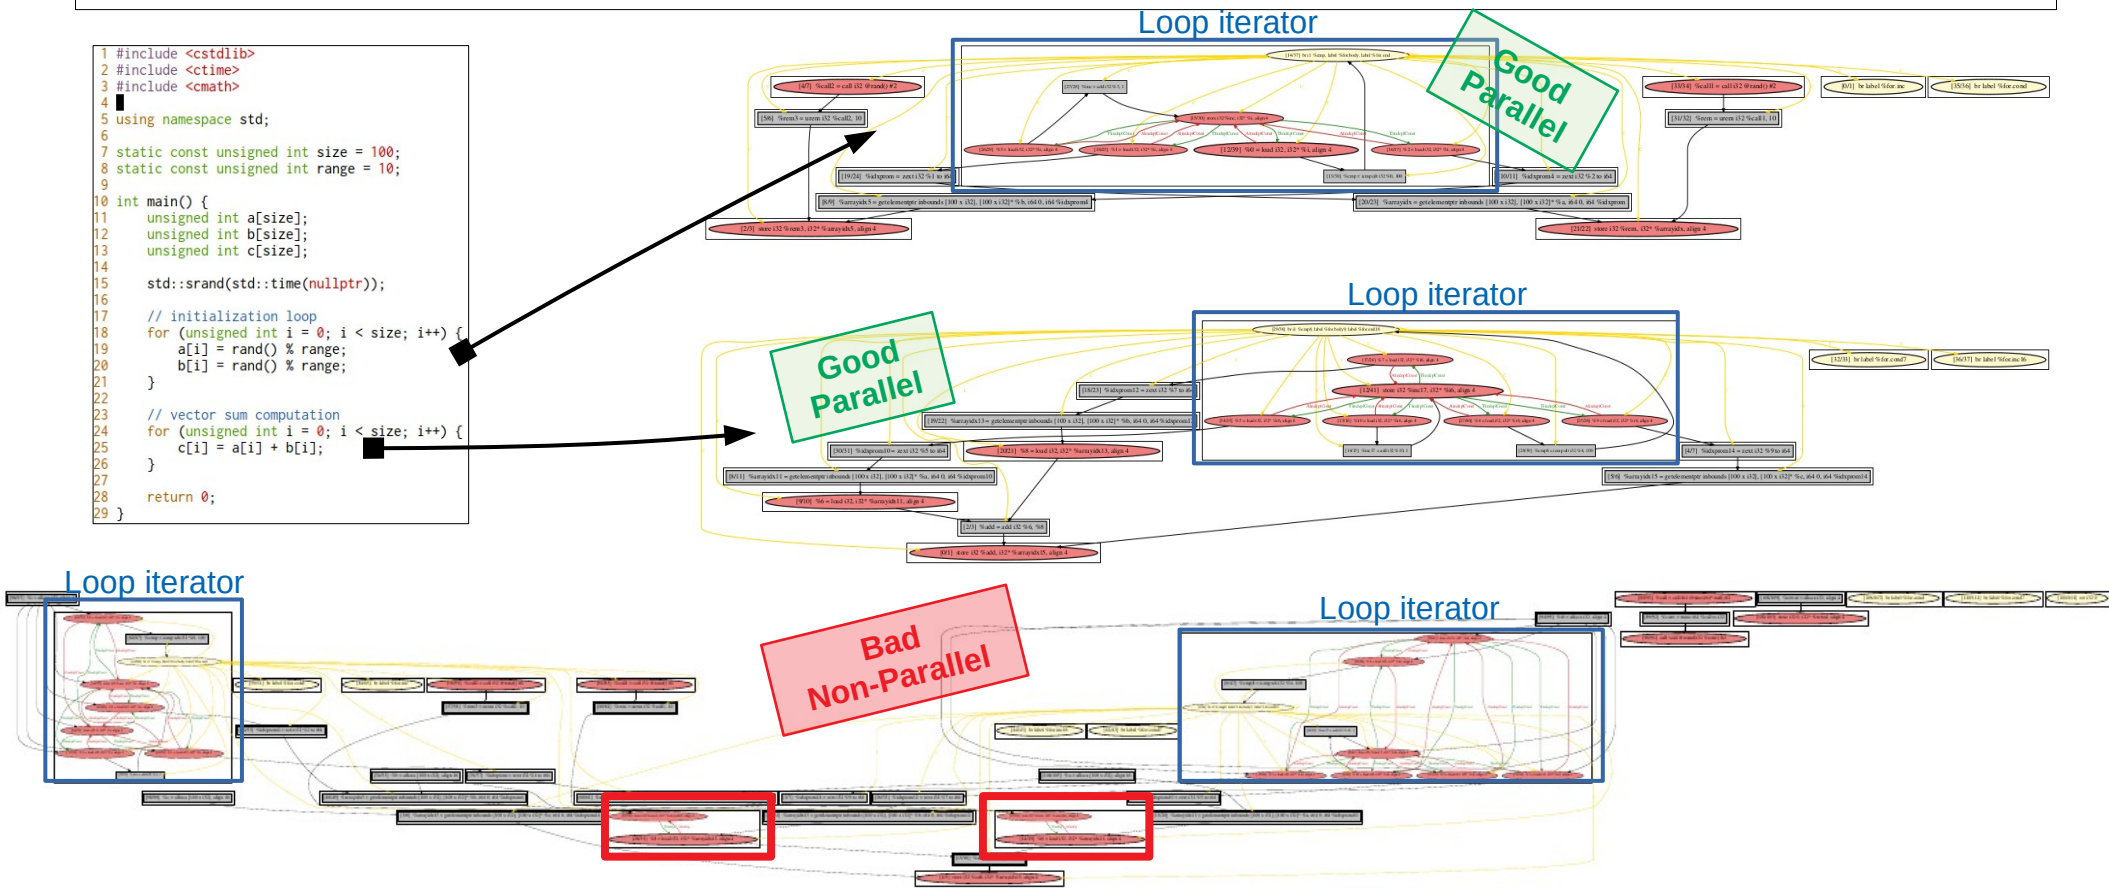


# General work observation

- To judge about true program parallelisability, we have to disassemble the program (its program dependence graph) to the smallest finest-grain pieces
- As it appears at the moment, the true indicator of the loop parallelisability is the absence of strongly connected components (besides the iterator one) of the size, greater than 1 instruction: there are cycles present in such strongly connected components, which imply dependence between instructions tying up/entangling loops and preventing them from parallelisation
- According to code comments, there is a work going on in LLVM on building `llvm::DependenceGraph` and breaking its edges with loop splitting and peeling.

# General work observation : manifestation [1]

- To judge about true program parallelisability, we have to disassemble the program (its program dependence graph) to the smallest finest-grain pieces



# General work observation : manifestation [1]