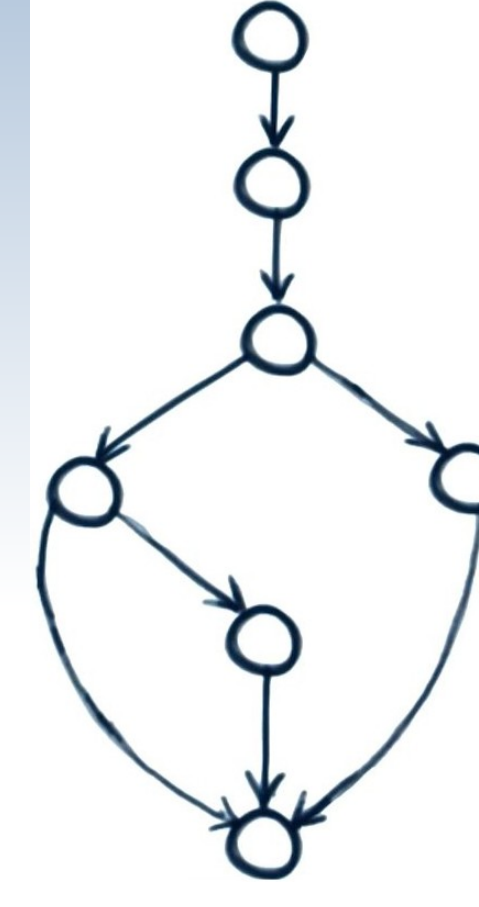


Project Overview

Despite decades of research into parallelising compiler technology, software parallelisation remains a largely manual task where the key resource is expert time. Typically, a programmer has to identify those program loops, which contribute the most to the ultimate program execution time, analyse if they are parallelisable and if so transform the program into a parallel form. All these tasks require a programmer to possess a set of extra skills and are time consuming. In this project we propose a **methodology and assistant tool [1]** which make better use of expert time by guiding their effort directly towards those loops, where the largest performance gains can be expected while keeping analysis and transformation effort at a minimum.

Our novel assistant tool is based on a machine learning (ML) model of loop parallelisability, which is captured by a vector of **static program features [2]**. Assistant takes a program profile as an input and provides a programmer with a ranking of all loops in a program based on their overall (worthwhile and feasible to parallelise) merit. The ranking function combines for each loop its potential contribution to the speedup (taken from a profile) and an estimated probability for its successful parallelisation (taken from a trained ML model). We use a set of 1415 SNU NAS Parallel Benchmark (NPB) loops as a data set for ML. Despite its limited size our model demonstrates a **prediction accuracy of greater than 90% [3]**. Being less conservative than a state-of-the-art compiler our assistant extends **the amount of discovered parallelism** in the SNU NPB suite from 81% to 96% [4]. That discovery comes at a price of 6,5% of cases being false positives.

We deploy our parallelisation assistant against separate SNU NAS benchmarks and show that our methodology achieves results comparable to those of expert programmers while requiring a programmer to examine fewer loops, i.e. it requires less expert time. On average, our assistant reduces programmer efforts required to reach an expert-level performance manually by 20% and **converges to the maximum achievable performance faster [5]**.

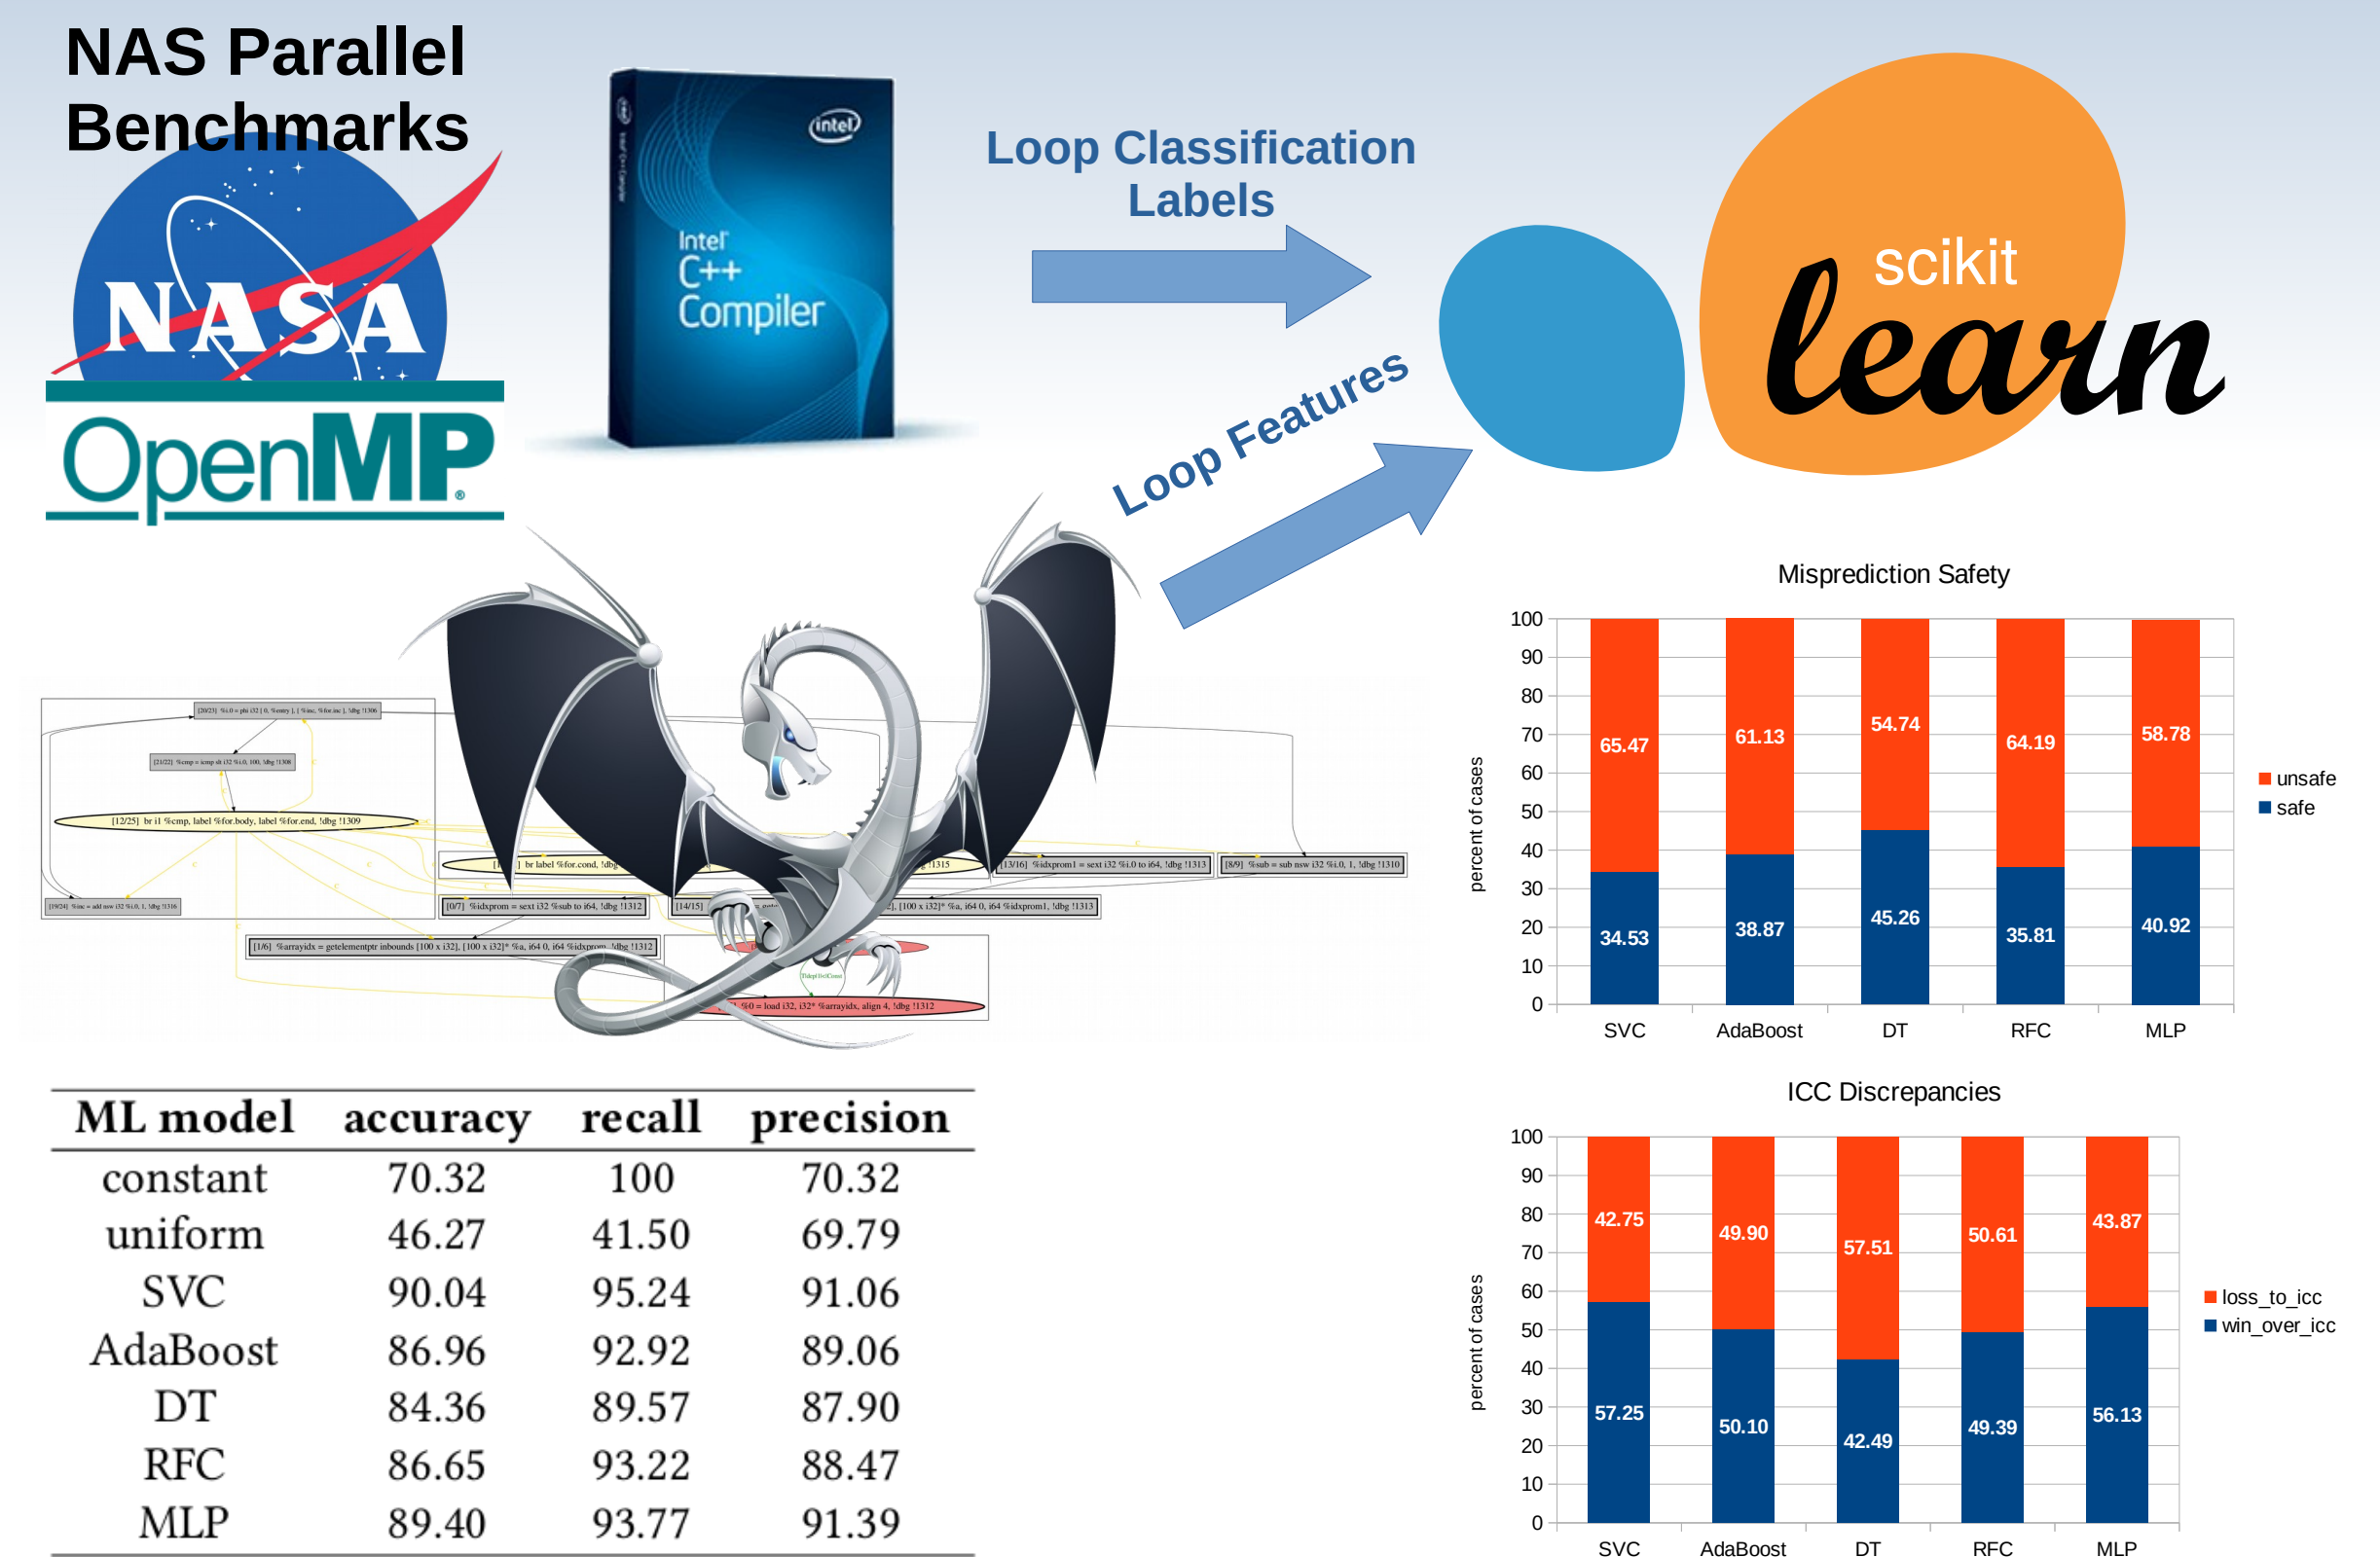


Software Quality Metrics

(the initial motivation for the project)

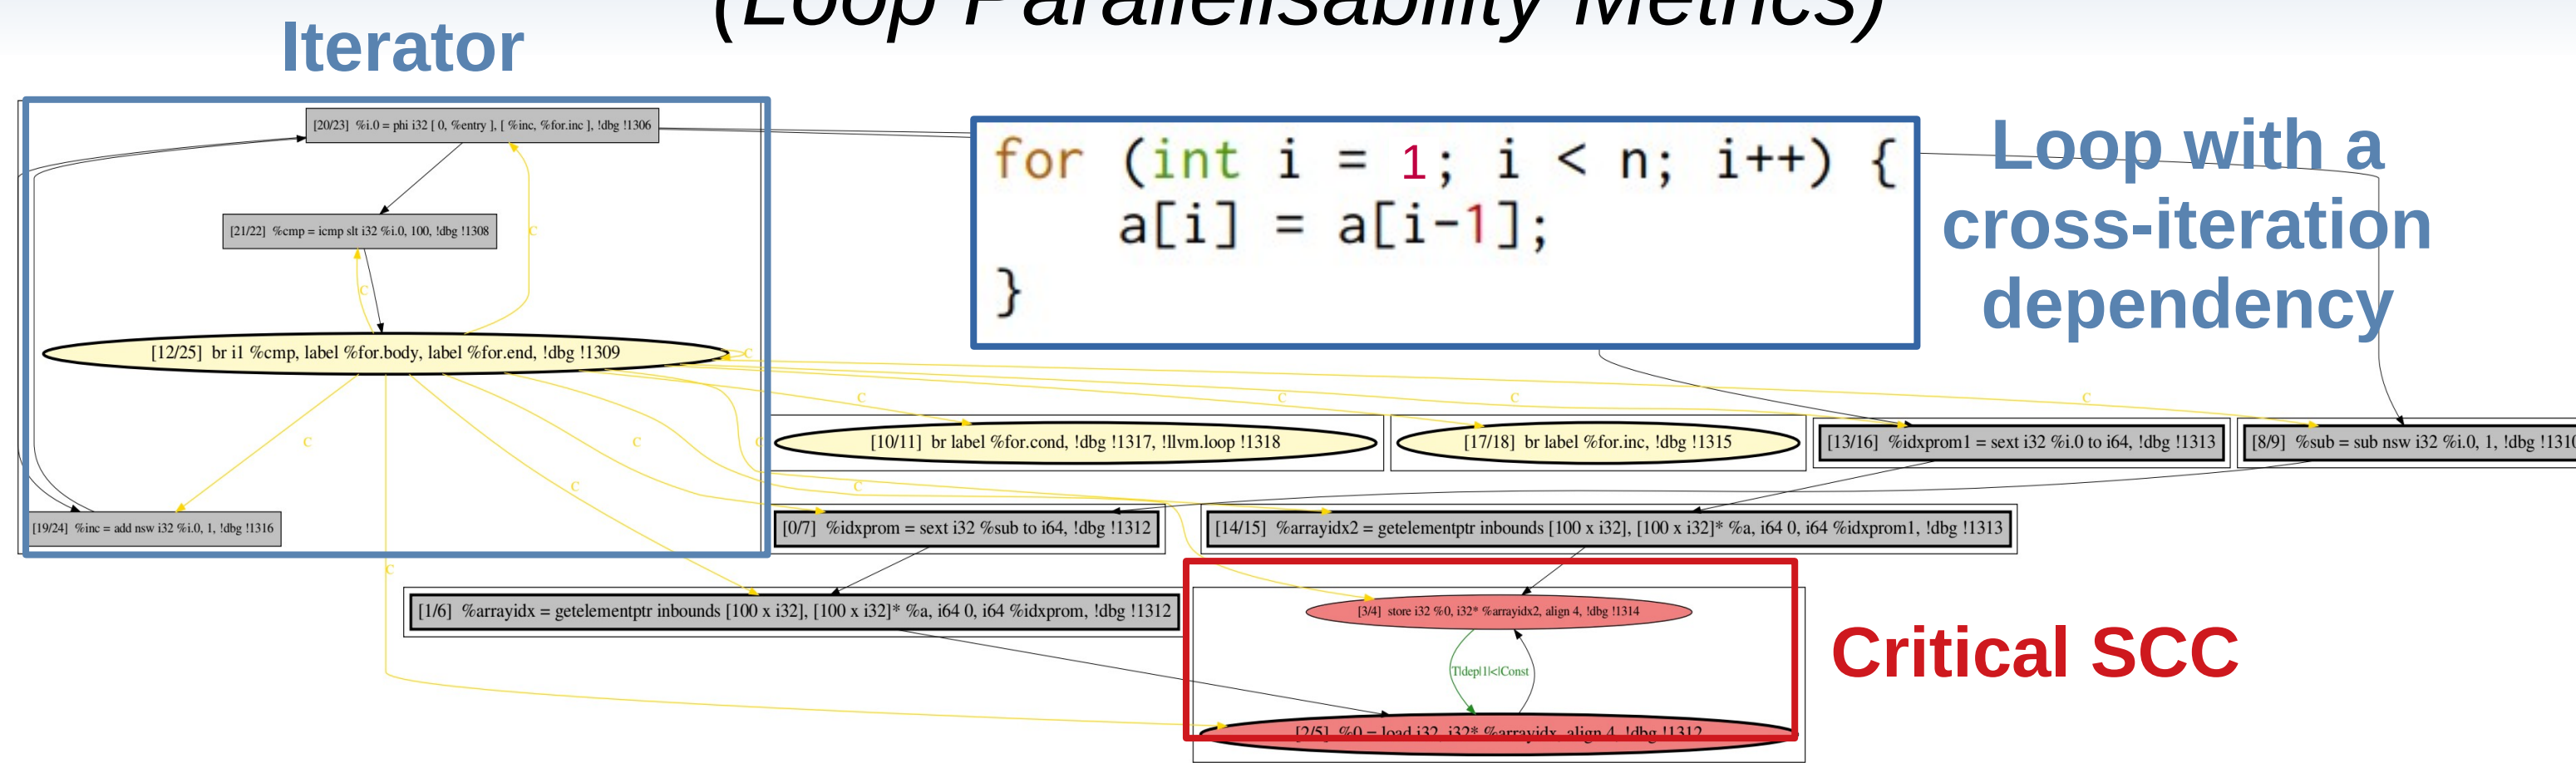
Cyclomatic Complexity (CC) (Thomas J. McCabe [1976]) is based on the CFG of the code section and represents the number of independent paths through it.

[3] Machine Learning Predictive Performance



[2] Machine Learning Loop Features

(Loop Parallelisability Metrics)

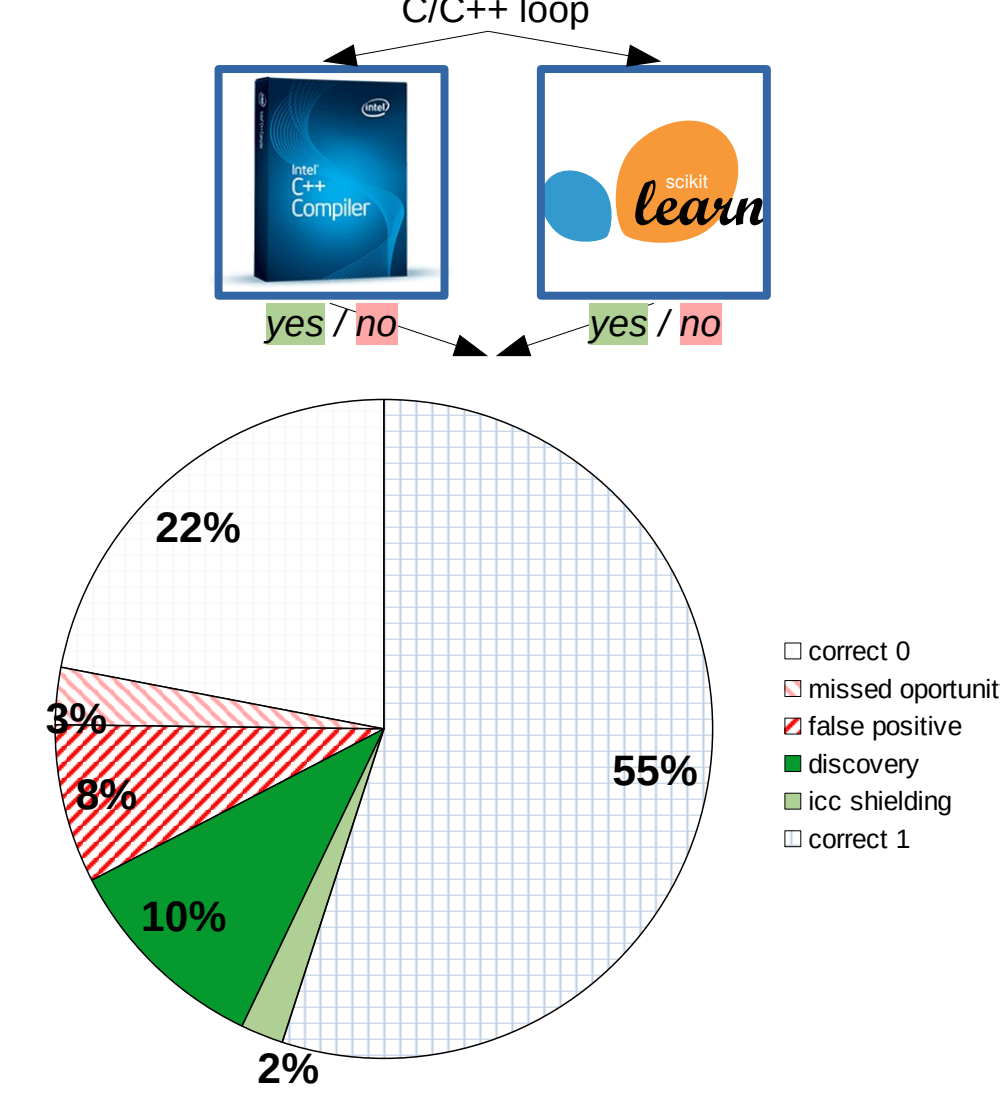


Software parallelisability metrics [74]

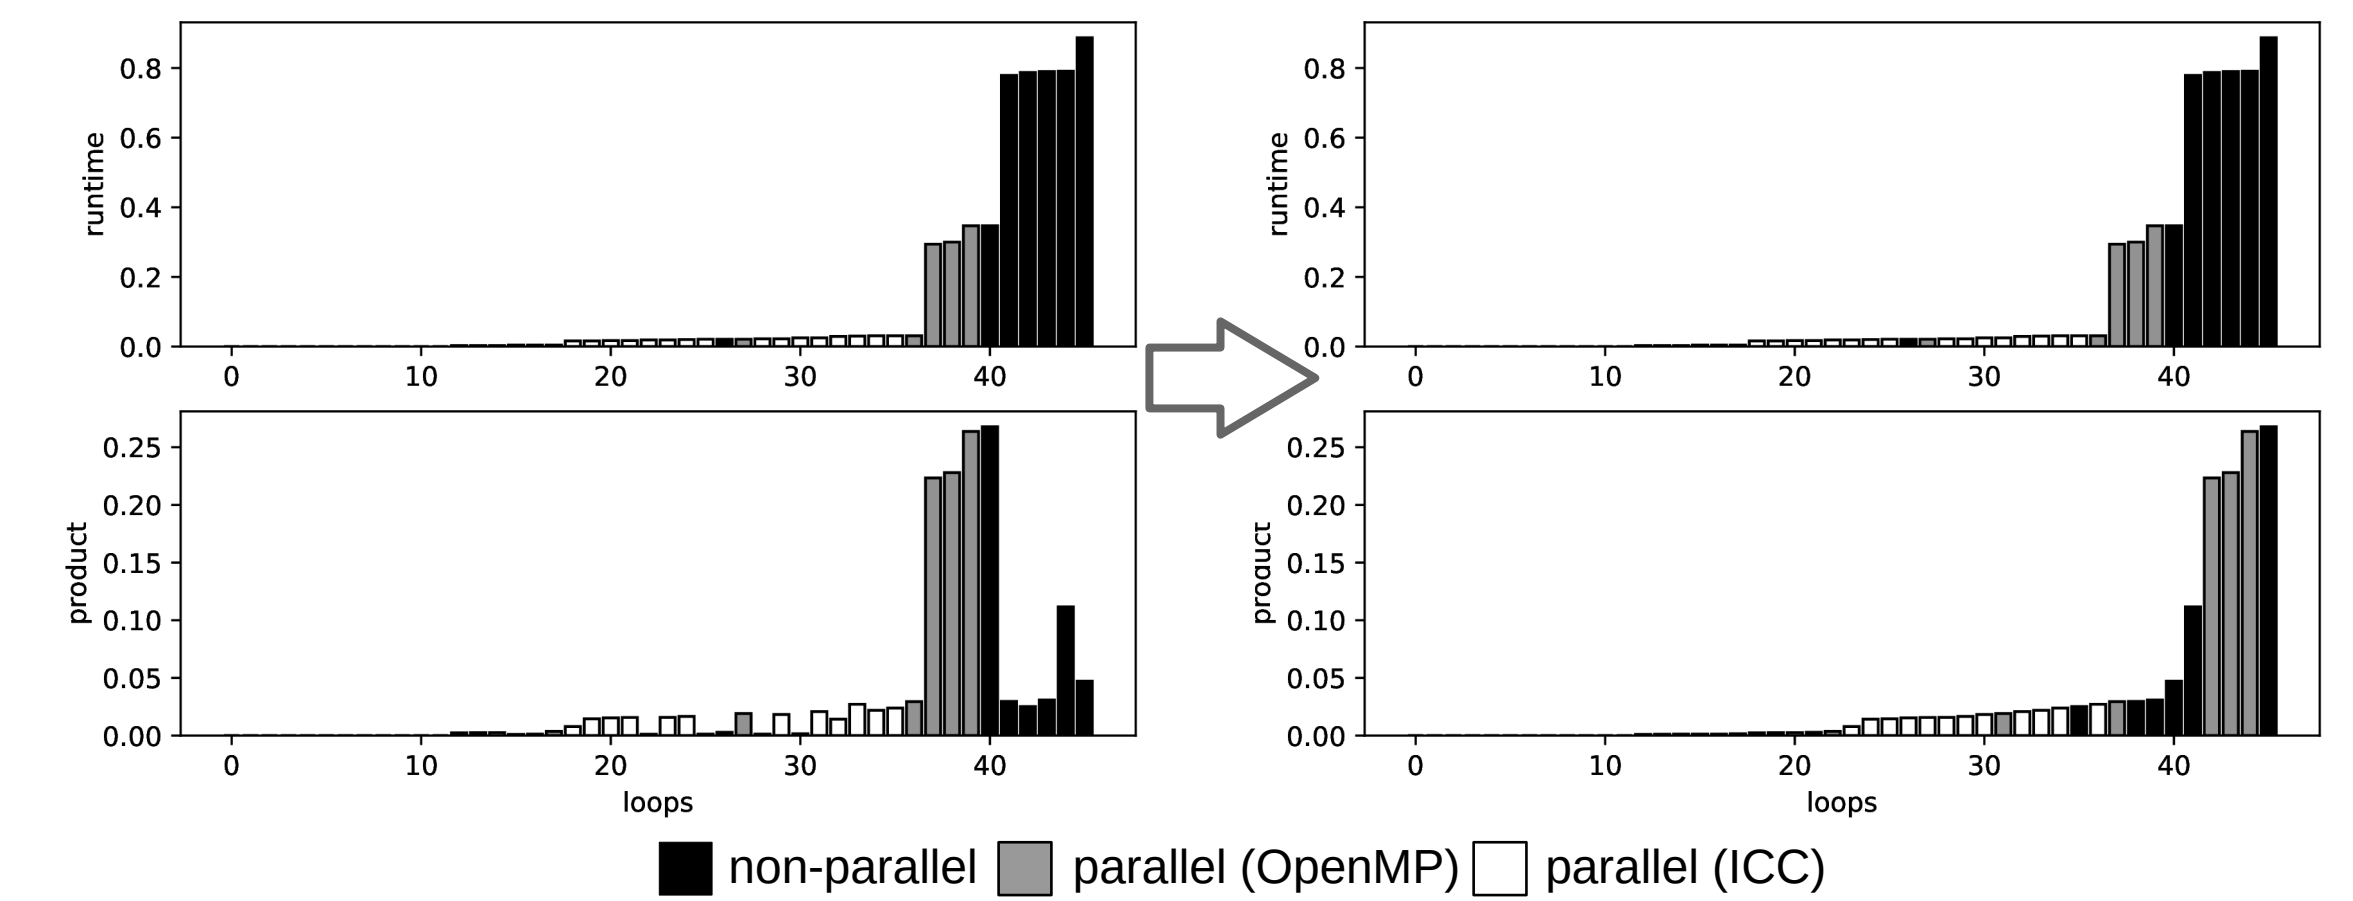
Metric Group	Metric	Metric Definition
Loop Proportions	Absolute Size	Number of LLVM IR instructions in a whole loop
	Payload Fraction	
	Proper SCCs Number	Number of SCCs with more than one LLVM IR instruction in a payload of a loop
	Critical Payload Fraction	
Loop Dependencies Number	Payload Dependencies Number	Number of PDG edges in a payload (True, Anti, Output and Total)
	Critical Payload Dependencies Number	
	Iterator/Payload Cohesion	Normalised number of edges between iterator and payload
Loop Cohesion	Critical/Regular Payload Cohesion	
	Call instructions count	The number of LLVM IR call instructions in a loop
Loop Instructions Nature	Branch instructions count	

[1] Methodology and assistant tool

[4] Parallelism Discovery Improvement

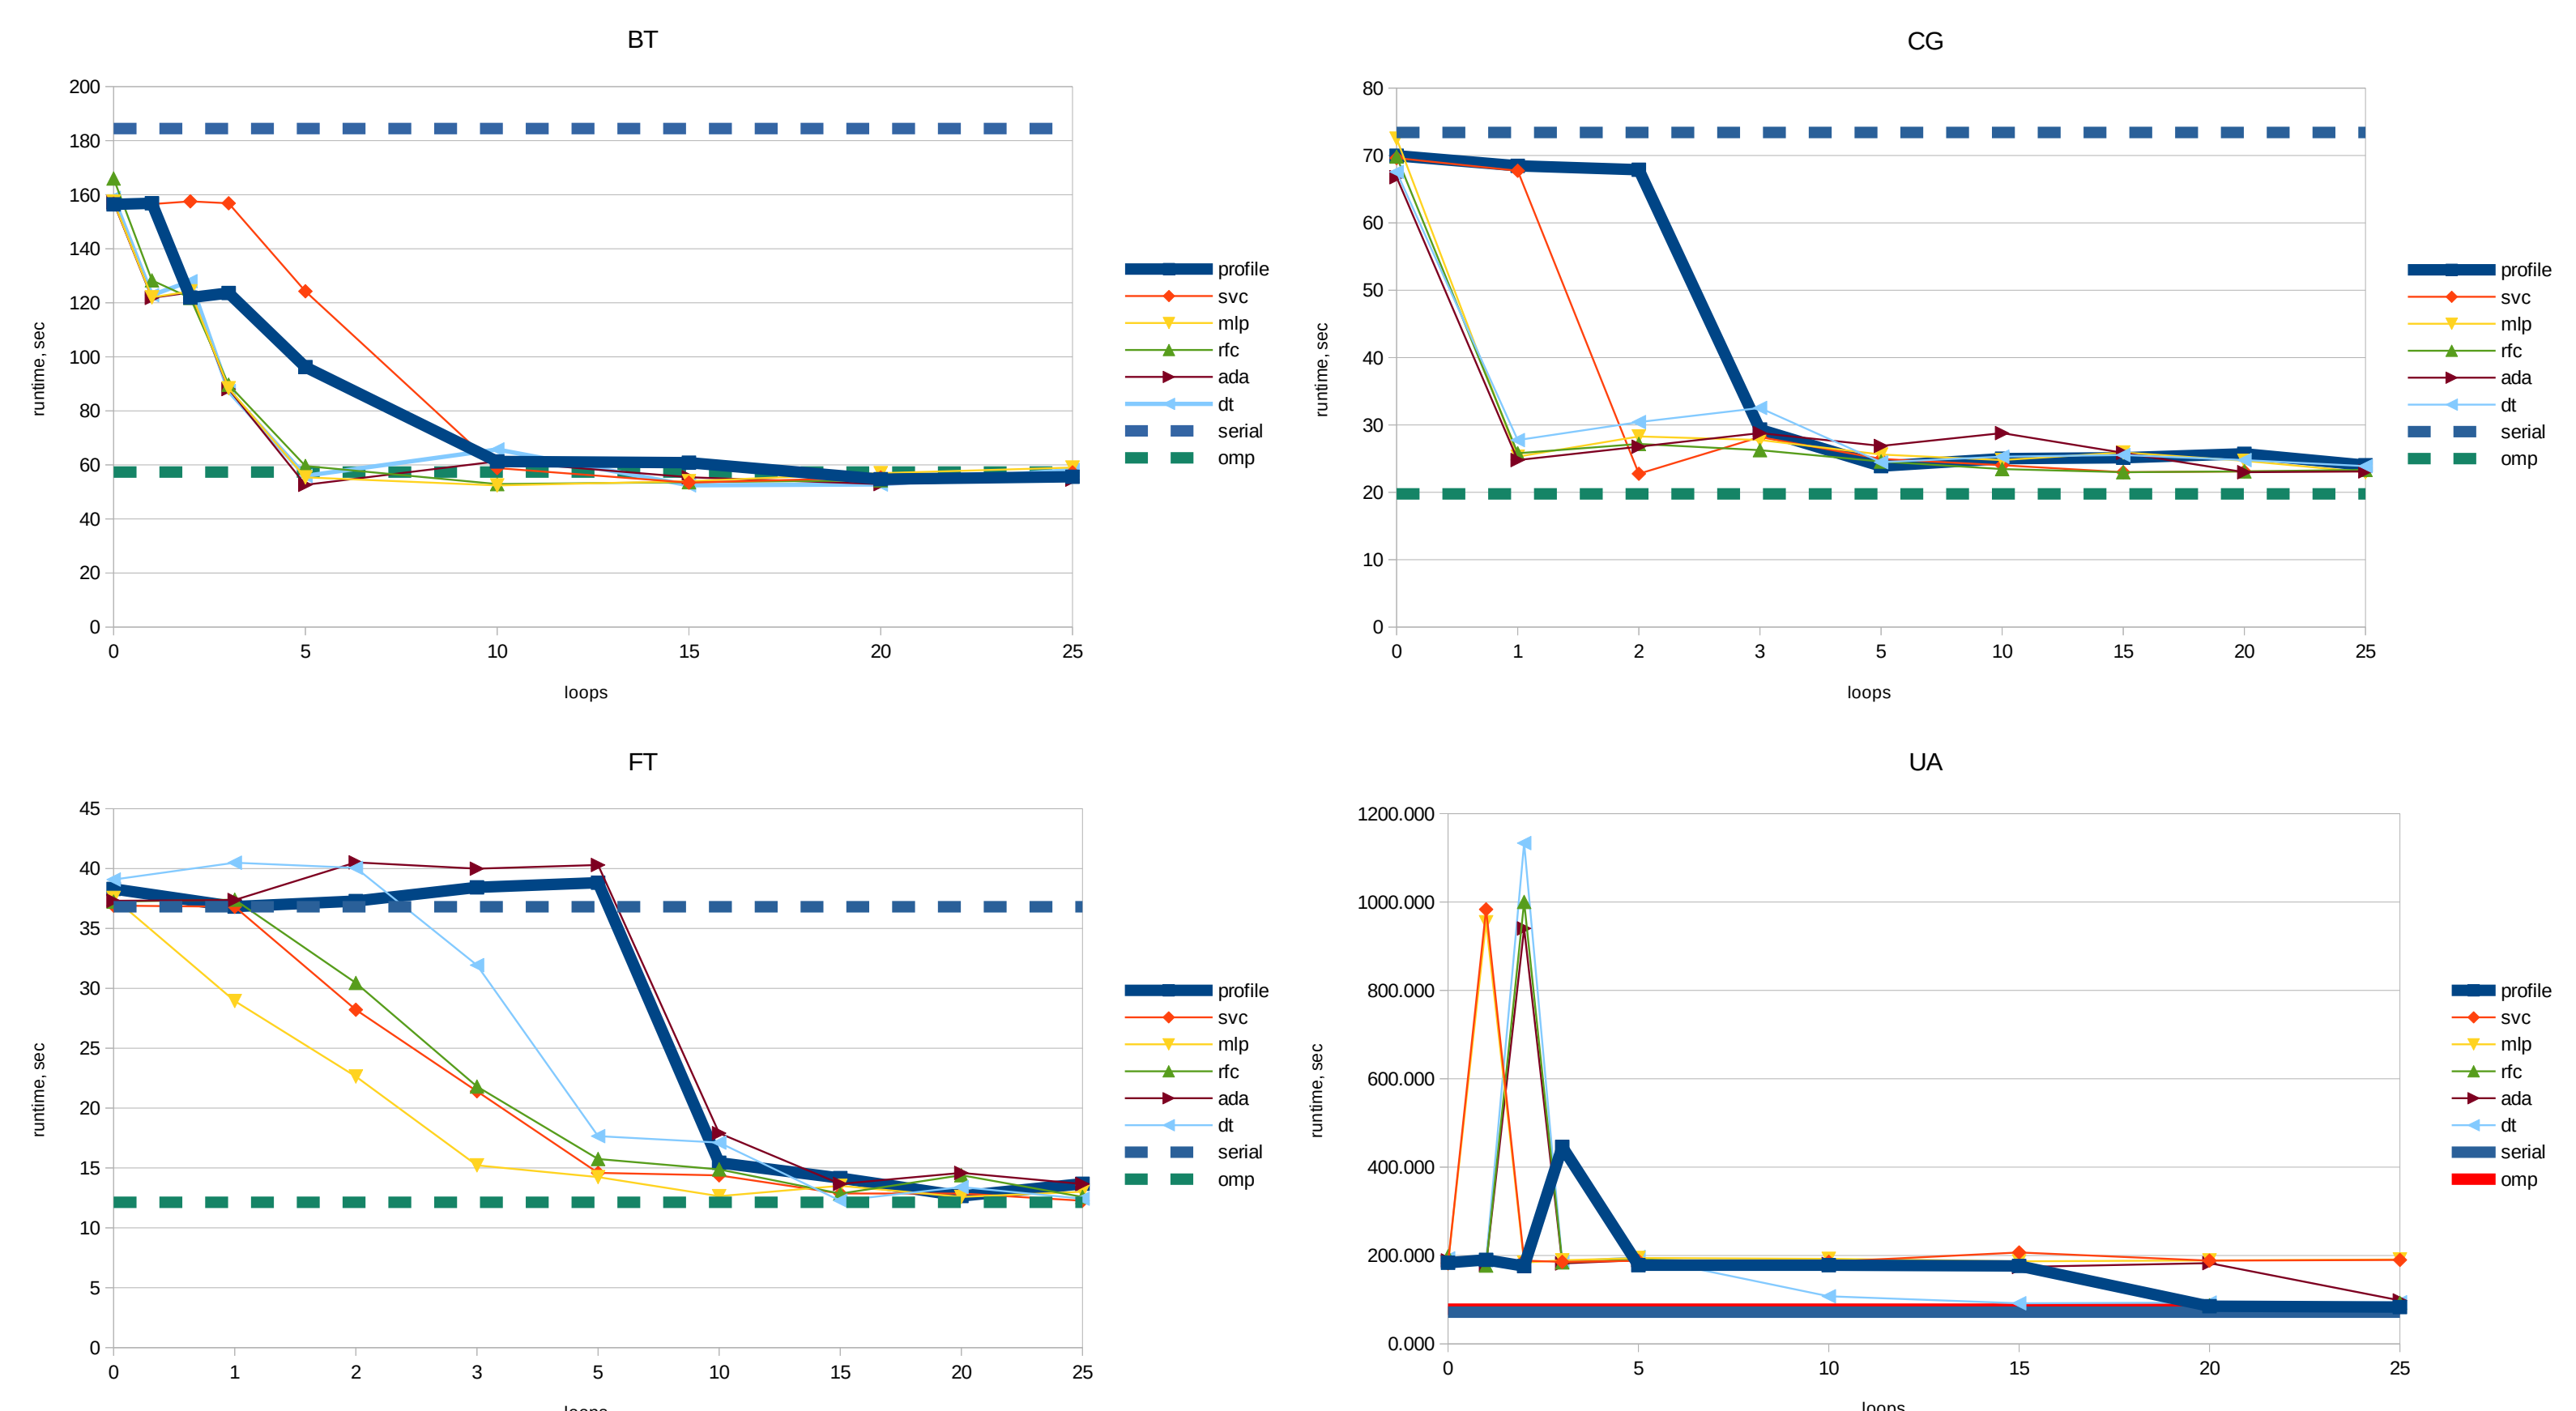


Loop Rankings: runtime order vs. product order



■ non-parallel ■ parallel (OpenMP) □ parallel (ICC)

[5] Performance convergence curves



EPSRC Centre for Doctoral Training in Pervasive Parallelism

icsa

Institute for Computing Systems Architecture



THE UNIVERSITY of EDINBURGH informatics